

```

404b30:    77 ae          ja     404ae0 <.omp_outlined..8+0x1c0>
404b32:    c5 f9 2e 04 25 48 60  vucomisd xmm0,QWORD PTR ds:0x406048
404b39:    40 00
404b3b:    75 02          jne    404b3f <.omp_outlined..8+0x21f>
404b3d:    7b a1          jnp    404ae0 <.omp_outlined..8+0x1c0>
404b3f:    c5 fb 11 54 24 28  vmovsd QWORD PTR [rsp+0x28],xmm2
404b45:    c5 fb 11 44 24 20  vmovsd QWORD PTR [rsp+0x20],xmm0
404b4b:    e8 60 d5 ff ff   call   4020b0 <log@plt>
404b50:    c5 fb 59 05 e8 14 00  vmulsd xmm0,xmm0,QWORD PTR [rip+0x14e8]      # 406040 <_I0_stdin_used+0x40>
404b57:    00
404b58:    c5 fb 5e 44 24 20  vdivsd xmm0,xmm0,QWORD PTR [rsp+0x20]
404b5e:    c5 f9 2e 04 25 48 60  vucomisd xmm0,QWORD PTR ds:0x406048
404b65:    40 00
404b67:    72 06          jb    404b6f <.omp_outlined..8+0x24f>
404b69:    c5 fb 51 c0          vsqrtsd xmm0,xmm0,xmm0
404b6d:    eb 05          jmp    404b74 <.omp_outlined..8+0x254>
404b6f:    e8 ac d5 ff ff   call   402120 <sqrt@plt>
404b74:    c5 fb 59 4c 24 10  vmulsd xmm1,xmm0,QWORD PTR [rsp+0x10]
404b7a:    c4 c1 7b 11 4e 10  vmovsd QWORD PTR [r14+0x10],xmm1
404b80:    41 c6 46 18 01   mov    BYTE PTR [r14+0x18],0x1
404b85:    c5 fb 59 44 24 28  vmulsd xmm0,xmm0,QWORD PTR [rsp+0x28]
404b8b:    c4 c1 7b 59 46 08  vmulsd xmm0,xmm0,QWORD PTR [r14+0x8]
404b91:    c4 c1 7b 58 06   vaddsd xmm0,xmm0,QWORD PTR [r14]
404b96:    49 8b 07          mov    rax,QWORD PTR [r15]
404b99:    c5 fb 11 04 e8   vmovsd QWORD PTR [rax+rpb*8],xmm0
404b9e:    41 c6 46 18 00   mov    BYTE PTR [r14+0x18],0x0
404ba3:    c4 c1 7b 10 46 10  vmovsd xmm0,QWORD PTR [r14+0x10]
404ba9:    31 c0          xor    eax,eax
404bab:    e9 3e fe ff ff   jmp    4049ee <.omp_outlined..8+0xce>
404bb0:    bf 68 81 40 00   mov    edi,0x408168
404bb5:    44 89 e6          mov    esi,r12d
404bb8:    e8 c3 d6 ff ff   call   402280 <__kmpc_for_static_fini@plt>
404bbd:    48 83 c4 38   add    rsp,0x38
404bc1:    5b          pop    rbx
404bc2:    41 5c          pop    r12
404bc4:    41 5d          pop    r13
404bc6:    41 5e          pop    r14
404bc8:    41 5f          pop    r15
404bca:    5d          pop    rbp
404bcb:    c3          ret
404bcc:    48 89 c7          mov    rdi,rax
404bcf:    e8 0c 00 00 00   call   404be0 <__clang_call_terminate>
404bd4:    48 89 c7          mov    rdi,rax
404bd7:    e8 04 00 00 00   call   404be0 <__clang_call_terminate>
404bdc:    0f 1f 40 00   nop    DWORD PTR [rax+0x0]

```

x86_64 executable binaries in action

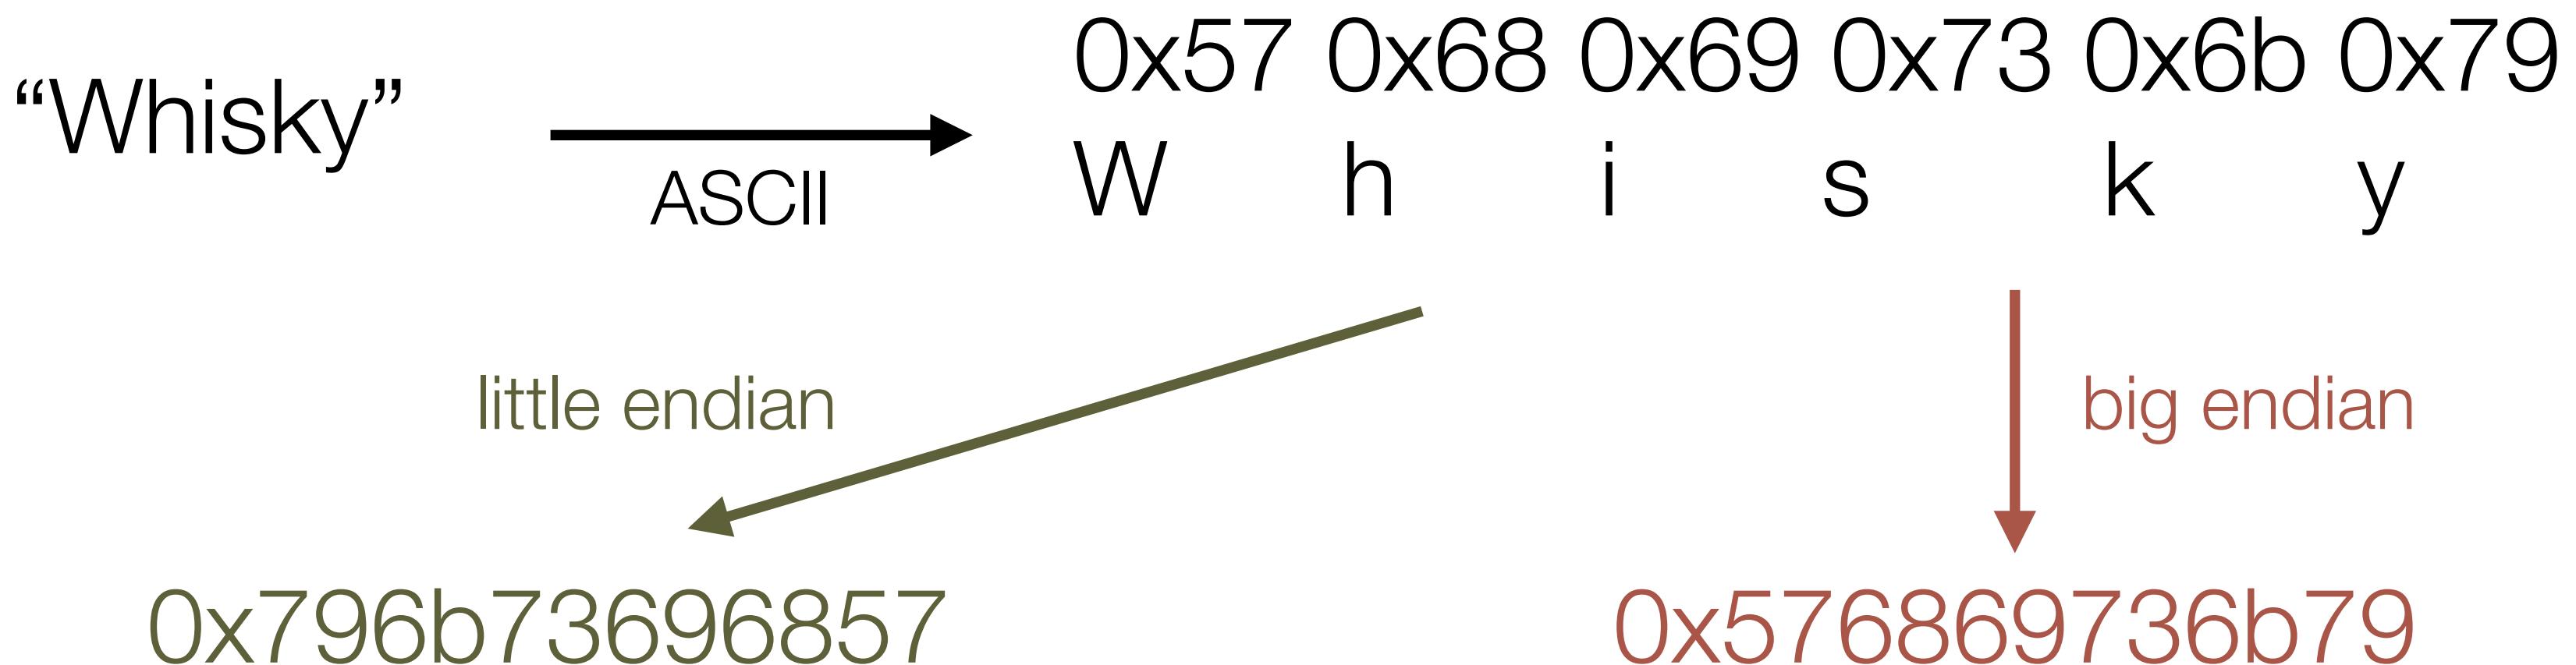
Introduction to executable binaries
Antonin Portelli

- Memory segmentation
- The stack
- Your best friend the debugger
- Let's play with stack frames
- The Executable and Linkable Format (ELF)
- Hello world again
- Wrap up

Memory segmentation

Endianness

- **Endianness** defines the order in which bytes are written in memory. Bytes not bits!
- Big endian: weakest bytes first “left to right” (e.g. PowerPC architecture)
- Little endian: strongest bytes first “right to left” (e.g. x86, x86_64 and ARM)

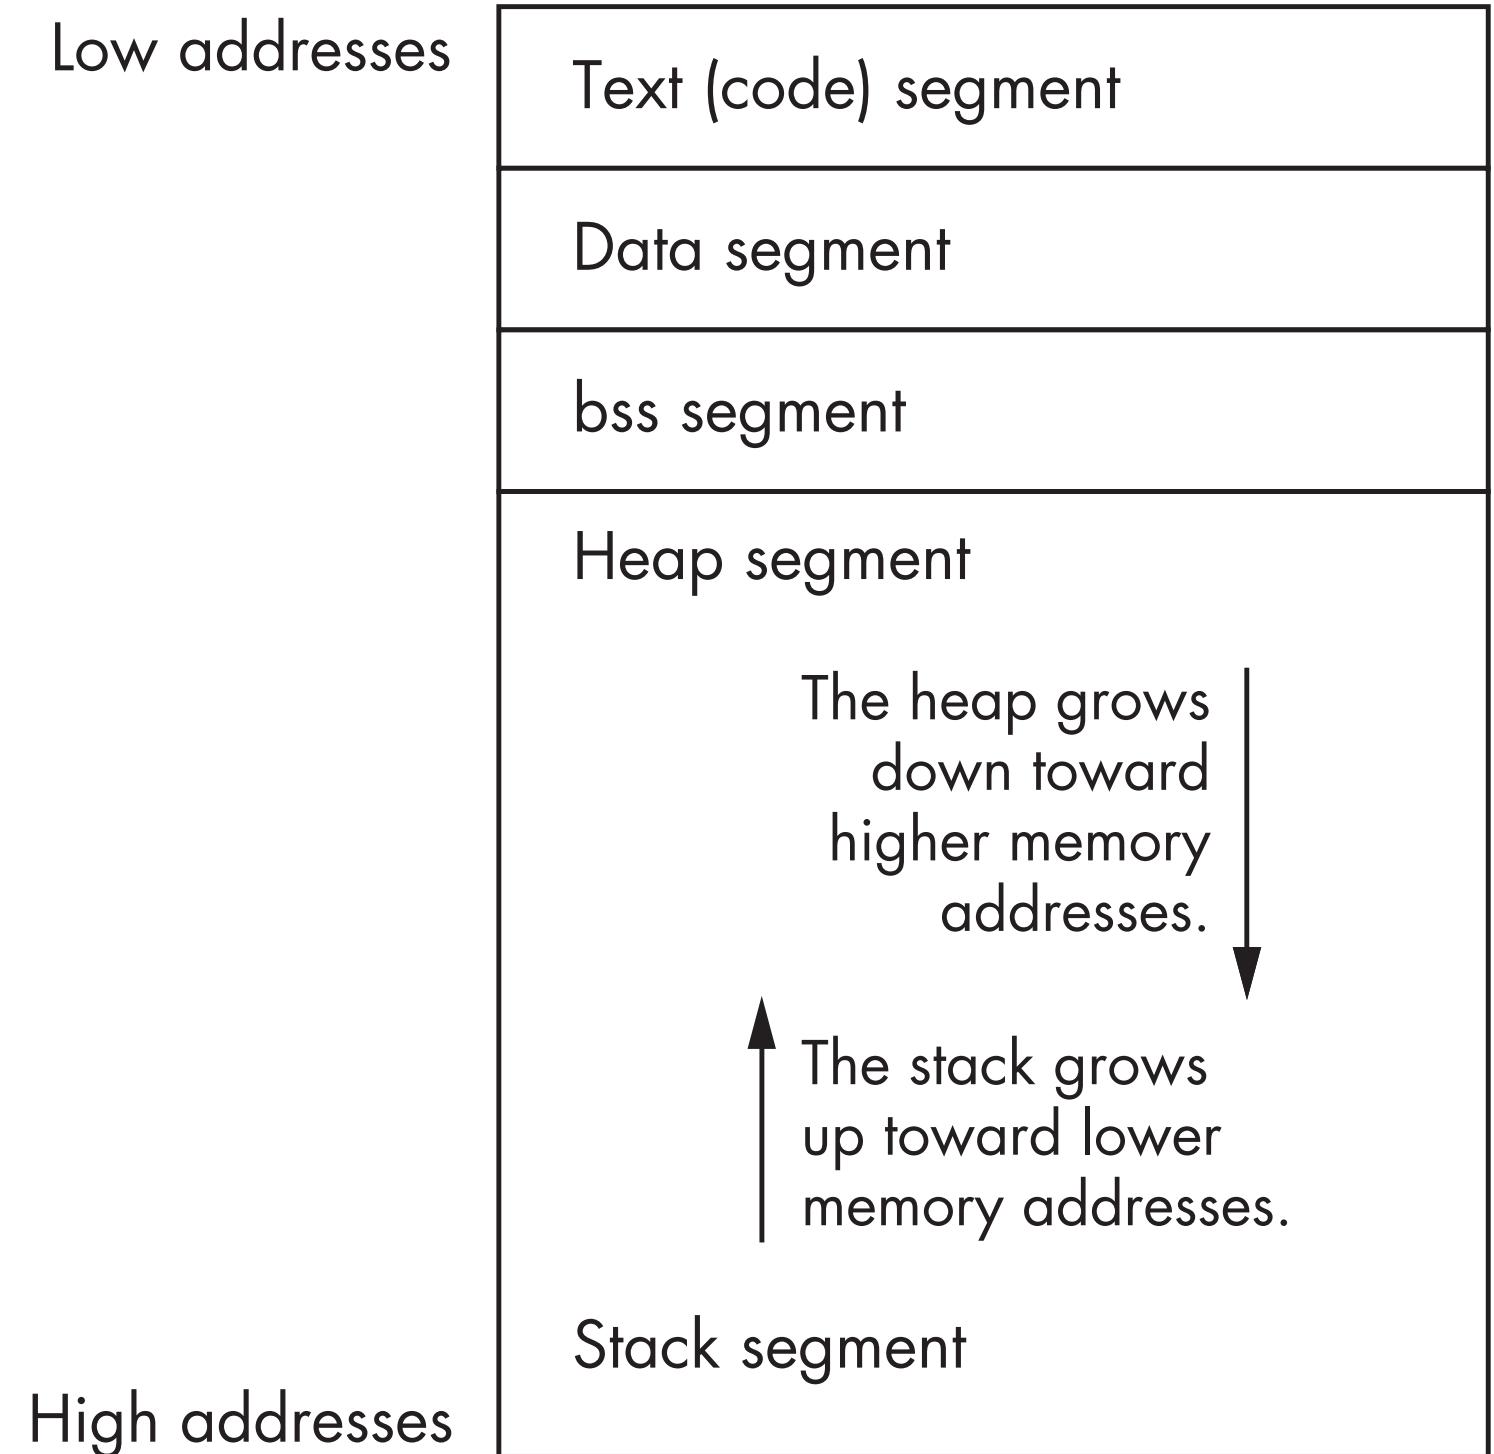


Memory segments

- A program memory is separated in mainly **5 segments**

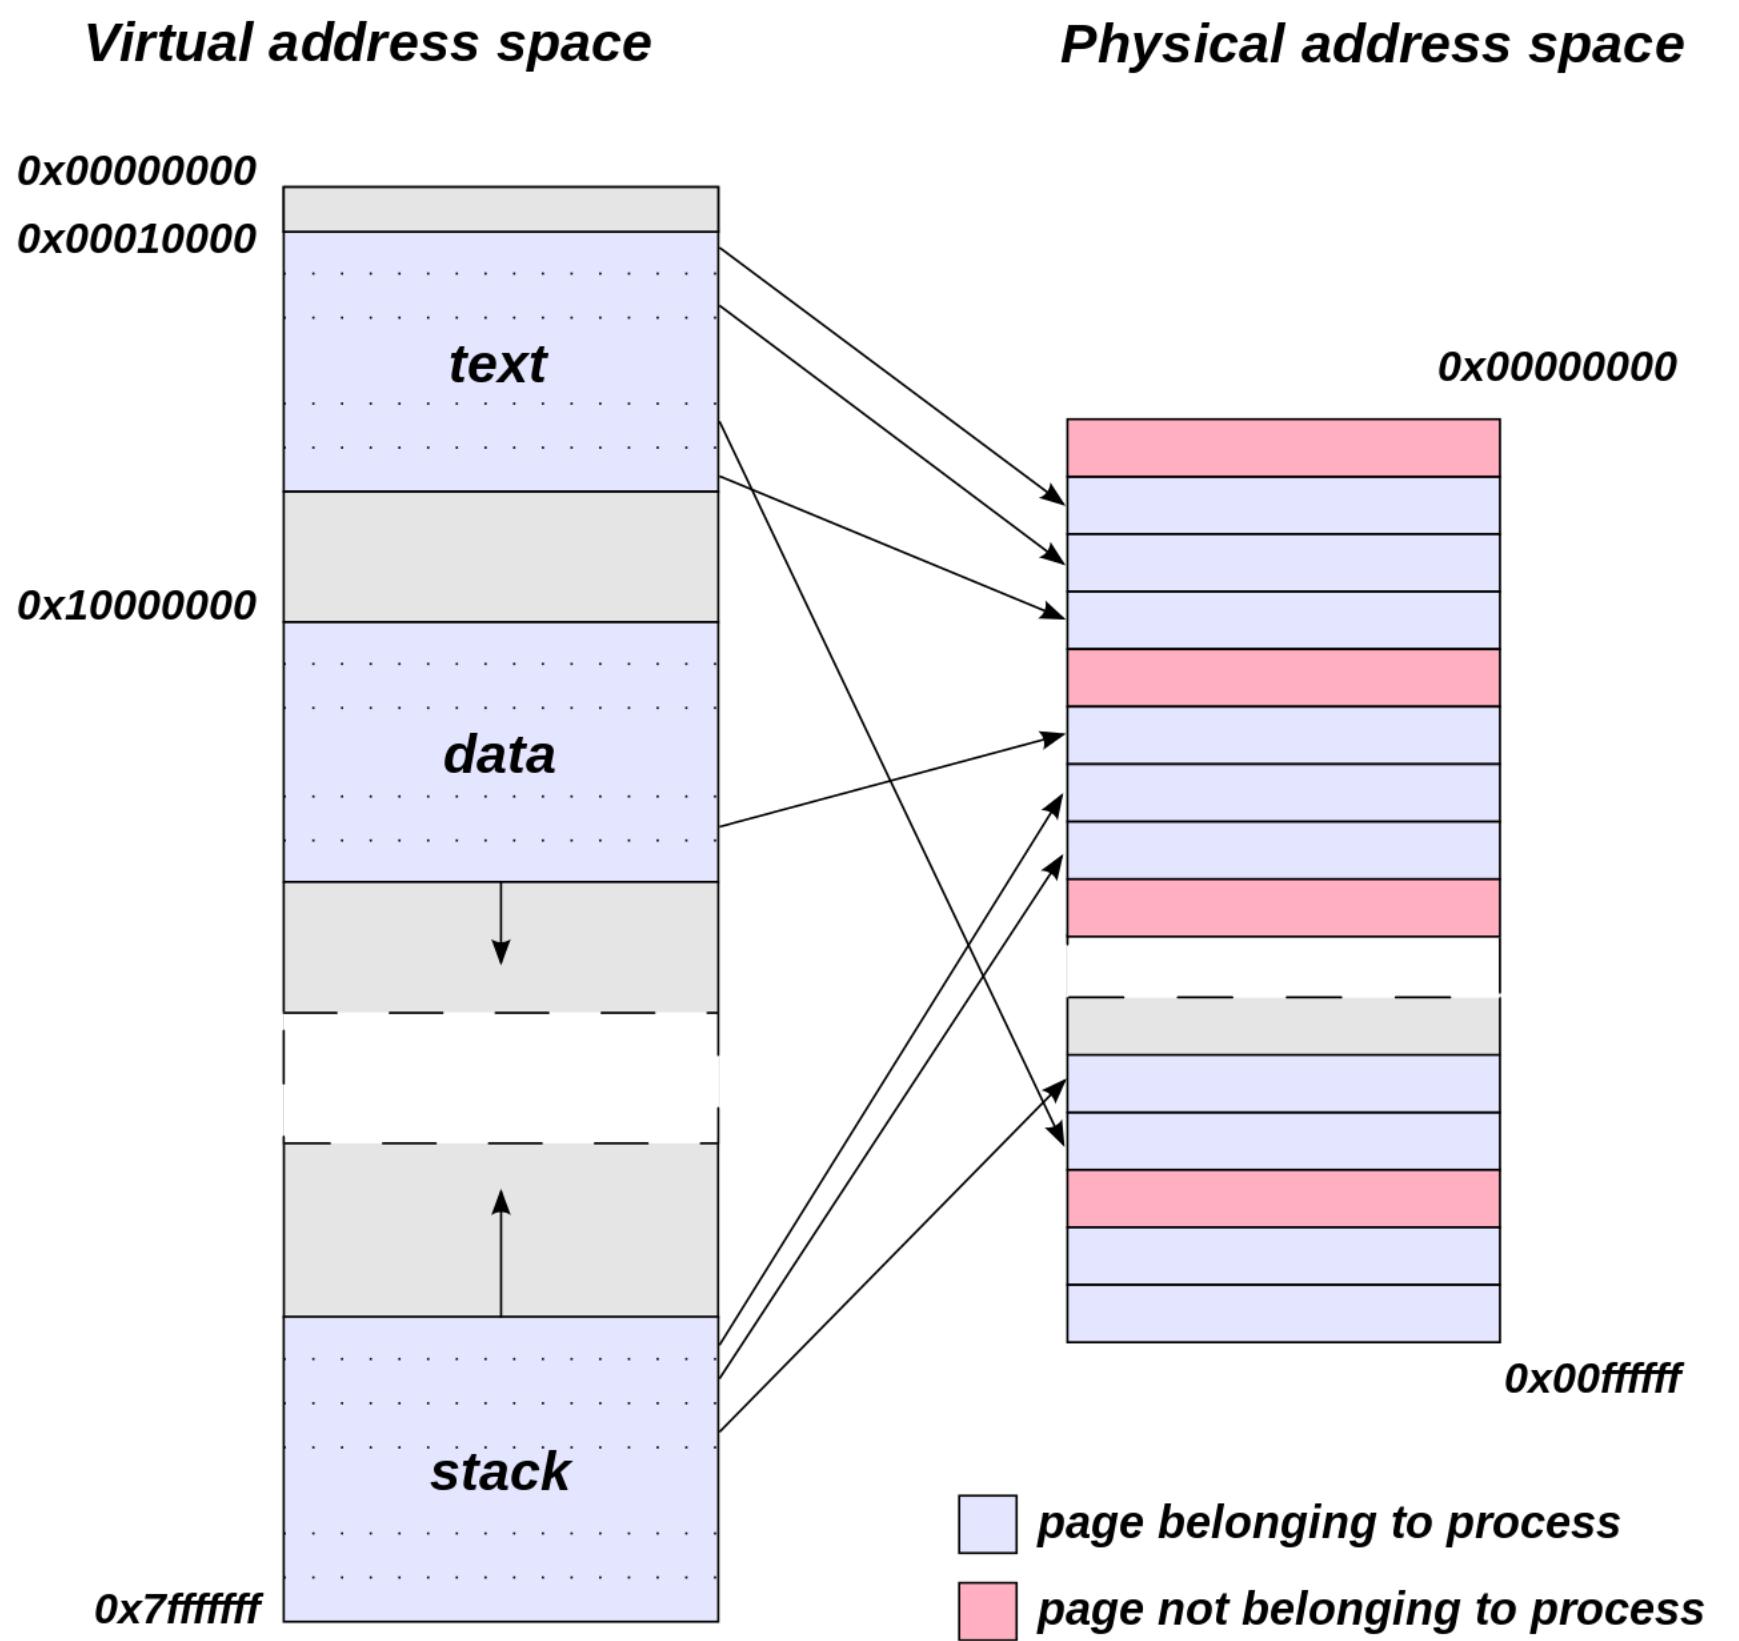
From low to high addresses

- 1- Text segment: program code
- 2- Data: initialised static variables
- 3- BSS: uninitialised static variables
- 4- Heap: global variable space
- 5- Stack: local “scratch pad”



Virtual memory in modern OSs

- Program address space is **not physical**
- It **may not be contiguous** in physical memory
- **Fragmented** in pages (typically 4 kB)
- Mapping to physical memory (page tables) maintained by the **kernel**



Text (code) segment

- Contains the program code in binary form
- Each instruction is code using an **opcode** and binary representation of the operands
- x86 has **variable size** instructions!

```
$ objdump -d -M intel ./hello-asm

...
48 c7 c0 01 00 00 00    mov    rax,0x1
48 c7 c7 01 00 00 00    mov    rdi,0x1
48 8b 75 d0              mov    rsi,QWORD PTR [rbp-0x30]
48 c7 c2 0d 00 00 00    mov    rdx,0xd
0f 05                   syscall
...
```

- This segment also contains constants

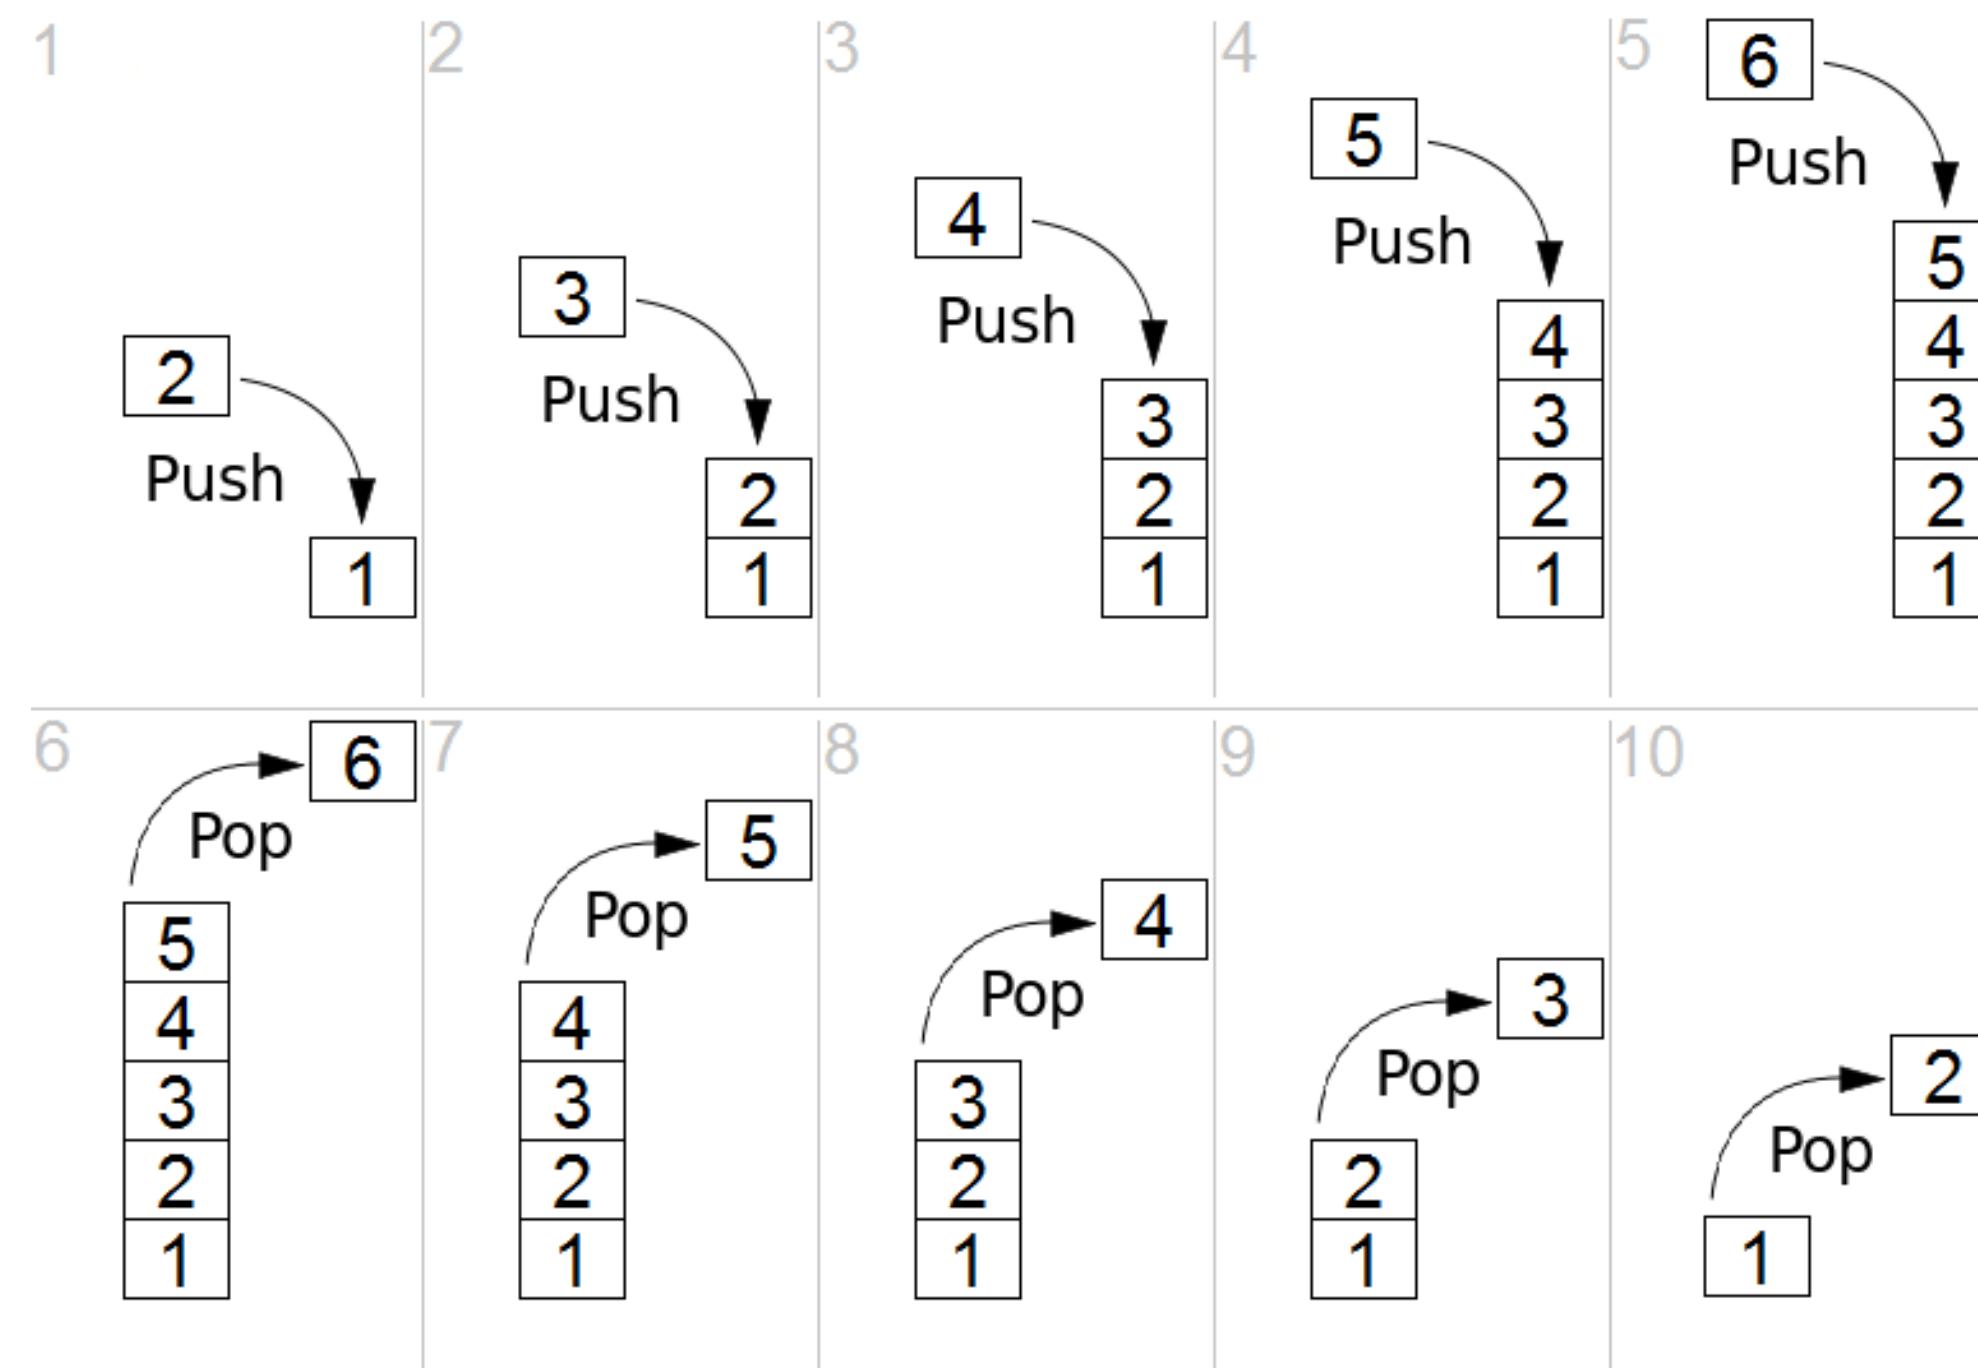
Heap & stack

- Heap: **extensible** segment meant to store persistent data
- The programs **needs to request** space on the heap, and needs to declare when the space isn't needed anymore
- In C one uses `malloc` & `free` to get heap space
In C++ it is done with the `new` & `delete` operators
- Stack: “**scratch pad**” for the program
- Also extensible, but with a **much simpler structure**
- Allocations are generally managed by the compiler

The stack

Stack generalities

- The stack is a FILO container. **First In Last Out**
- Accessed through **push** & **pop** operations



Stack manipulation in ASM

- Two registers are used to track the stack

rbp: base pointer (bottom of the stack address)

rsp: stack pointer (top of the stack address)



Stack manipulation in ASM

- **push src** : add src on the stack & update rsp

equivalent to (if src is 8B)

```
sub rsp, 8  
mov qword ptr [rsp], src
```

- **pop dest** : move top element to dest & update rsp

equivalent to (if dest is 8B)

```
mov dest, qword ptr [rsp]  
add rsp, 8
```

Function calls

- Functions can be called with **call address**
Next instruction address (**rip**) is pushed on the stack
Then execution jump to **address**
- Functions can return using **ret**
Pop top stack element as an address and jumps to it
- Argument passing is defined though a convention called
Application Binary Interface (ABI)
- For x86_64 on Linux, **System V AMD64 ABI**
<https://bit.ly/3kwiMxH>

System V AMD64 ABI

Figure 3.4: Register Usage

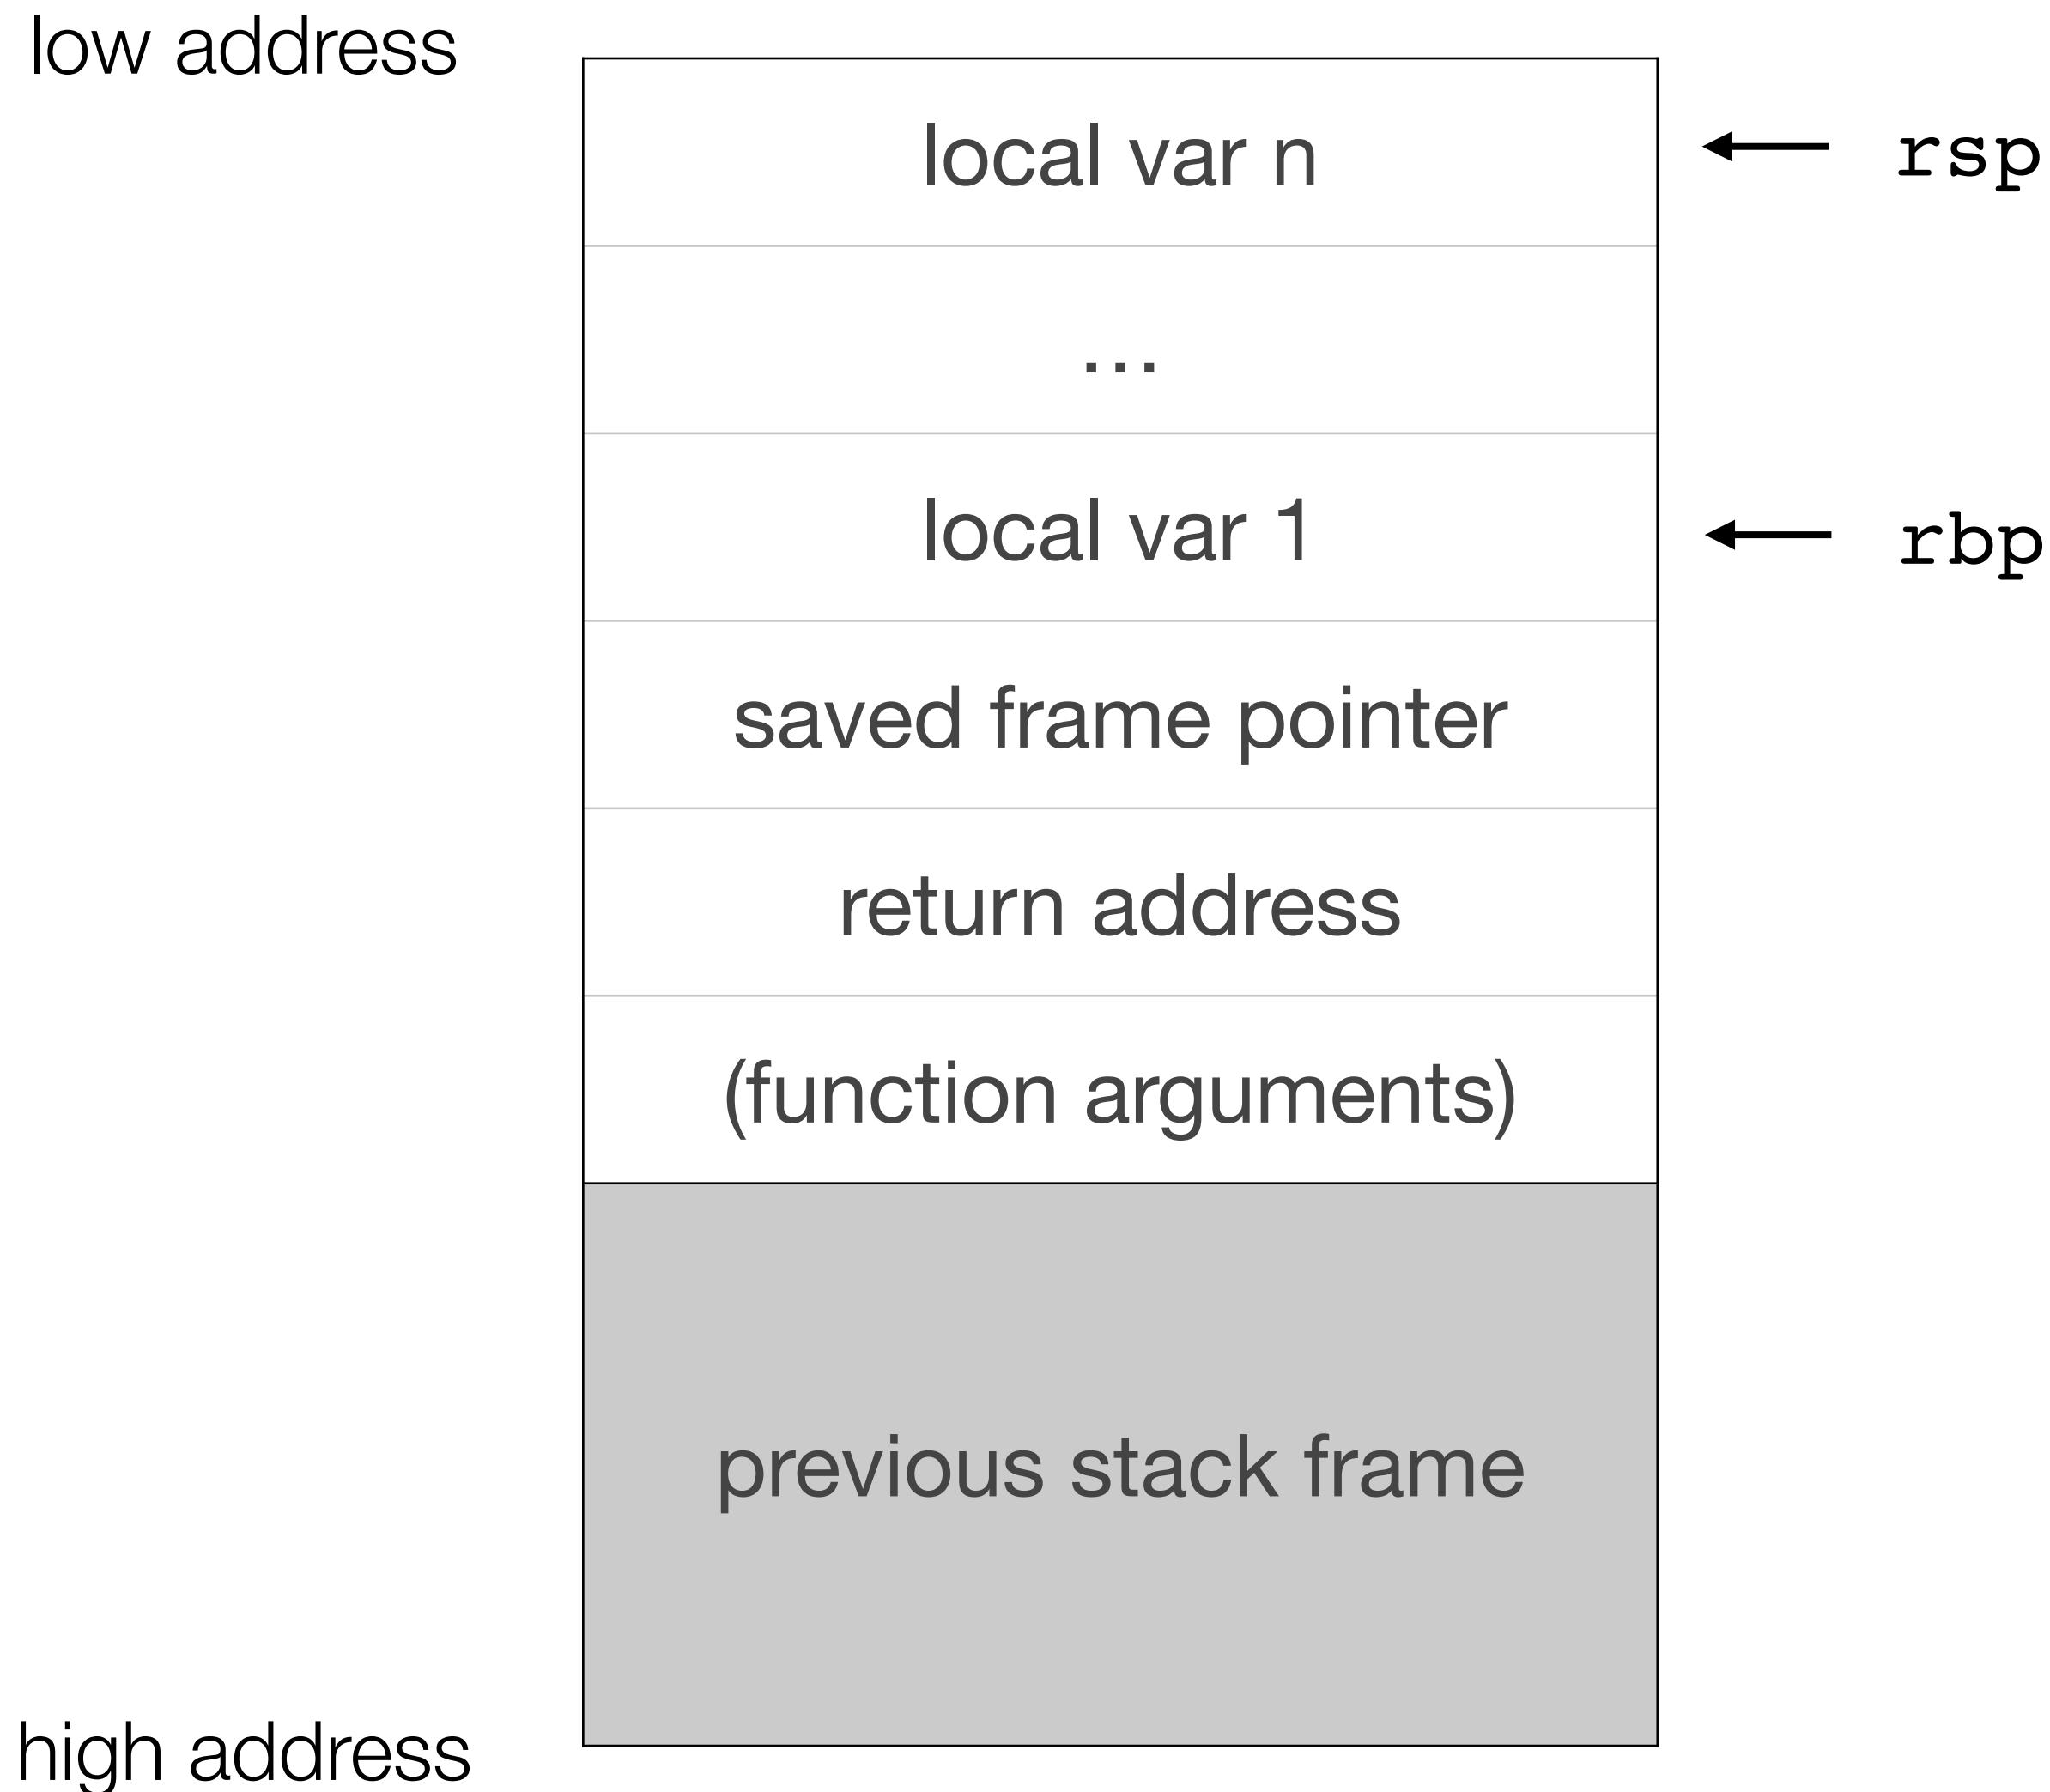
| Register | Usage | Preserved across function calls |
|-------------|---|---------------------------------|
| %rax | temporary register; with variable arguments passes information about the number of vector registers used; 1 st return register | No |
| %rbx | callee-saved register | Yes |
| %rcx | used to pass 4 th integer argument to functions | No |
| %rdx | used to pass 3 rd argument to functions; 2 nd return register | No |
| %rsp | stack pointer | Yes |
| %rbp | callee-saved register; optionally used as frame pointer | Yes |
| %rsi | used to pass 2 nd argument to functions | No |
| %rdi | used to pass 1 st argument to functions | No |
| %r8 | used to pass 5 th argument to functions | No |
| %r9 | used to pass 6 th argument to functions | No |
| %r10 | temporary register, used for passing a function's static chain pointer | No |
| %r11 | temporary register | No |
| %r12–r14 | callee-saved registers | Yes |
| %r15 | callee-saved register; optionally used as GOT base pointer | Yes |
| %xmm0–%xmm1 | used to pass and return floating point arguments | No |
| %xmm2–%xmm7 | used to pass floating point arguments | No |

More complex arguments (e.g. structs) passed through the stack

Stack frames

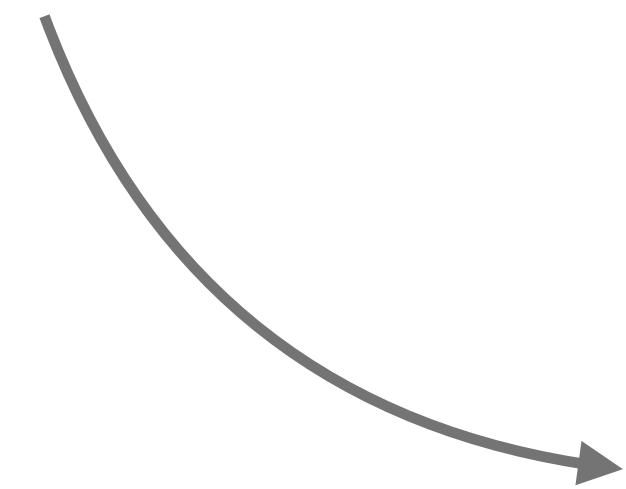
- **Local stack space during function call**
- When the function starts:
 - 1) the return address is pushed (`call` does that)
 - 2) the previous base pointer `rbp` is pushed for backup
 - 3) the base pointer is set to the current position `rsp`
 - 4) `rsp` is decreased (stack grows) the size of local variables
 - 5) space for local variables is between `rbp` and `rsp`
- This set of data is the **stack frame**
- If a function is called within the function, a new stack frame is created on top of the current one, etc...

Stack frame visualisation



Stack frame ASM implementation

```
1 void function(void)
2 {
3     int a = 1, b = 2;
4
5     printf("%d\n", a+b);
6
7     return;
8 }
```

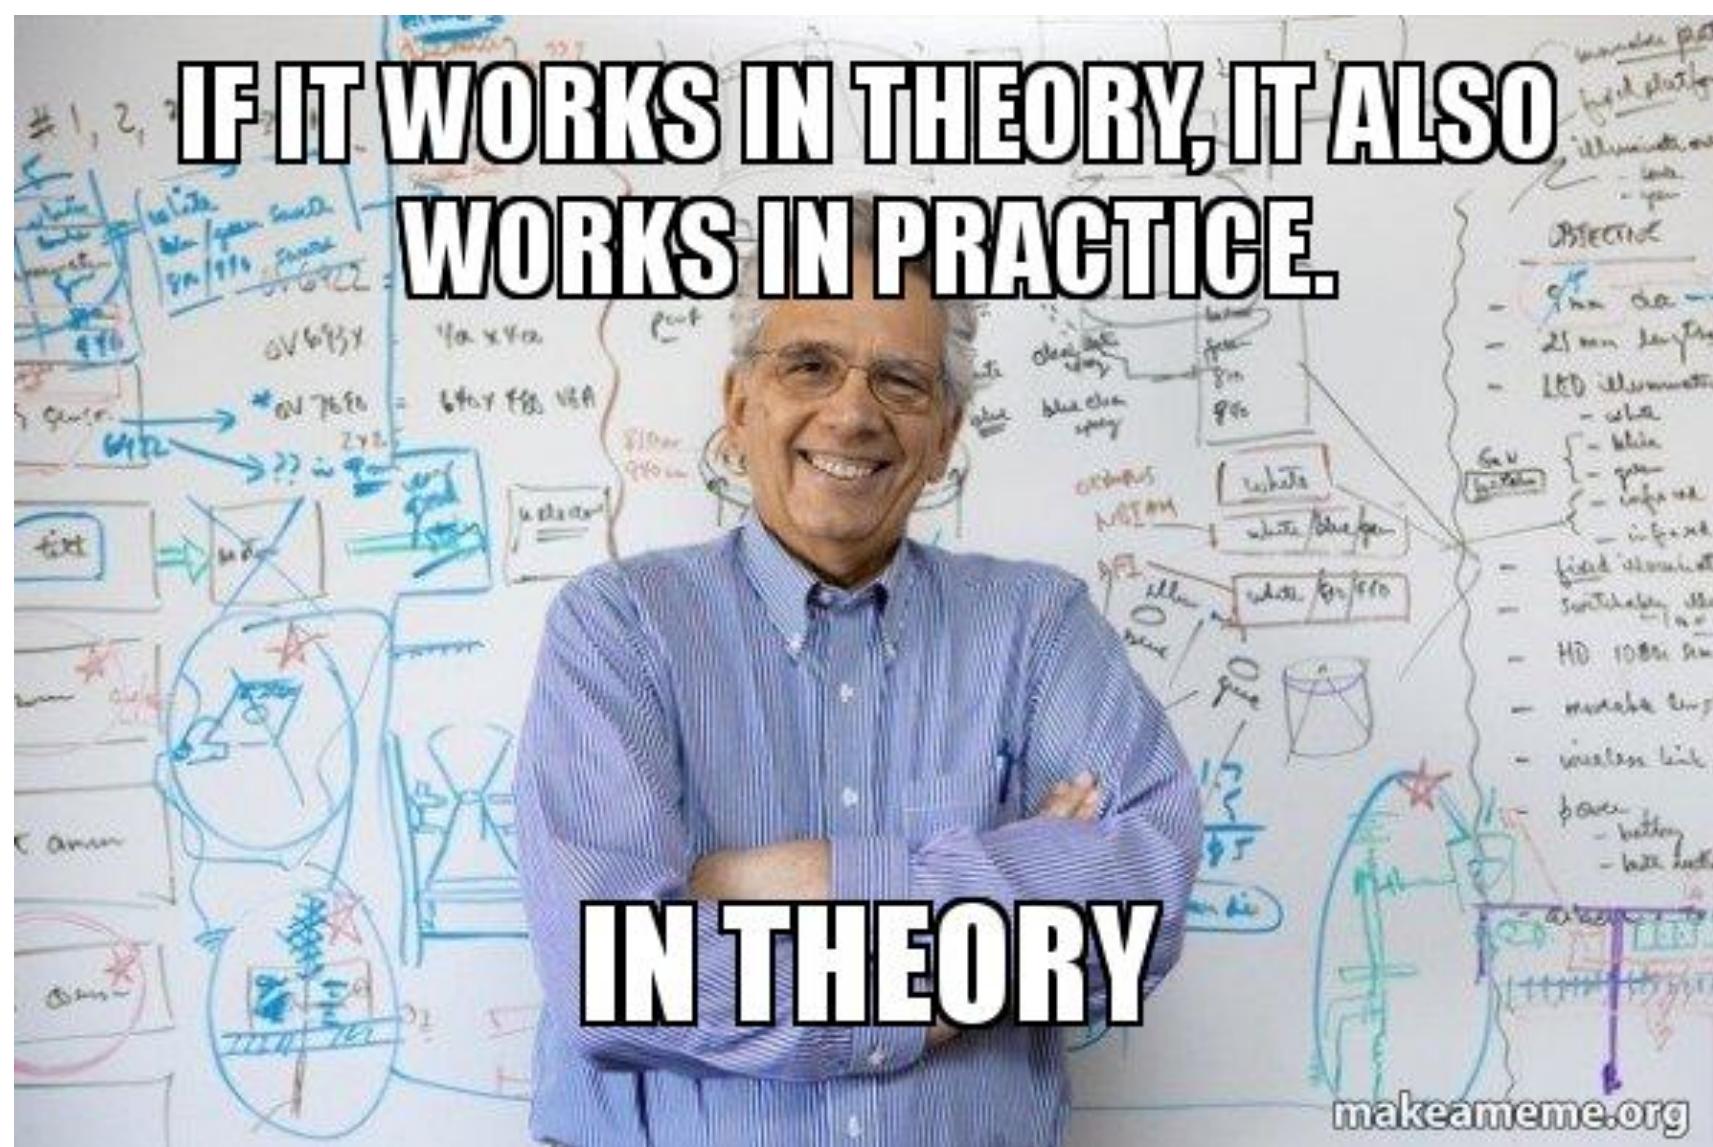


```
1 function: # @function
2 ; save previous frame pointer
3 push rbp
4 mov rbp, rsp
5 ; allocate 16B on the stack
6 sub rsp, 16
7 ; push a on the stack
8 mov dword ptr [rbp - 4], 1
9 ; push b on the stack
10 mov dword ptr [rbp - 8], 2
11
12 ; ... printf call ...
13
14 ; remove 16B from the stack
15 add rsp, 16
16 ; restore saved base pointer
17 pop rbp
18 ; return (will pop return address)
19 ret
```

<https://godbolt.org/z/nPE7sK>

Remember...

- **These are all conventions! You do what you want!**
- But this is what compilers are going to generate respecting the System V AMD64 ABI
- Nothing replaces practice for these things...



Your best friend the debugger

Generalities on debuggers

- Debuggers allow you to step line-by-line (or instruction-by-instruction) through a program execution
- They generally offer possibilities to disassemble the code and investigate memory and registers
- **They are invaluable tools to understand runtime issues when testing a program**
- Most common is probably GDB, here I will use **LLDB** from the LLVM project.

LLDB vs. GDB syntax: <https://bit.ly/3q7dme4>

Starting LLDB

- Compiling C/C++ code with the `-g` option allow debuggers to be aware of the source code

```
$ clang -g hello.c -o hello
$ ./hello
Hello world
$ lldb hello
(lldb) target create "hello"
Current executable set to '/mnt/hgfs/binary-hack/lectures/1-x86_64/hello' (x86_64).
(lldb) l
 5  {
 6      printf("Hello world\n");
 7
 8      return EXIT_SUCCESS;
 9  }
10
(lldb) r
Process 61856 launched: '/mnt/hgfs/binary-hack/lectures/1-x86_64/hello' (x86_64)
Hello world
Process 61856 exited with status = 0 (0x00000000)
```

Breakpoints

- Breakpoints are locations (code or binary) where the execution will be paused

```
(lldb) b 6
Breakpoint 1: where = hello`main + 15 at hello.c:6:3, address = 0x000000000040113f
(lldb) r
Process 61989 launched: '/mnt/hgfs/binary-hack/lectures/1-x86_64/hello' (x86_64)
Process 61989 stopped
* thread #1, name = 'hello', stop reason = breakpoint 1.1
  frame #0: 0x000000000040113f hello`main at hello.c:6:3
  3
  4     int main(void)
  5 {
→ 6     printf("Hello world\n");
  7
  8     return EXIT_SUCCESS;
  9 }
(lldb) c
Process 61989 resuming
Hello world
Process 61989 exited with status = 0 (0x00000000)
```

Breakpoints

- Breakpoints are locations (code or binary) where the execution stops

```
(lldb) d -n main
hello`main:
0x401130 <+0>: push   rbp
0x401131 <+1>: mov    rbp, rsp
0x401134 <+4>: sub    rsp, 0x10
0x401138 <+8>: mov    dword ptr [rbp - 0x4], 0x0
0x40113f <+15>: movabs rdi, 0x402004
0x401149 <+25>: mov    al, 0x0
0x40114b <+27>: call   0x401030          ; symbol stub for: printf
0x401150 <+32>: xor    ecx, ecx
0x401152 <+34>: mov    dword ptr [rbp - 0x8], eax
0x401155 <+37>: mov    eax, ecx
0x401157 <+39>: add    rsp, 0x10
0x40115b <+43>: pop    rbp
0x40115c <+44>: ret

(lldb) b 0x40114b
Breakpoint 2: where = hello`main + 27 at hello.c:6:3, address = 0x000000000040114b
(lldb) r
Process 62132 launched: '/mnt/hgfs/binary-hack/lectures/1-x86_64/hello' (x86_64)
Process 62132 stopped
* thread #1, name = 'hello', stop reason = breakpoint 1.1
  frame #0: 0x000000000040113f hello`main at hello.c:6:3
(lldb) c
4      int main(void)
5  {
Hello world → 6      printf("Hello world\n");
Process 62132 stopped
7
8      return EXIT_SUCCESS;
9 }
```

Backtraces

- LLDB can provide the stack of calls that led to the current code position

```
$ lldb functions
(lldb) target create "functions"
Current executable set to '/mnt/hgfs/binary-hack/lectures/2-elf/functions' (x86_64).
(lldb) l
20  {
21      printf("%f\n", function2(2.));
22
23      return EXIT_SUCCESS;
24  }
(lldb) b function1
Breakpoint 1: where = functions`function1 + 10 at functions.c:7:11, address = 0x000000000040113a
(lldb) r
Process 107292 launched: '/mnt/hgfs/binary-hack/lectures/2-elf/functions' (x86_64)
Process 107292 stopped
* thread #1, name = 'functions', stop reason = breakpoint 1.1
    frame #0: 0x000000000040113a functions`function1(a=2, b=2) at functions.c:7:11
4
5      int function1(int a, int b)
6  {
→ 7      int c = a + b;
8
9      return c;
10 }
(lldb) bt
* thread #1, name = 'functions', stop reason = breakpoint 1.1
* frame #0: 0x000000000040113a functions`function1(a=2, b=2) at functions.c:7:11
frame #1: 0x000000000040116c functions`function2(a=2) at functions.c:14:19
frame #2: 0x00000000004011ac functions`main at functions.c:21:18
frame #3: 0x00007ffff7de60b3 libc.so.6`__libc_start_main + 243
frame #4: 0x000000000040106e functions`_start + 46
```

Investigating registers and memory

- Registers and memory can be investigated with a large choice of formats

```
(lldb) re r
General Purpose Registers:
    rax = 0x0000000000401190  functions`main at functions.c:20
    rbx = 0x00000000004011d0  functions`__libc_csu_init
    rcx = 0x00000000004011d0  functions`__libc_csu_init
    rdx = 0x0007fffffffde8
    rdi = 0x0000000000000002
    rsi = 0x0000000000000002
    rbp = 0x0007fffffffdea0
    rsp = 0x0007fffffffdea0
    r8 = 0x0000000000000000
    r9 = 0x0007ffff7fe0d50  ld-2.31.so`__lldb_unnamed_symbol59$$ld-2.31.so
    r10 = 0x0000000000000002
    r11 = 0x0000000000000000
    r12 = 0x0000000000401040  functions`_start
    r13 = 0x0007fffffffdf0
    r14 = 0x0000000000000000
    r15 = 0x0000000000000000
    rip = 0x000000000040113a  functions`function1 + 10 at functions.c:7:11
rflags = 0x000000000000202
    cs = 0x000000000000033
    fs = 0x0000000000000000
    gs = 0x0000000000000000
    ss = 0x000000000000002b
    ds = 0x0000000000000000
    es = 0x0000000000000000
(lldb) x/x 0x0007fffffffde9c
0x7fffffffde9c: 0x00000002
(lldb) x/d 0x0007fffffffde9c
0x7fffffffde9c: 2
```

- LLDB is actually a great calculator for programming!

LLDB demonstrations

- Complete commented logs

<https://bit.ly/2NFGEmS>

<https://bit.ly/3kudPFT>

Let's play with stack frames

Test program with nested calls

```
1 #include <stdio.h>
2 #include <stdint.h>
3 #include <stdlib.h>
4
5 int function1(int a, int b)
6 {
7     int c = a + b;
8
9     return c;
10}
11
12 double function2(double a)
13 {
14     double b = 3.14*function1(a, a);
15
16     return b;
17}
18
19 int main(void)
20 {
21     printf("%f\n", function2(2.));
22
23     return EXIT_SUCCESS;
24}
```

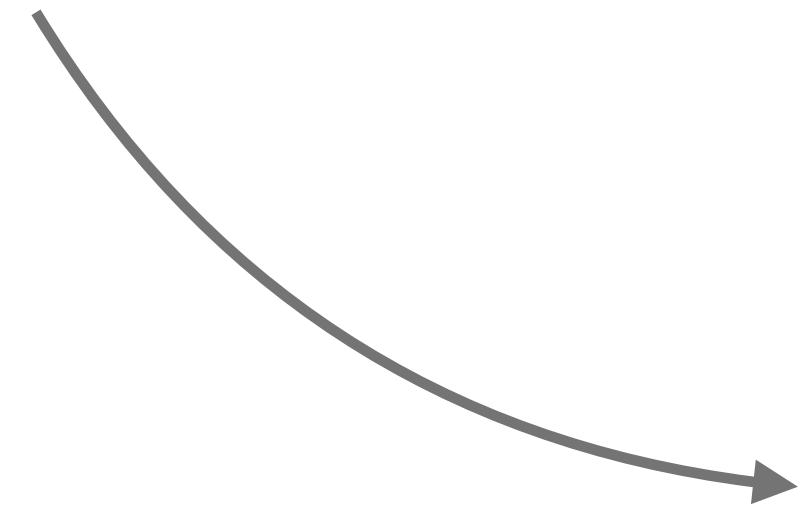
<https://godbolt.org/z/T17j3e>

Some code analysis first

```
1 int function1(int a, int b)
2 {
3     int c = a + b;
4
5     return c;
6 }
```

Should be quite clear at this point

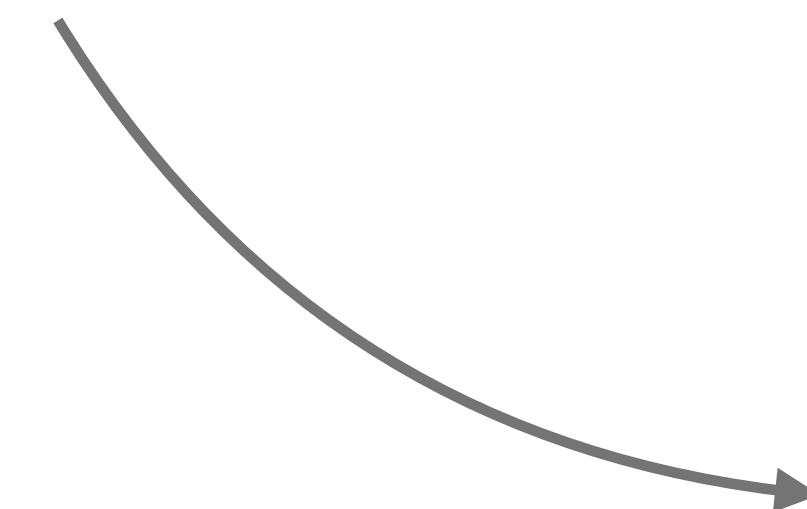
Do you see how sub-optimal this is?
(no compiler optimisations)



```
1 function1:
2     push rbp
3     mov rbp, rsp
4     mov dword ptr [rbp - 4], edi
5     mov dword ptr [rbp - 8], esi
6     mov eax, dword ptr [rbp - 4]
7     add eax, dword ptr [rbp - 8]
8     mov dword ptr [rbp - 12], eax
9     mov eax, dword ptr [rbp - 12]
10    pop rbp
11    ret
```

Some code analysis first

```
1 double function2(double a)
2 {
3     double b = 3.14*function1(a, a);
4
5     return b;
6 }
```



Notes:

- `cvt*` convert `double <>> int`
- `rip + .LCPI_0` will be the content at `.LCPI_0` (cf. RIP-relative addressing)
- 4614.... is 3.14 in binary/decimal!
- you can spot some useless statements

```
1 .LCPI1_0:
2     .quad 4614253070214989087
3 function2:
4     push rbp
5     mov rbp, rsp
6     sub rsp, 16
7     movsd qword ptr [rbp - 8], xmm0
8     cvttsd2si edi, qword ptr [rbp - 8]
9     cvttsd2si esi, qword ptr [rbp - 8]
10    call function1
11    movsd xmm0, qword ptr [rip + .LCPI1_0]
12    cvtsi2sd xmm1, eax
13    mulsd xmm0, xmm1
14    movsd qword ptr [rbp - 16], xmm0
15    movsd xmm0, qword ptr [rbp - 16]
16    add rsp, 16
17    pop rbp
18    ret
```

Stack frame from debugger

- Complete commented log

<https://bit.ly/3aYM318>

Make your own: extracting stake frame info

- Stack frame info can be directly extracted in C

```
1 void stackFrameInfo(void)
2 {
3     uint64_t **rbp; // it does not have to be a double pointer
4                     // it is just convenient for what follows
5
6     __asm__ volatile
7     (
8         "movq %%rbp, %0;"
9         : "=r"(rbp)
10    );
11
12    // here rbp is the base pointer of stackFrameInfo
13    printf("==== STACK FRAME INFO ====\n");
14    // 8B after is the return address of stackFrameInfo
15    // i.e. the caller of stackFrameInfo
16    printf("Called from %p\n", *(rbp + 1));
17    // we actually want the base pointer of the caller...
18    // remember, this is the first stack element (at the bottom)
19    rbp = (uint64_t **)(*rbp);
20    // again, return address is 8B after
21    printf("Stack base %p -- return address %p\n", rbp, *(rbp + 1));
22    printf("======\n");
23 }
```

Make your own: extracting stake frame info

- Let's insert it in all functions of the previous example

```
$ ./stackframes
==== STACK FRAME INFO ====
Called from 0x401254
Stack base 0x7ffc25cac000 -- return address 0x7fb9062a70b3
=====
==== STACK FRAME INFO ====
Called from 0x4011dc
Stack base 0x7ffc25cabfc0 -- return address 0x40120c
=====
==== STACK FRAME INFO ====
Called from 0x401226
Stack base 0x7ffc25cabfe0 -- return address 0x401261
=====
12.560000
```

- addr2line can convert addresses into code positions

```
$ addr2line -fp -e ./stackframes -a 0x401254 0x4011dc 0x401226
0x0000000000401254: main at /mnt/hgfs/binary-hack/lectures/2-elf/stackframes.c:49
0x00000000004011dc: function1 at /mnt/hgfs/binary-hack/lectures/2-elf/stackframes.c:35
0x0000000000401226: function2 at /mnt/hgfs/binary-hack/lectures/2-elf/stackframes.c:44
```

Make your own: backtrace

- You can walk the stack frame and get a backtrace

```
1 void backtrace(void)
2 {
3     uint64_t      **rbp;
4     unsigned int  level = 0;
5
6     __asm__ volatile
7     (
8         "movq %%rbp, %0;"
9         : "=r"(rbp)
10    );
11
12    printf("==== BACKTRACE =====\n");
13    // keep walking until rbp is NULL
14    // no guarantee this is a safe way to do thing
15    while (rbp != NULL)
16    {
17        // get the return address 8B after base pointer
18        printf("frame -: %p\n", level, *(rbp + 1));
19        // set rbp to be next frame pointer
20        rbp = (uint64_t **)(*rbp);
21        level++;
22    }
23    printf("=====\\n");
24 }
```

Make your own: backtrace

- Call backtrace from function1

```
$ ./backtrace
==== BACKTRACE ====
frame 0: 0x4011cc
frame 1: 0x4011fc
frame 2: 0x40123c
frame 3: 0x7f4e7d29c0b3
=====
12.560000
$ addr2line -fp -e ./backtrace -a 0x4011cc 0x4011fc 0x40123c 0x7f4e7d29c0b3
0x000000004011cc: function1 at /mnt/hgfs/binary-hack/lectures/2-elf/backtrace.c:36
0x000000004011fc: function2 at /mnt/hgfs/binary-hack/lectures/2-elf/backtrace.c:41
0x0000000040123c: main at /mnt/hgfs/binary-hack/lectures/2-elf/backtrace.c:48
0x00007f4e7d29c0b3: ?? ???:0
```

- Why the last address is so different?
What are the ?? ?? ?
Why is it random between executions?
I'll explain later

Frame pointer optimisation

- Actually **you don't need to save rbp...**
- You can keep **everything relative to rsp** and reduce **rsp** when needed to make the stack bigger
- This allow one more **free register** and **reduce code size**
- GCC/clang do that with **-fomit-frame-pointer**, it is included in **-O1** and higher
- Of course this is going to break the code before!

Frame pointer optimisation

```
$ clang -g -fomit-frame-pointer -o stackframes ./stackframes.c
antonin@ubuntu /mnt/hgfs/binary-hack/lectures/2-elf $ ./stackframes
===== STACK FRAME INFO =====
[1] 89547 segmentation fault (core dumped) ./stackframes
$ clang -g -fomit-frame-pointer -o backtrace backtrace.c
antonin@ubuntu /mnt/hgfs/binary-hack/lectures/2-elf $ ./backtrace
===== BACKTRACE =====
=====
12.560000
```

- There is a POSIX API for backtraces...
<https://bit.ly/3r3NEYZ> but I hope you learned something!
- And... (from the man page above)
 - * Omission of the frame pointers (as implied by any of `gcc(1)`'s nonzero optimization levels) may cause these assumptions to be violated.
- So in general:
Mind your optimisation level while debugging

The Executable and Linkable Format (ELF)

Remaining open questions

- How does the OS know how to start a program?
- How are jumps and calls handled if the program is divided across several executable binaries (i.e. libraries)?

Executable binary format

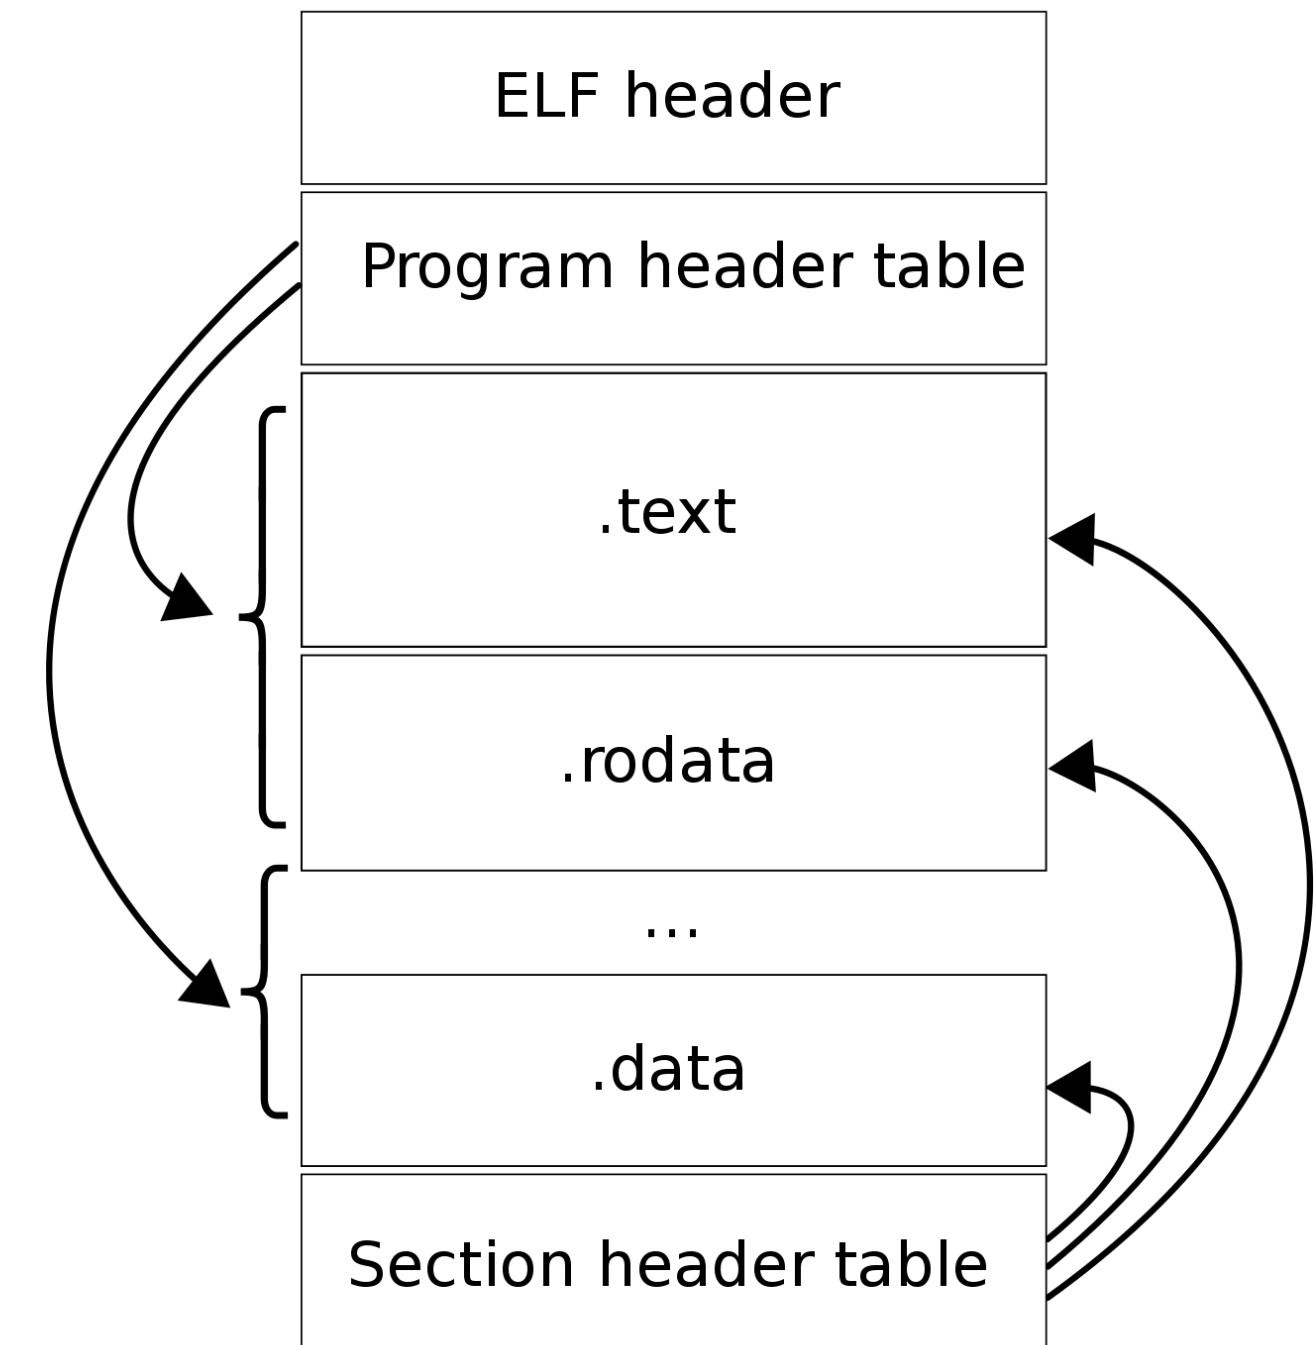
- When an executable is generated, its compiled code is wrapped up with all sorts of data explain the kernel how to run the program
- This data is generated by the **linker**
- On Linux, the adopted format is the **Executable and Linkable Format (ELF)**

```
$ file ./hello
./hello: ELF 64-bit LSB executable, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/ld-linux-x86-64.so.2,
BuildID[sha1]=113ce67758fa247a837becd3bac8a314e8c6bcf5, for GNU/Linux 3.2.0, with debug_info, not stripped
```

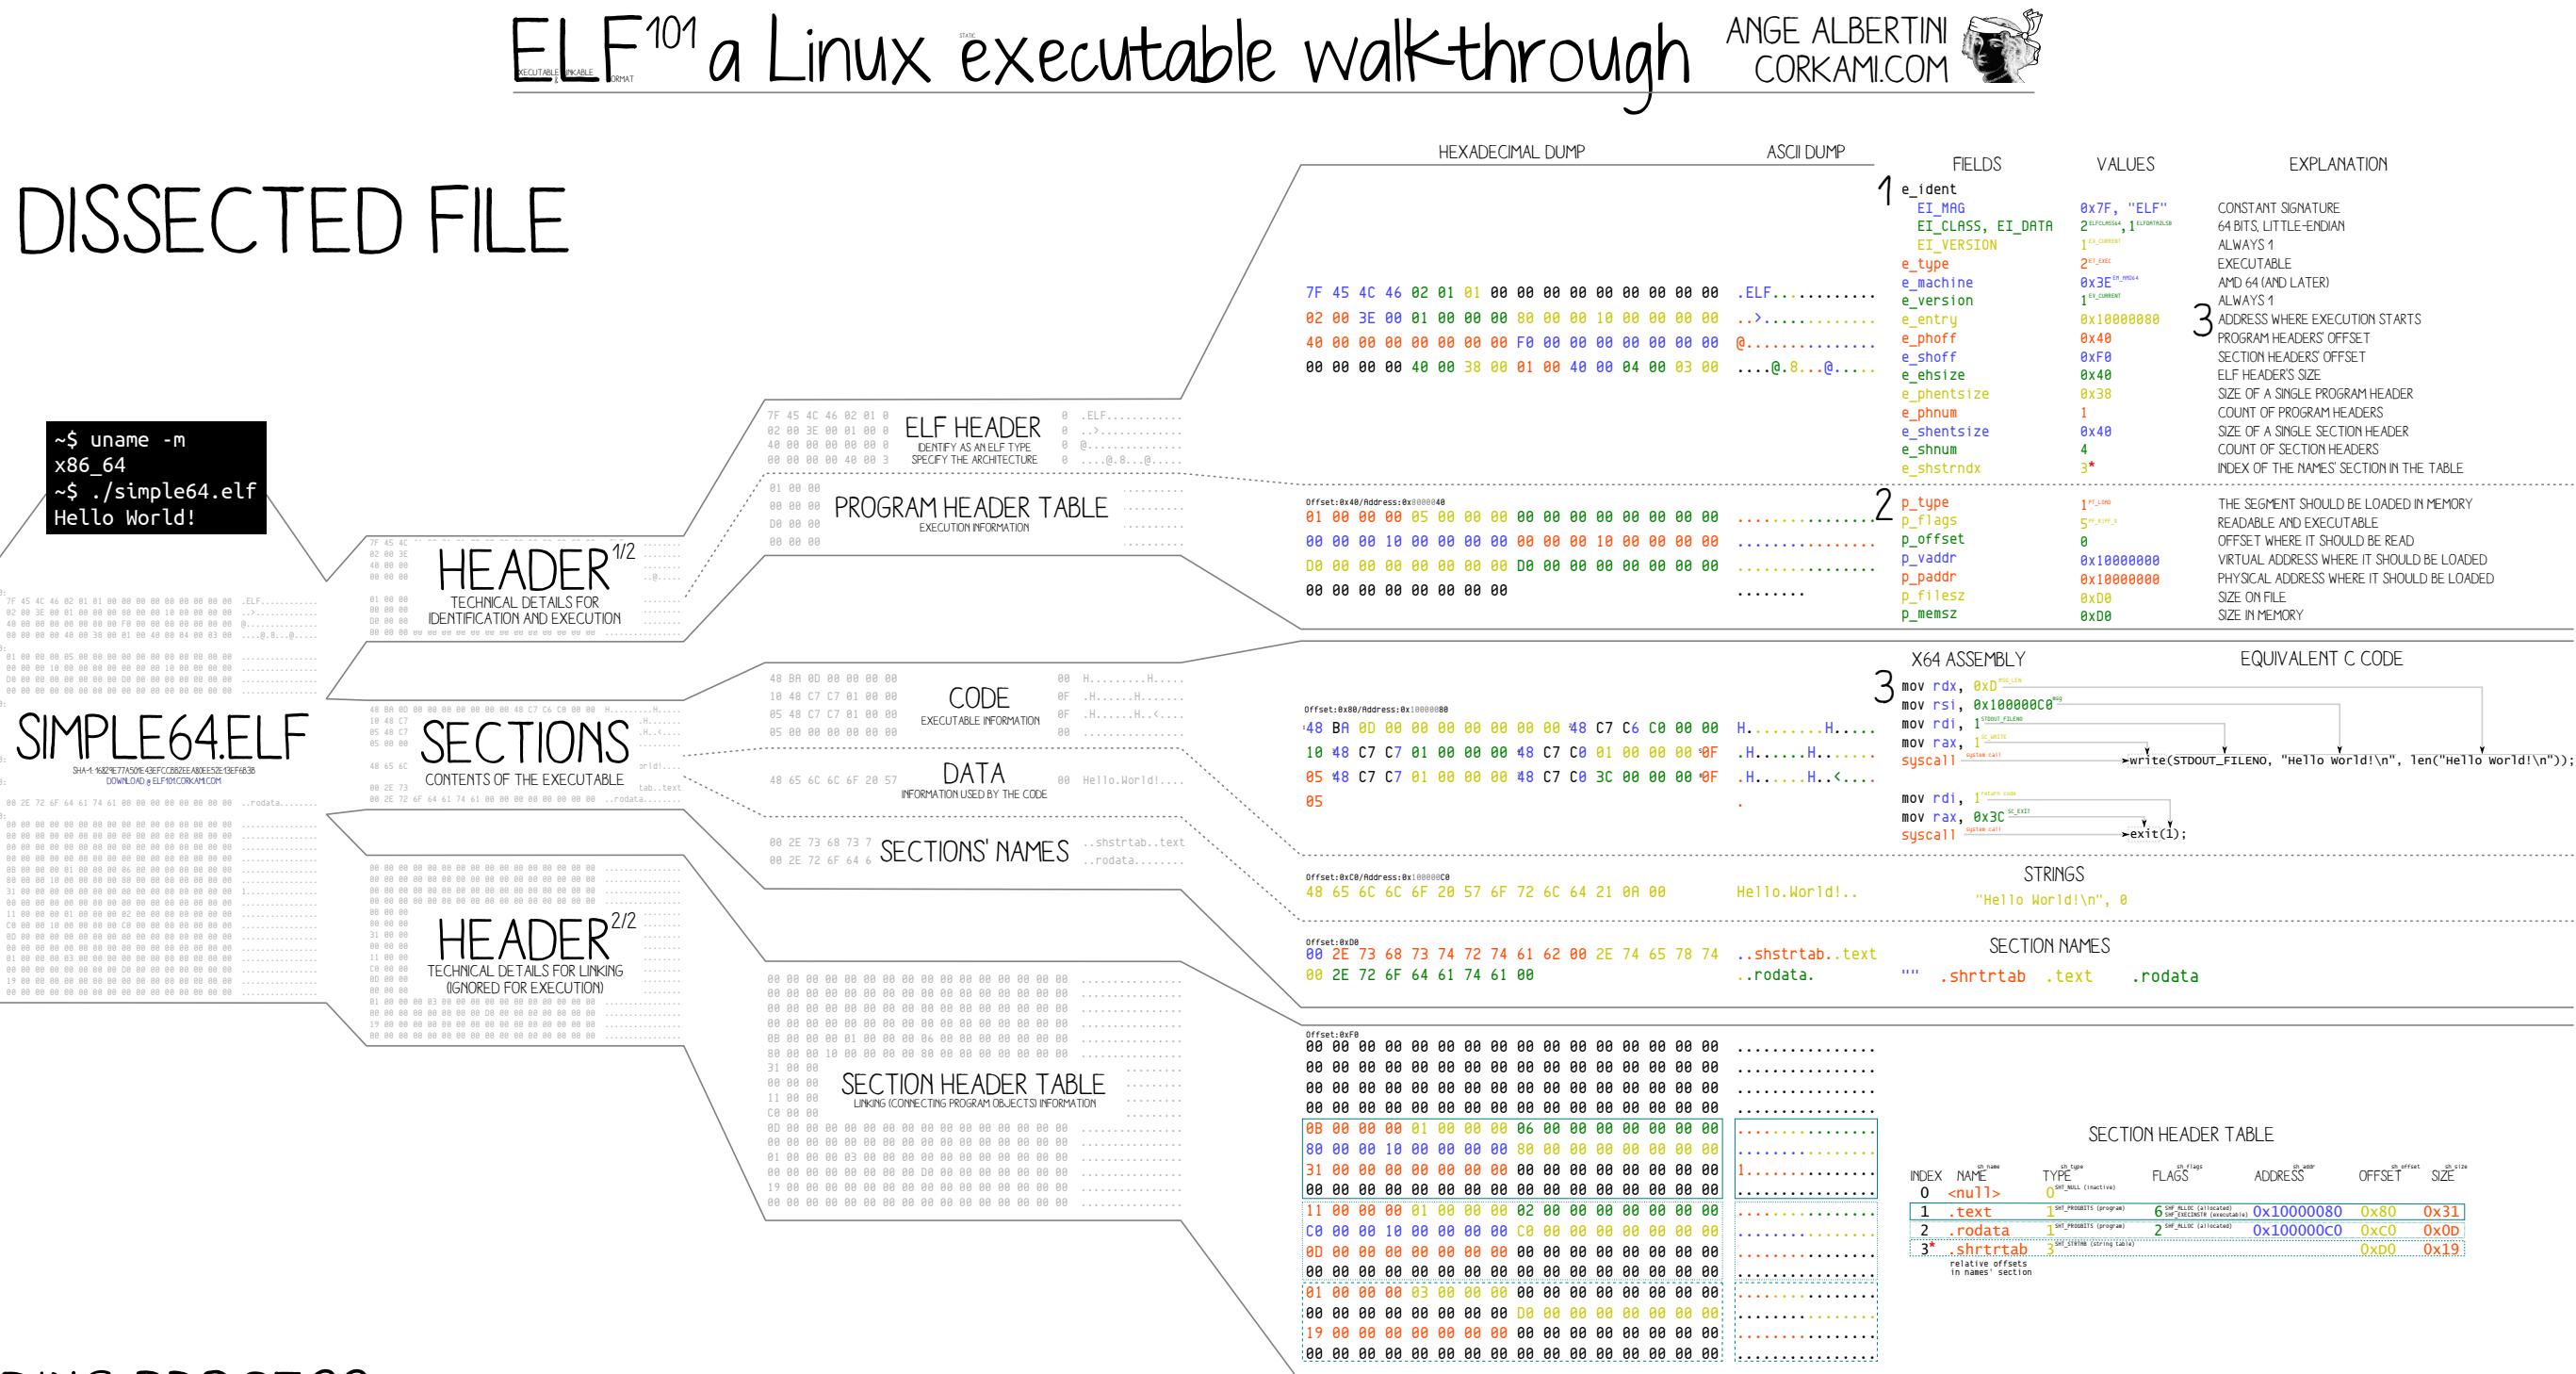
- This is **completely OS-dependent**, although ELF is common across a number of platforms (macOS/iOS uses Mach-O, Windows uses PE)

Quick description of ELF

1. A **ELF header** containing global properties of the binary
2. A **program table** describing the memory segments to create
3. **Sections** containing the content of the binary (code, data, ...)
4. A **section header table** describing how to link against the binary



ELF in one chart



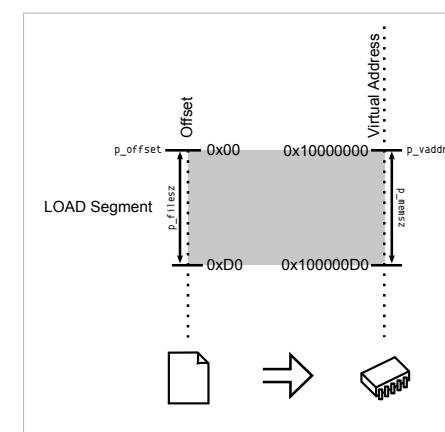
LOADING PROCESS

1 HEADER

THE ELF HEADER IS PARSED
THE PROGRAM HEADER IS PARSED
(SECTIONS ARE NOT USED)

2 MAPPING

THE FILE IS MAPPED IN MEMORY
ACCORDING TO ITS SEGMENT(S)



3 EXECUTION

ENTRY IS CALLED
SYSCALLS ARE ACCESSED VIA:
- SYSCALL NUMBER IN THE RAX REGISTER
- CALLING INSTRUCTION SYSCALL

TRIVIA

THE ELF WAS FIRST SPECIFIED BY U.S. L. AND U.I.
FOR UNIX SYSTEM V, IN 1989

THE ELF IS USED, AMONG OTHERS, IN:

- LINUX, ANDROID, *BSD, SOLARIS, BEOS
- PSP, PLAYSTATION 2-4, DREAMCAST, GAMECUBE, WII
- VARIOUS OSES MADE BY SAMSUNG, ERICSSON, NOKIA,
- MICROCONTROLLERS FROM ATMEL, TEXAS INSTRUMENTS

Binary tools

- The `readelf` command reads the ELF information
- The `objdump` command can read the content of the binary and disassemble code
- The `nm` command gives the symbols table (functions, constants, we'll see that in a min)
- The `strings` command gives the constant strings

Linking

- Remember, **function calls are just jumps**
- The linker implements in the binary all the necessary tools to find the addresses of functions when called
- ```
$ clang ./hello.c -o ./hello
```

actually does 2 things: compiling & linking

- Use **-c** to just compile, the result is an **object file**

```
$ clang -c ./hello.c -o ./hello.o
$ file ./hello.o
./hello.o: ELF 64-bit LSB relocatable, x86-64, version 1 (SYSV), not stripped
```

the linker assemble all object files into a properly formed ELF executable binary

# Linking

- The default linker on Linux is GNU LD
- **Calling the linker by hand is tricky...** there is a lot of information to provide on how to link to system libraries
- Using the compiler as a linker automatises that part

```
$ clang hello.o -o hello
$./hello
Hello world
```

- Use **-v** to know what the linking command was, here it is

```
/usr/bin/ld -z relro --hash-style=gnu --build-id --eh-frame-hdr -m elf_x86_64 -dynamic-linker /lib64/ld-linux-x86-64.so.2 -o ./hello /usr/bin/../lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu/crt1.o /usr/bin/../lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu/crti.o /usr/bin/../lib/gcc/x86_64-linux-gnu/9/crtbegin.o -L/usr/bin/../lib/gcc/x86_64-linux-gnu/9 -L/usr/bin/../lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu -L/usr/bin/../lib/gcc/x86_64-linux-gnu/9/../../../../lib64 -L/lib/x86_64-linux-gnu -L/lib/./lib64 -L/usr/lib/x86_64-linux-gnu -L/usr/lib/llvm-10/bin/./lib -L/lib -L/usr/lib /tmp/hello-a175e2.o -lgcc --as-needed -lgcc_s --no-as-needed -lc -lgcc --as-needed -lgcc_s --no-as-needed /usr/bin/../lib/gcc/x86_64-linux-gnu/9/crtend.o /usr/bin/../lib/gcc/x86_64-linux-gnu/9/../../../../x86_64-linux-gnu/crtn.o
```

!?

# Static linking

---

- During **static linking**, the binary code of the external function is **relocated and copied** inside the final binary
- Pro: the executable binary is **independent** from the library and therefore **more portable**
- Con: if the library is updated, all statically linked executables **need to be relinked**
- Con: statically linked executables are **larger**, and **redundant** information is stored
- Linux static libraries have extension `.a` and can be assembled from objects using `ar`

# Dynamic linking

---

- Linking is done at runtime: when the program starts, the linked code is loaded in the program memory space
- The system and the executable contain informations on how to find the dynamically linked libraries at runtime (**dynamic links**)
- Pro: **smaller** executables
- Pro: if the library is updated, nothing needs to be done to dynamically linked programs (assuming same interface)
- Con: executables are **less portable** as they will always need the external libraries when loaded

# Dynamic linking environment

- When dynamic linking, only references to the libraries files are generated in the binary, code is only linked at runtime
- You can see the dynamic link resolution using `ldd`

```
$ ldd latan-sample-read
 linux-vdso.so.1 (0x00007ffe04767000)
 libLatAnalyze.so.0 => /home/antonin/Desktop/Plateau/dependencies/prefix/lib/libLatAnalyze.so.0 (0x00007fe1dbb70000)
 libstdc++.so.6 => /lib/x86_64-linux-gnu/libstdc++.so.6 (0x00007fe1db97a000)
 libm.so.6 => /lib/x86_64-linux-gnu/libm.so.6 (0x00007fe1db82b000)
 libgcc_s.so.1 => /lib/x86_64-linux-gnu/libgcc_s.so.1 (0x00007fe1db810000)
 libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007fe1db61e000)
 libgsl.so.25 => /home/antonin/Desktop/Plateau/dependencies/prefix/lib/libgsl.so.25 (0x00007fe1db33f000)
 libgslcblas.so.0 => /home/antonin/Desktop/Plateau/dependencies/prefix/lib/libgslcblas.so.0 (0x00007fe1db2fb000)
 libnlopt.so.0 => /home/antonin/Desktop/Plateau/dependencies/prefix/lib/libnlopt.so.0 (0x00007fe1db258000)
 libhdf5_cpp.so.103 => /home/antonin/Desktop/Plateau/dependencies/prefix/lib/libhdf5_cpp.so.103 (0x00007fe1db1e5000)
 libhdf5.so.103 => /home/antonin/Desktop/Plateau/dependencies/prefix/lib/libhdf5.so.103 (0x00007fe1dae01000)
 libMinuit2.so.0 => /home/antonin/Desktop/Plateau/dependencies/prefix/lib/libMinuit2.so.0 (0x00007fe1dad59000)
 /lib64/ld-linux-x86-64.so.2 (0x00007fe1dbd39000)
 libz.so.1 => /lib/x86_64-linux-gnu/libz.so.1 (0x00007fe1dad3b000)
```

- Path resolution sequence depends on the environment and information in the binary <https://bit.ly/3bojyZJ>

# Linking troubleshooting checklist

```
/usr/bin/ld: undefined reference to `H5T_NATIVE_INT32_g'
/usr/bin/ld: undefined reference to `H5Pcreate'
/usr/bin/ld: undefined reference to `H5Tcopy'
/usr/bin/ld: undefined reference to `H5T_NATIVE_UINT32_g'
/usr/bin/ld: undefined reference to `H5Sget_simple_extent_dims'
/usr/bin/ld: undefined reference to `H5Eset_auto2'
/usr/bin/ld: undefined reference to `H5Sget_simple_extent_ndims'
/usr/bin/ld: undefined reference to `H5Aget_type'
/usr/bin/ld: undefined reference to `H5Fcreate'
/usr/bin/ld: undefined reference to `H5Sclose'
/usr/bin/ld: undefined reference to `H5Aopen_name'
/usr/bin/ld: undefined reference to `H5Gunlink'
/usr/bin/ld: undefined reference to `H5Tget_native_type'
/usr/bin/ld: undefined reference to `H5Acreate2'
/usr/bin/ld: undefined reference to `H5Aclose'
/usr/bin/ld: undefined reference to `H5Pset_deflate'
/usr/bin/ld: undefined reference to `H5T_C_S1_g'
/usr/bin/ld: undefined reference to `H5Tset_precision'
/usr/bin/ld: undefined reference to `H5Dopen2'
/usr/bin/ld: undefined reference to `H5Dget_type'
/usr/bin/ld: undefined reference to `H5P_CLS_LINK_ACCESS_ID_g'
/usr/bin/ld: undefined reference to `H5T_NATIVE_FLOAT_g'
/usr/bin/ld: undefined reference to `H5P_CLS_DATASET_CREATE_ID_g'
/usr/bin/ld: undefined reference to `H5Lunpack_elink_val'
/usr/bin/ld: undefined reference to `H5Dwrite'
/usr/bin/ld: undefined reference to `H5Lcreate_external'
/usr/bin/ld: undefined reference to `H5Screate_simple'
/usr/bin/ld: undefined reference to `H5Aread'
/usr/bin/ld: undefined reference to `H5check_version'
/usr/bin/ld: undefined reference to `H5Iget_name'
/usr/bin/ld: undefined reference to `H5Lis_registered'
/usr/bin/ld: undefined reference to `H5Fget_obj_ids'
/usr/bin/ld: undefined reference to `H5Gcreate2'
/usr/bin/ld: undefined reference to `H5Giterate'
/usr/bin/ld: undefined reference to `H5Tget_size'
/usr/bin/ld: undefined reference to `H5Lexists'
```



*error while loading shared libraries: libhdf5.so.10: cannot open shared object file: No such file or directory*

# More about the linking process

---

- If a function is not implemented in the code (e.g. just declared in a header), it is compiled as an **undefined symbol**

```
$ nm ./hello.o
0000000000000000 T main
U printf
```

- When an executable is linked **all undefined symbols must be resolved** user linker options (e.g. `-l` & `-L` flags)
- This is a **recursive process**, libraries themselves might have undefined symbols
- This is true for both static & dynamic linking

# C++ name mangling

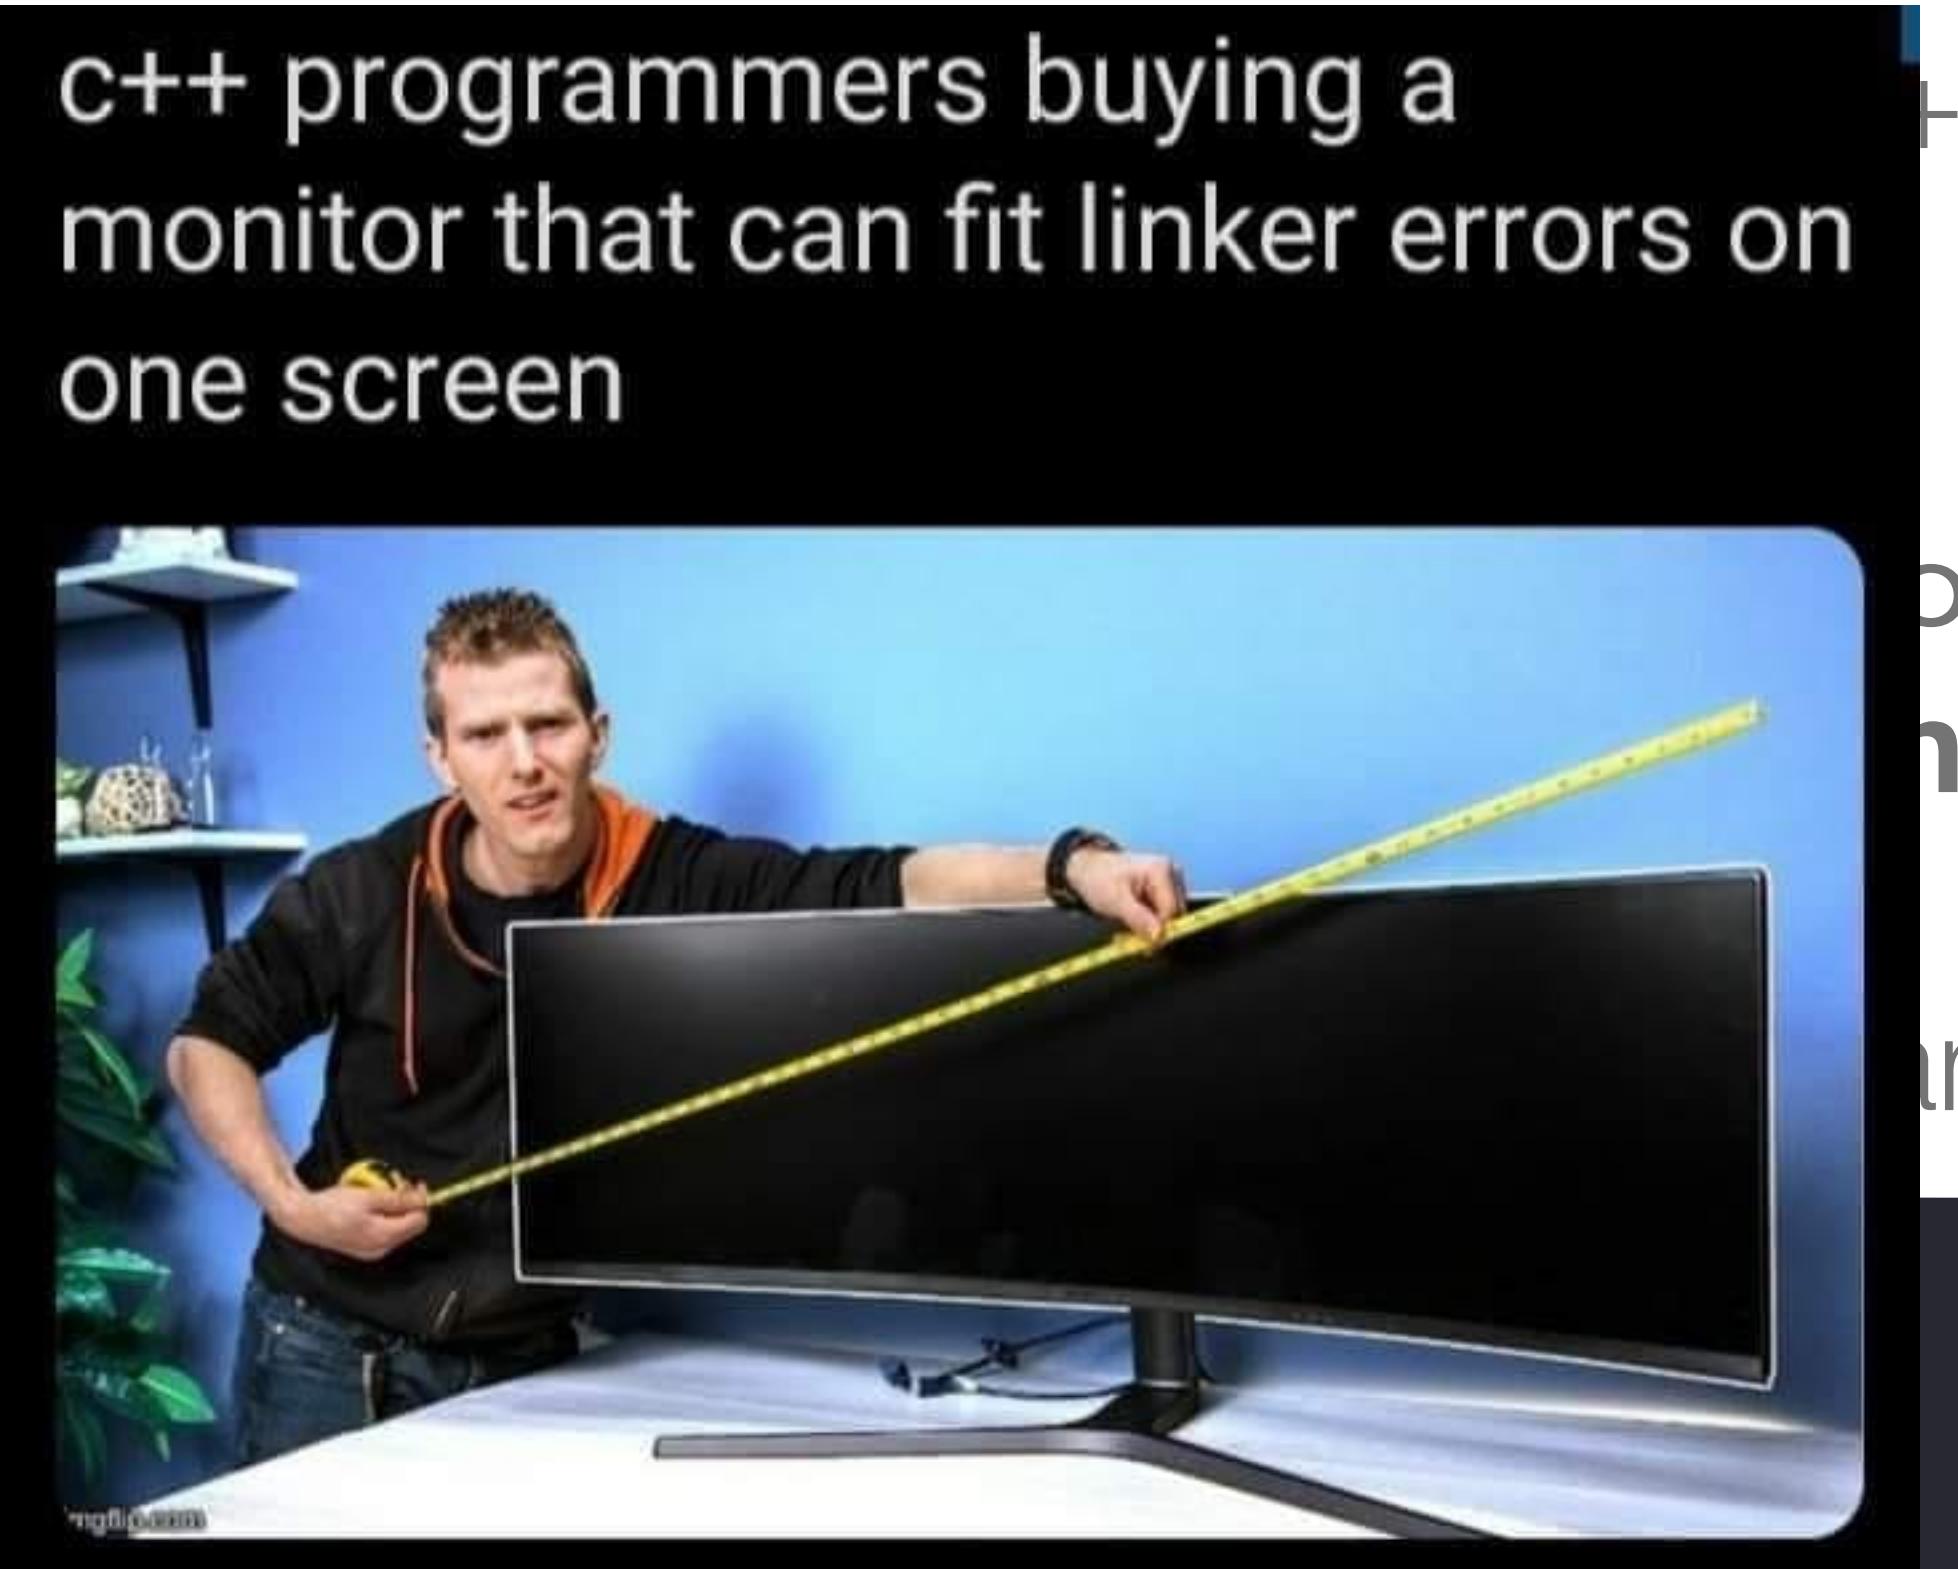
---

- Symbol names are compiler-generated, C/C++ standards do not impose any convention
- However, C symbols are generally function and C++ symbols are **mangled** according to **ABI standards**
- **c++filt** can be used to de-mangle C++ names

```
$ cat mangling.cpp
#include
std::string function(std::string s) {return "hello " + s;}
$ clang++ -c mangling.cpp -o mangling.o
$ nm mangling.o
0000000000000000 r GCC_except_table1
 U __gxx_personality_v0
 U strlen
 U _Unwind_Resume
0000000000000000 T _Z8functionNSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEEE
 U _ZNKSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEE4sizeEv
0000000000000000 W _ZNSt11char_traitsIcE6lengthEPKc
 U _ZNSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEE6appendEPKcm
 U _ZNSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEE6appendERKS4_
 U _ZNSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEE7reserveEm
 U _ZNSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEEC1Ev
 U _ZNSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEED1Ev
0000000000000000 W _ZStplIcSt11char_traitsIcESaIcEENSt7__cxx1112basic_stringIT_T0_T1_EEPKS5_RKS8_
$ echo _Z8functionNSt7__cxx1112basic_stringIcSt11char_traitsIcESaIcEEE | c++filt
function(std::__cxx11::basic_string, std::allocator >)
```

# C++ name mangling

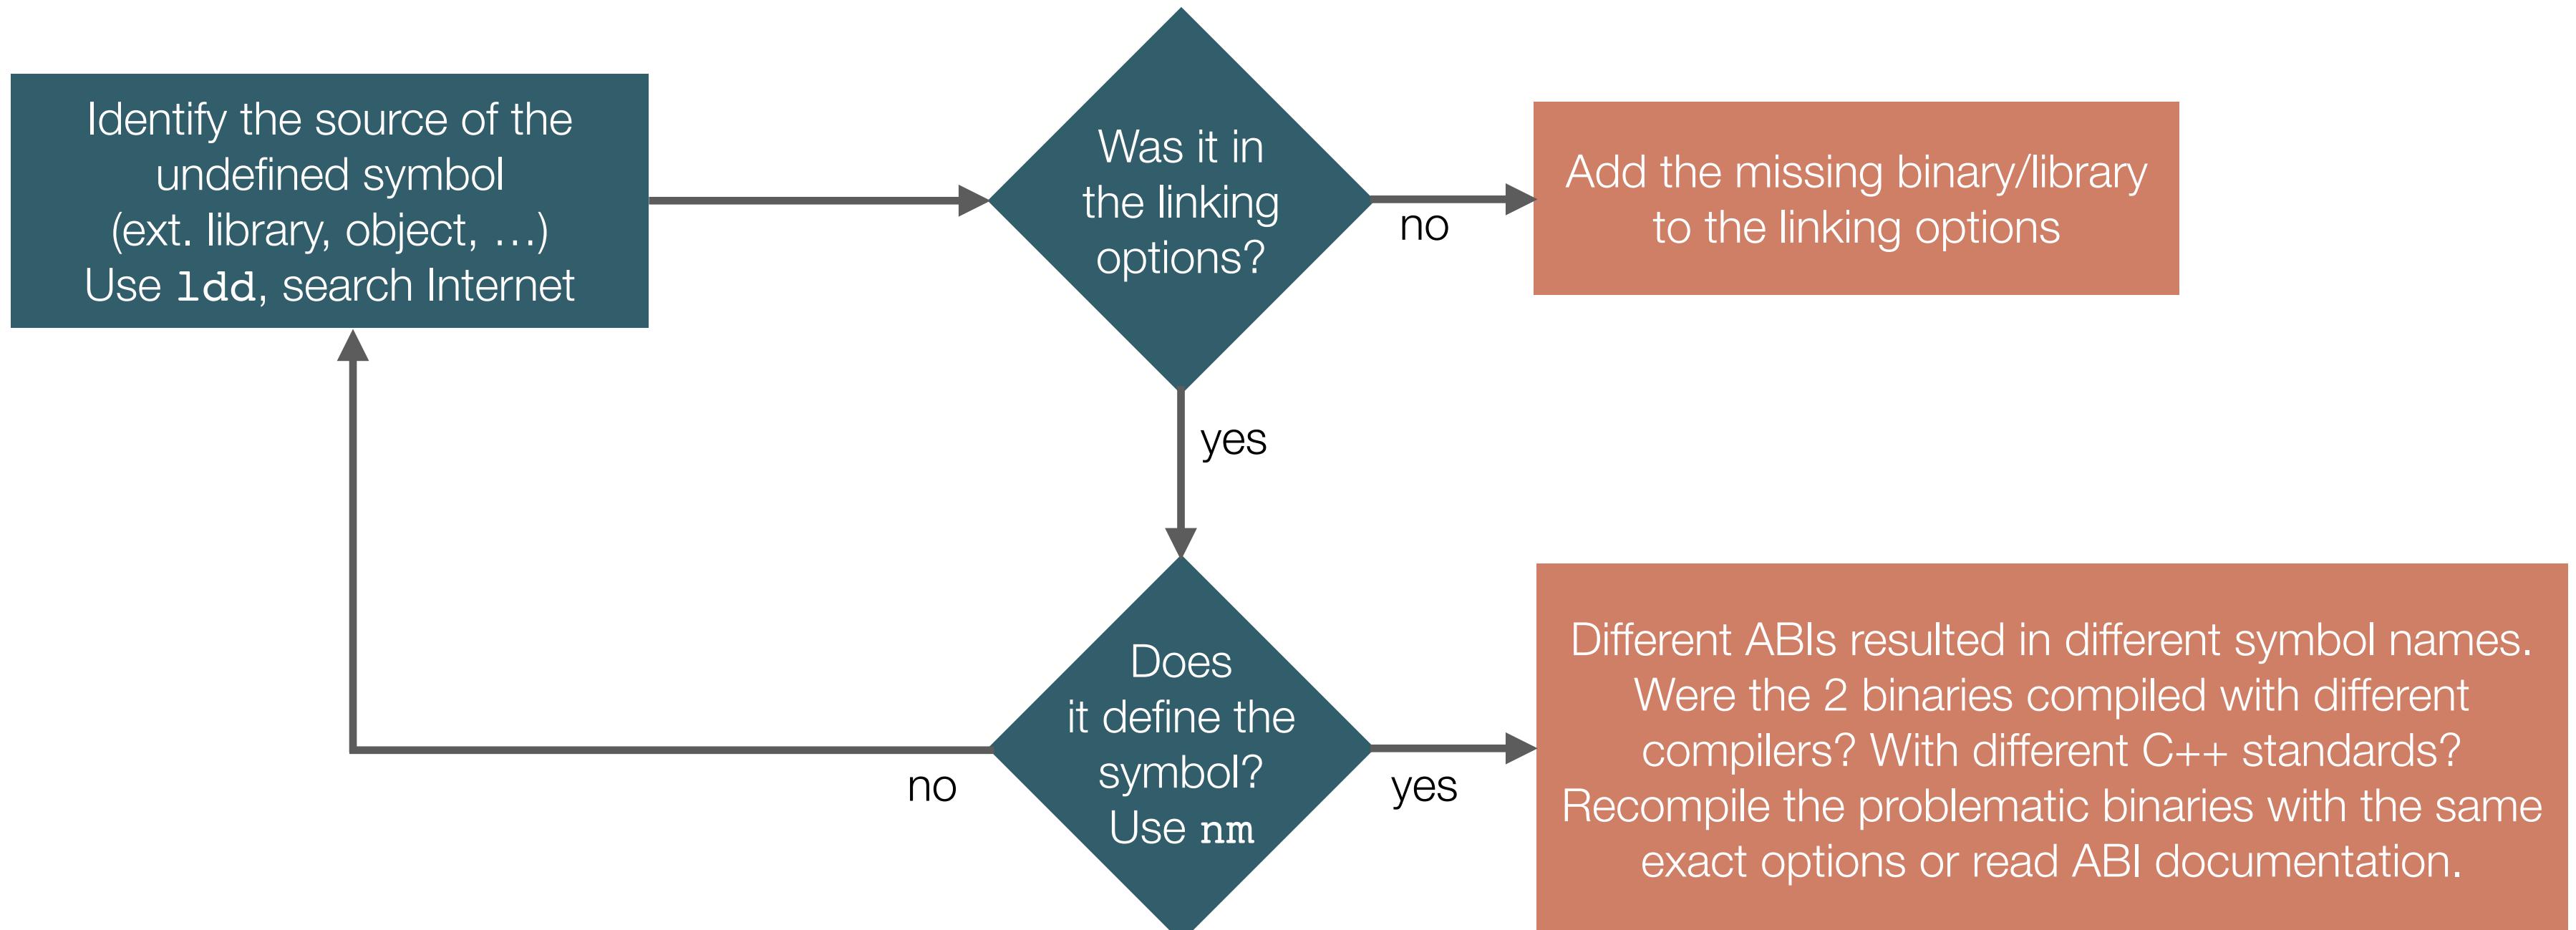
- Symbol names in C++ programs buying a monitor that can fit linker errors on one screen do not impact C++ standards
- However, C++ symbols are not part of C++ standards
- `c++filt` can help



```
$ cat main.cpp
#include <string>
std::string s;
$ clang++ -fno-rtti -fno-exceptions main.cpp
$ nm main | c++filt
0000000000000000 U strlen
0000000000000000 U __Unwind_Resume
0000000000000000 T __Z8functionNSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEEE
0000000000000000 U __ZNKSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEE4sizeEv
0000000000000000 W __ZNSt11char_traitsIcE6lengthEPKc
0000000000000000 U __ZNSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEE6appendEPKcm
0000000000000000 U __ZNSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEE6appendERKS4_
0000000000000000 U __ZNSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEE7reserveEm
0000000000000000 U __ZNSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEEC1Ev
0000000000000000 U __ZNSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEED1Ev
0000000000000000 W __ZStplIcSt11char_traitsIcESaIcEENSt7__cxx11basic_stringIT_T0_T1_EEPKS5_RKS8_
$ echo __Z8functionNSt7__cxx11basic_stringIcSt11char_traitsIcESaIcEEE | c++filt
function(std::__cxx11::basic_string, std::allocator >)
```

# ‘undefined reference to’ error resolution

- Both static & dynamic, a symbol is still undefined after all linking options have been considered



- Rinse and repeat until it links

# ‘cannot open shared object file’ error resolution

---

- The binary was compiled with a dynamic link, but the linked library cannot be found at runtime.  
Use `ldd`, you’ll likely see ‘not found’ on one of the links.
- Have you messed with `rpath` or `runpath` linker options? If yes double-check that
- You can append the problematic library path to the environment variable `LD_LIBRARY_PATH`
- Very useful: `LD_DEBUG=libs ldd` shows the whole path resolution process
- Useful post for more info <https://bit.ly/3axmcwQ>

Loose ends

# Make the “Hello world” shell code

- Position-independent ASM

```
1 section .text
2 global _start
3
4 _start:
5 jmp msg ; jump & call trick
6 print:
7 mov rax, 1
8 mov rdi, 1
9 pop rsi
10 mov rdx, 12
11 syscall
12 mov rax, 60
13 xor edi, edi
14 syscall
15 msg:
16 call print
17 db "Hello world",0xa
```

# Make the “Hello world”

- Position-independent AS

- call will push the address of the next instruction on the stack

- here is it the address of the string!

```
1 section .text
2 global _start
3
4 _start:
5 jmp msg
6 print:
7 mov rax, 1
8 mov rdi, 1
9 pop rsi
10 mov rdx, 12
11 syscall
12 mov rax, 60
13 xor edi, edi
14 syscall
15 msg:
16 call print
17 db "Hello world",0xa
```

- it can be popped later in print

- print actually never returns

# You can write machine code directly in C!

“shellcode”

```
1 #include <sys/mman.h>
2 #include <string.h>
3 #include <stdlib.h>
4
5 const unsigned char code [] =
6 "\xeb\x1b\xb8\x01\x00\x00\x00\xbf\x01\x00\x00\x00\x5e\xba\x0c"
7 "\x00\x00\x00\x0f\x05\xb8\x3c\x00\x00\x00\x31\xff\x0f\x05\xe8"
8 "\xe0\xff\xff\xff\x48\x65\x6c\x6c\x6f\x20\x77\x6f\x72\x6c\x64\x0a";
9
10 int main(int argc, char **argv)
11 {
12 // It used to be easier to break computers...
13 // the commented code will likely end up in SIGSEGV on modern OSs
14 // for security data is tagged non-executable
15 /////////////////////////////////
16 // void (*hello)(void);
17 //
18 // hello = (void(*)(void))code;
19 // hello();
20 ///////////////////////////////
21
22 char *buf;
23 int prot = PROT_READ | PROT_WRITE | PROT_EXEC;
24 int flags = MAP_PRIVATE | MAP_ANONYMOUS;
25
26 buf = mmap(0, sizeof(code), prot, flags, -1, 0);
27 memcpy(buf, code, sizeof(code));
28
29 ((void (*)())buf)();
30
31 return EXIT_SUCCESS;
32 }
33
```

## You can write machine code directly in C!

- `mmap` create a new memory region with executable permissions
- The shellcode is then copied in this new region
- Look at `/proc/<pid>/maps` before and after `mmap` is called, you will see the new region being created!
- In general, be curious about `/proc/<pid>`, it contains all the information about the program memory space
- Log from debugger investigation <https://bit.ly/304b9VP>

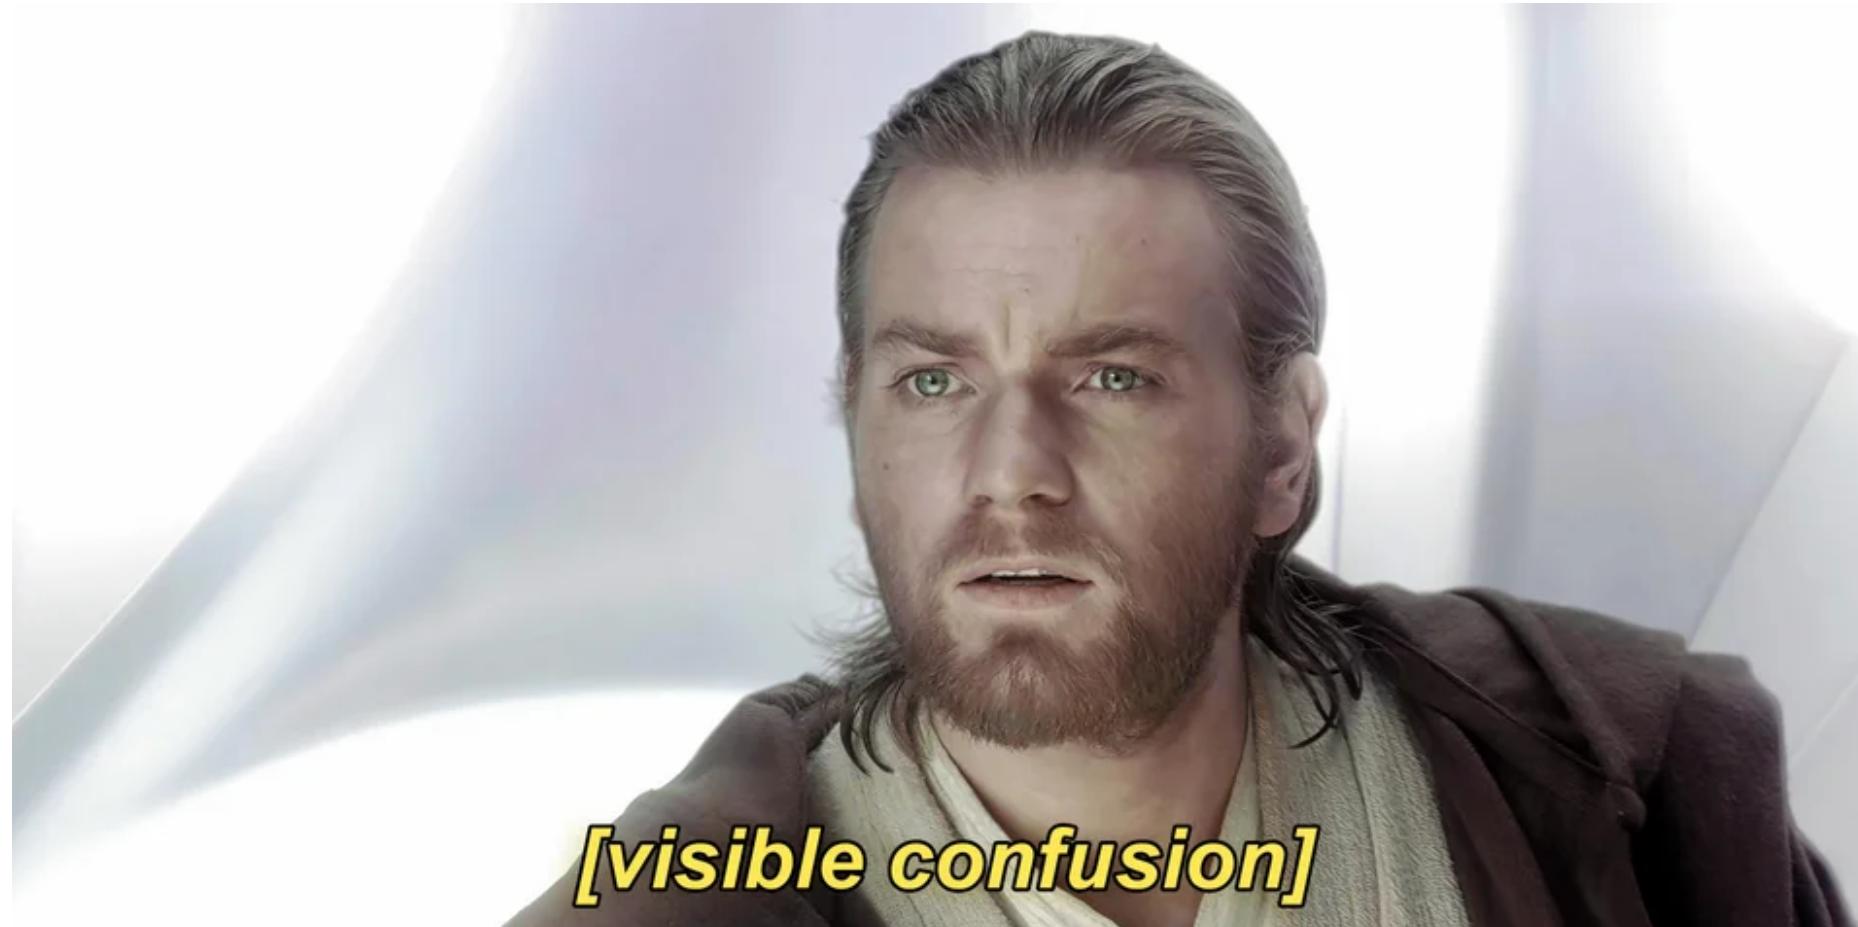
```
21
22 char *buf;
23 int prot = PROT_READ | PROT_WRITE | PROT_EXEC;
24 int flags = MAP_PRIVATE | MAP_ANONYMOUS;
25
26 buf = mmap(0, sizeof(code), prot, flags, -1, 0);
27 memcpy(buf, code, sizeof(code));
28
29 ((void (*)(void))buf)();
30
31 return EXIT_SUCCESS;
32 }
33
```

# Making the “Hello world” shell code

---

- Assemble and transform into a binary string

```
nasm -f elf64 hello-pi.asm -o hello-pi.o # assemble
ld hello-pi.o -o hello-pi # link
objcopy -O binary hello-pi hello-pi.bin # extract binary section
hexdump -v -e '"\\"x" 1/1 "%02x" "' hello-pi.bin # convert to C string
```



- Probably better after next lecture

# Making the “Hello world” shell code

- Assemble and transform into a binary string

```
nasm -f elf64 hello-pi.asm -o hello-pi.o # assemble
ld hello-pi.o -o hello-pi # link
objcopy -O binary hello-pi hello-pi.bin # extract binary section
hexdump -v -e '"\\"x" 1/1 "%02x" "' hello-pi.bin # convert to C string
```

- `objcopy` extract the loadable part of an ELF binary (i.e. code & constants)
- Use `objdump` on the result to see!
- The `hexdump` command convert bytes e.g. 0xa3 to the notation \xa3 used in C strings
- Probably better after next lecture

# Make your own: backtrace

- Call backtrace from function1

```
$./backtrace
==== BACKTRACE ====
frame 0: 0x4011cc
frame 1: 0x4011fc
frame 2: 0x40123c
frame 3: 0x7f4e7d29c0b3
=====
12.560000
$ addr2line -fp -e ./backtrace -a 0x4011cc 0x4011fc 0x40123c 0x7f4e7d29c0b3
0x0000000004011cc: function1 at /mnt/hgfs/binary-hack/lectures/2-elf/backtrace.c:36
0x0000000004011fc: function2 at /mnt/hgfs/binary-hack/lectures/2-elf/backtrace.c:41
0x00000000040123c: main at /mnt/hgfs/binary-hack/lectures/2-elf/backtrace.c:48
0x00007f4e7d29c0b3: ?? ???:0
```

- Why the last address is so different?  
What are the ?? ?? ?  
Why is it random between executions?  
I'll explain later

# Make your own: backtrace

---

- Call `backtrace` from `function1`
  - The mysterious address was a live memory address where dynamically linked code was loaded (it was the system loader calling `main`)
  - This address can't be resolved after the program terminated, hence the ???
  - It is randomised for security reasons, that's called Address Space Layout Randomization (ASLR)
  - ASLR makes buffer overflow exploits much harder to create

I'll explain later

Wrap up

# Summary

---

- The program memory space is split into different segments with different purposes
- The stack stores temporary data related to local variables
- Debuggers are invaluable development tools which allow to have full visibility & control of the program's execution
- The ELF binary format splits code and data into sections which are mapped onto the segments during execution
- The ELF format allows linking, i.e. the modular creation of an executable image from different binaries

# Next lecture

---

- More about the debugger, this time from an actual debugging point of view
- The horrors of memory bugs and some strategies to try to systematically deal with them
- Code profiling and optimisation strategy
- The roofline performance model
- Practical profiling and performance measurement for two simple linear algebra routines