

```

} while ((uint64_t)uVar15 < CONCAT44(in_register_0000000c, arg4));
} else {
    iVar13 = CONCAT44(in_register_0000000c, arg4) * uVar12;
    uVar7 = arg1 + iVar13 * 8;
    puVar10 = (uint64_t *)((uVar12 * 8 + 8) * CONCAT44(in_register_0000000c, arg4) + arg1);
    if (puVar5 < puVar10 && uVar7 < CONCAT44(in_register_0000000c, arg4) * 8 * uVar12 + arg2 + 1) {
        uVar15 = 0;
        uVar8 = 0;
        goto code_r0x004015c0;
    }
    uVar15 = 0;
    uVar8 = 0;
    if (uVar7 < (uint64_t)(CONCAT44(in_register_00000014, arg3) + CONCAT44(in_register_0000000c, arg4) * 8)
        && ((uint64_t *)CONCAT44(in_register_00000014, arg3) <= puVar10 &&
              puVar10 != (uint64_t *)CONCAT44(in_register_00000014, arg3))) goto code_r0x004015c0;
    if (uVar2 == 0) {
        iVar13 = 0;
    } else {
        iVar16 = arg1 + 0xe0 + iVar13 * 8;
        auVar17 = vbroadcastsd_avx(*puVar5);
        iVar13 = 0;
        uVar7 = uVar14 & 0xfffffffffffffe;
        do {
            auVar18 = vmulpd_avx(auVar17, *(undefined (*) [32])
                (CONCAT44(in_register_00000014, arg3) + iVar13 * 8));
            auVar19 = vmulpd_avx(auVar17, *(undefined (*) [32])
                (CONCAT44(in_register_00000014, arg3) + 0x20 + iVar13 * 8));
            auVar20 = vmulpd_avx(auVar17, *(undefined (*) [32])
                (CONCAT44(in_register_00000014, arg3) + 0x40 + iVar13 * 8));
            auVar21 = vmulpd_avx(auVar17, *(undefined (*) [32])
                (CONCAT44(in_register_00000014, arg3) + 0x60 + iVar13 * 8));
            auVar18 = vmoveupd_avx(auVar18);
            *(undefined (*) [32])(iVar16 + -0xe0 + iVar13 * 8) = auVar18;
            auVar18 = vmoveupd_avx(auVar19);
            *(undefined (*) [32])(iVar16 + -0xc0 + iVar13 * 8) = auVar18;
            auVar18 = vmoveupd_avx(auVar20);
            *(undefined (*) [32])(iVar16 + -0xa0 + iVar13 * 8) = auVar18;
            auVar18 = vmoveupd_avx(auVar21);
            *(undefined (*) [32])(iVar16 + -0x80 + iVar13 * 8) = auVar18;
            auVar18 = vmulpd_avx(auVar17, *(undefined (*) [32])
                (CONCAT44(in_register_00000014, arg3) + 0x80 + iVar13 * 8));
            auVar19 = vmulpd_avx(auVar17, *(undefined (*) [32])
                (CONCAT44(in_register_00000014, arg3) + 0xa0 + iVar13 * 8));
            auVar20 = vmulpd_avx(auVar17, *(undefined (*) [32])

```

# Reverse engineering and some hacking

*Introduction to executable binaries*  
Antonin Portelli

- Reverse Engineering (RE)
- Introduction to code injection hacks
- Wrap up

# Reverse Engineering (RE)

# RE overview

---

- Software reverse engineering consists in **reconstructing a software's source code** from the binary only
- This is **very non-trivial** because close to no high-level structure survives the compilation process
- Some applications:
  1. **Offensive security**: understand and circumvent security measures in computer systems
  2. **Defensive security**: understand and protect against offensive softwares (malware, viruses, ...)
  3. **Open-source development**: develop open-source variants of proprietary software (nouveau driver, Samba protocol, ...)

# Connection to HPC?

---

- Not direct, mainly for your own curiosity, however...
- The RE community provides a lot of high-quality materials on binary analysis very useful for low-level development
  - cf. for example <https://bit.ly/3st1C6Y>
- RE tools are quite powerful to just browse and understand an executable binary
- In an HPC context “reverse-engineering” your own programs can be quite instructive regarding the optimisations the compiler made

# Radare2/Rizin

---



- **Open-source**, 15 years old community based **RE toolset**
- Powerful **command-line swiss-army knife** to navigate and understand disassembled binaries
- Radare2 & Rizin are (for the moment) **essentially the same software...**  
... a lot of politics beyond the recent separation



# Ghidra

---

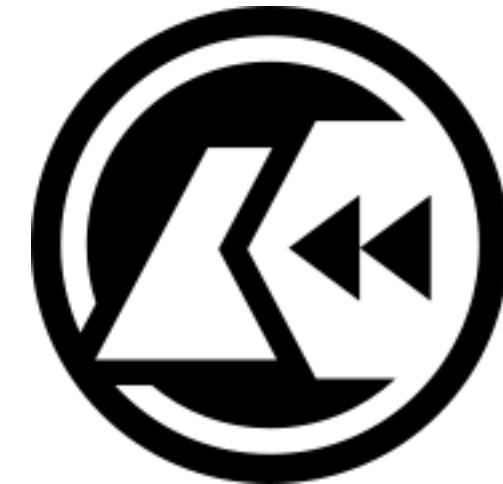


- **NSA RE framework**, first mentioned in a 2017 WikiLeaks publication, then **declassified** in 2019
- Now **open-source** and freely available
- Provide many tools only available in > \$100 commercial RE software, in particular a **powerful decompiler**

# Cutter

---

- Best of both worlds: **GUI for Rizin** and packs the **Ghidra decompiler**
- Multi-platform, ambitious with very active development
- Provide a **graph view of assembly code**, very useful to follow jumps, loops, conditionals, ...
- Provide **visual representations** of many other aspects of the binary (data, segments, sections, ...)
- Let's try!



# Introduction to code injection hacks

# What is code injection

---

- Code injection consists in constraining an executable binary to execute external arbitrary code at run time
- It can be used to “hook” functions, i.e. to wrap-up an existing function in additional code, for example for instrumentation purposes
- It can be used to add features to a closed-source program
- It can be used to build exploits: force a program to reveal sensitive memory content, circumvent a protection...

# Hooking dynamically linked functions

---

- By far the easiest
- If you set the environment variable `LD_PRELOAD` to the filename of a shared library, this library will be dynamically linked when a program runs before any other
- If this library contains a symbol called in the program, it will override the original dynamical link
- In short you can replace any function that comes from a shared library by whatever you want!
- Of course the system will ignore it if the executable has any form of privilege, that would be a big vulnerability

# Basic example: code

---

- Main program

```
1 int main(void)
2 {
3     func(15431);
4     return EXIT_SUCCESS;
5 }
```

- Target function dynamically linked to the main program

```
1 void func(const int id)
2 {
3     printf("calling original with id %d\n", id);
4 }
```

- Replacement function in another library

```
1 void func(const int id)
2 {
3     printf("calling variant with id %d\n", id);
4 }
```

# Basic example: execution

---

- Use `LD_PRELOAD` to replace function

```
$ clang -g -fPIC -shared func1.c -o libfunc1.so
$ clang -g callfunc.c -o callfunc -lfunc1 -L.
$ ./callfunc
./callfunc: error while loading shared libraries: libfunc1.so: cannot open shared object file: No such file or directory
$ export LD_LIBRARY_PATH=$(pwd)
$ ./callfunc
calling original with id 15431
$ LD_PRELOAD=./libfunc2.so ./callfunc
calling variant with id 15431
```

- But what if we want to **still execute the original function?** e.g. just to wrap code around it

# Basic example: function wrapper

```
1 // RTLD_NEXT is a GNU extension
2 #define _GNU_SOURCE
3
4 #include <dlfcn.h>
5 #include <stdio.h>
6 #include <stdlib.h>
7
8 void func(const int id)
9 {
10    // static pointer to the target
11    static void (*target_func)(const int) = NULL;
12
13    if (!target_func)
14    {
15        // if the target pointer is NULL, load it using dlsym from libdl
16        target_func = dlsym(RTLD_NEXT, "func");
17        printf("Target loaded at %p\n", target_func);
18    }
19
20    // wrapper
21    printf("Do some things before func call...\n");
22    target_func(id);
23    printf("Do some things after func call...\n");
24 }
```

- libdl (POSIX) allows to get pointer on dynamic symbols

# Basic example: function wrapper

---

- The wrapper in action

```
$ clang -g -fPIC -shared funchook.c -o libfunchook.so -ldl
$ LD_PRELOAD=./libfunchook.so ./callfunc
Target loaded at 0x7f7014cc9110
Do some things before func call...
calling original with id 15431
Do some things after func call...
```

- This trick can be extremely useful to **track, debug or instrument calls to functions**

# Example: tracing memory allocations

---

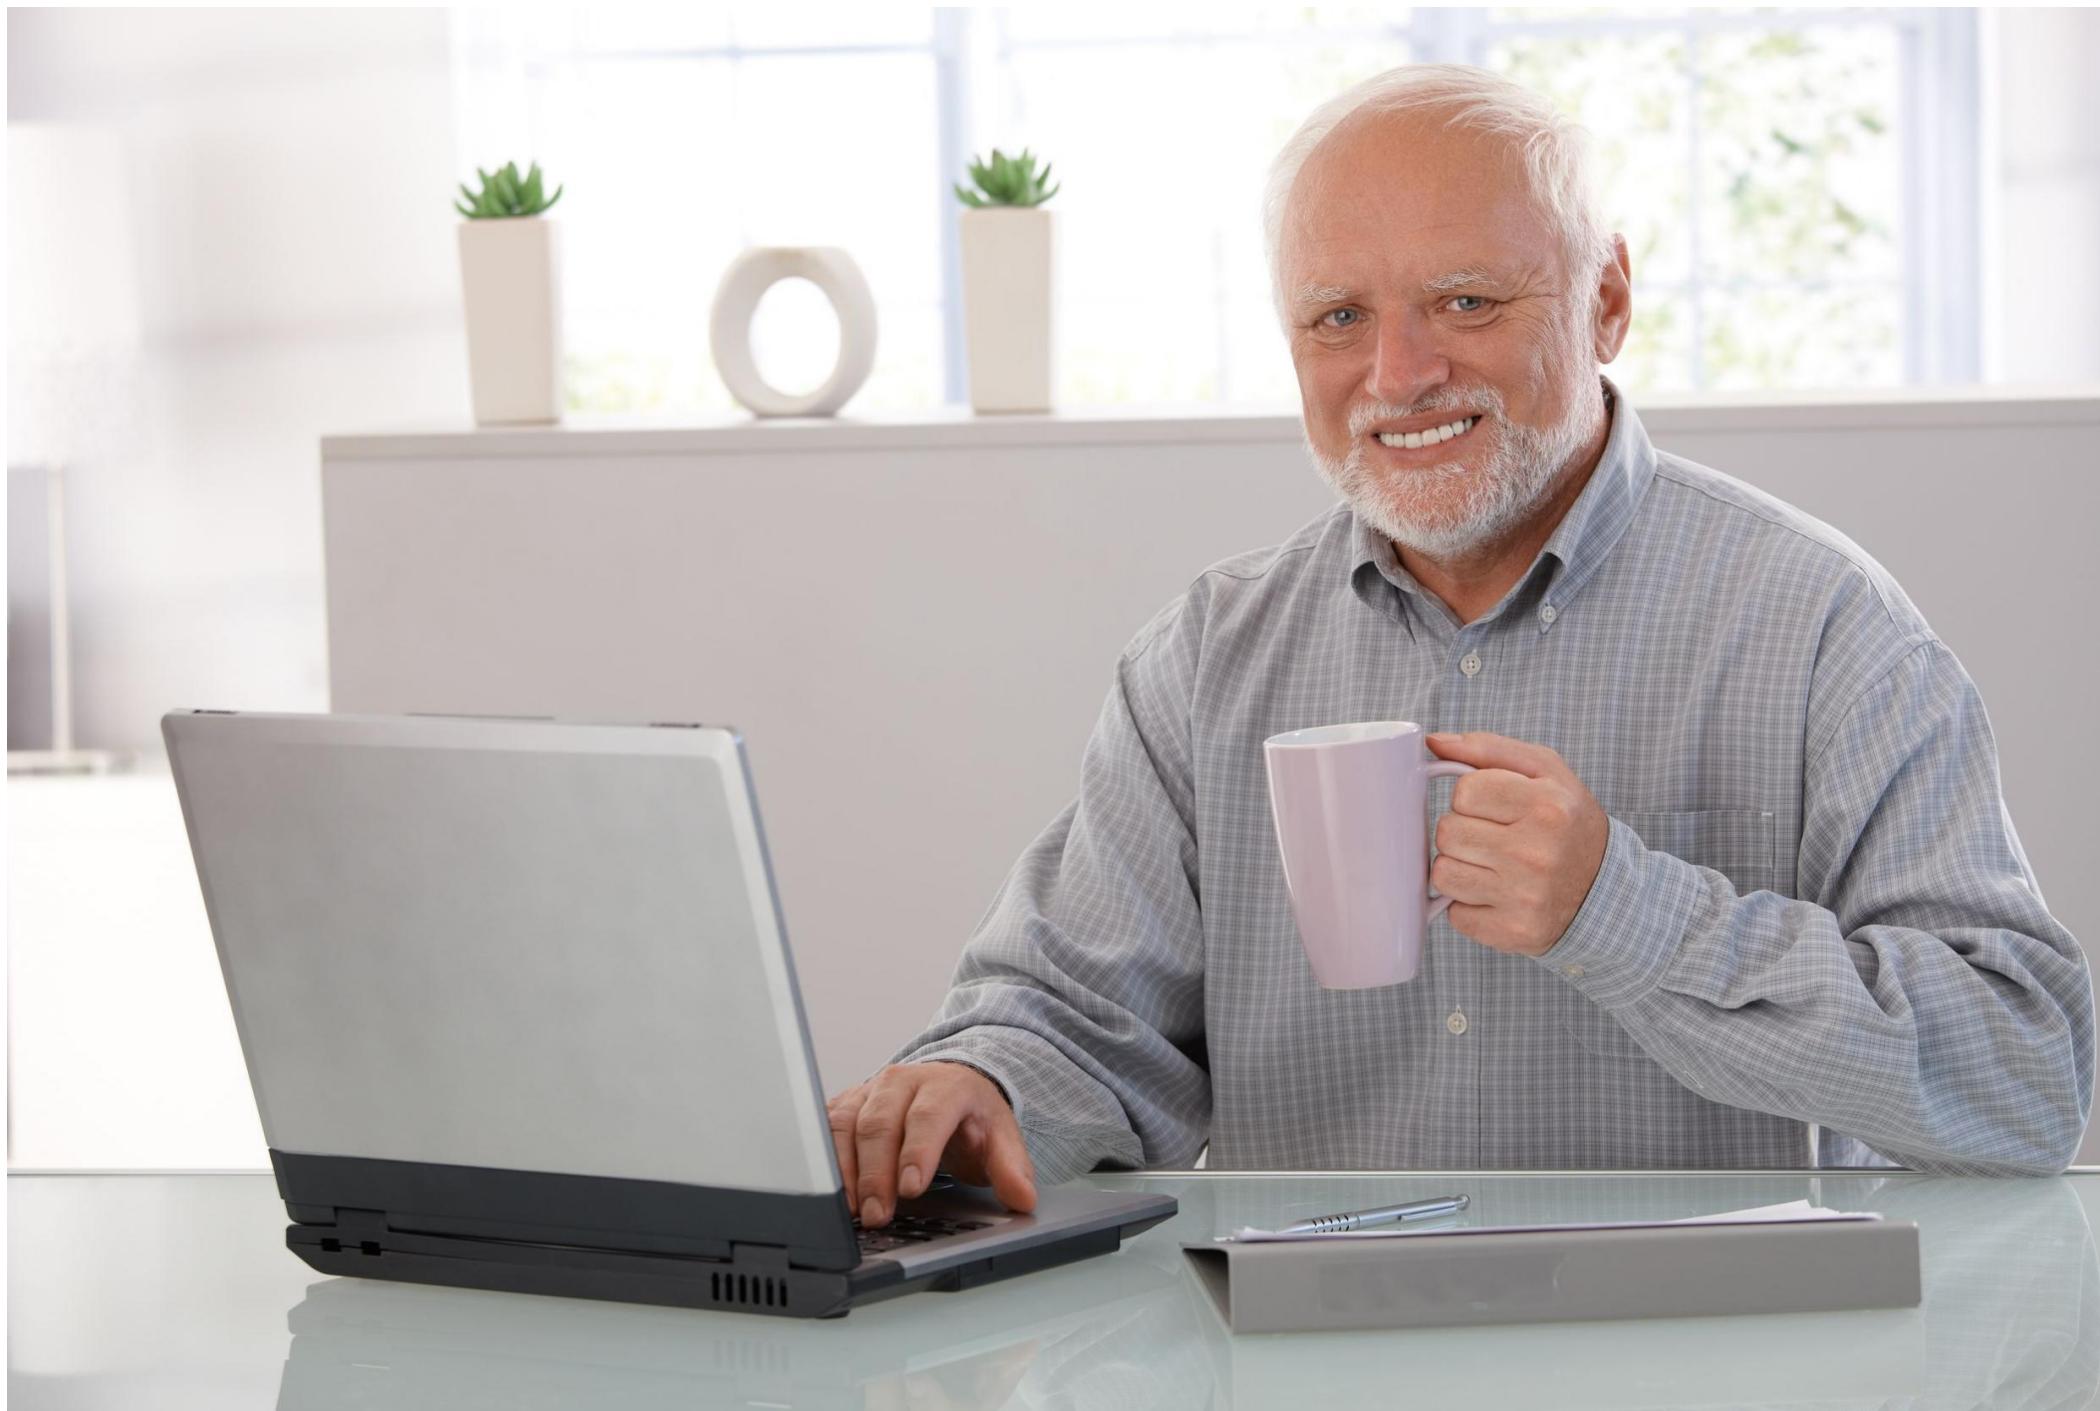
- Let's track `malloc` and `free` calls to produce statistics for memory allocations
- Wrapper

```
1 static size_t peak = 0, current = 0, cumul = 0;
2
3 void *malloc(size_t size)
4 {
5     static void *(*real_malloc)(size_t) = NULL;
6
7     if (!real_malloc)
8     {
9         real_malloc = dlsym(RTLD_NEXT, "malloc");
10        printf("Target function malloc loaded at %p\n", real_malloc);
11    }
12
13    void *p = real_malloc(size);
14    size_t real_size = malloc_usable_size(p);
15
16    current += real_size;
17    cumul += real_size;
18    peak = (current > peak) ? current : peak;
19    printf("==== MEM DEBUG === Allocated %zdB at %p (current %zdB, peak %zdB, cumulative %zdB)\n",
20          real_size, p, current, peak, cumul);
21
22    return p;
23}
24
```

# Example: tracing memory allocations

- Ok let's see the allocations in the `ls` command...

```
$ LD_PRELOAD=./libmemhookv1.so ls  
[1] 55639 segmentation fault (core dumped) LD_PRELOAD=./libmemhookv1.so ls --color=tty
```



# Example: tracing memory allocations

---

- The issue is that `printf` calls `malloc`, and now we put a `printf` call *inside* `malloc`, so **infinite recursion**
- We need to protect the `printf` call

```
static int in_hook = 0;
if (!in_hook)
{
    size_t real_size = malloc_usable_size(p);

    in_hook = 1;
    current += real_size;
    cumul += real_size;
    peak = (current > peak) ? current : peak;
    printf("==== MEM DEBUG ==== Allocated %zdB at %p (current %zdB, peak %zdB, cumulative %zdB)\n",
           real_size, p, current, peak, cumul);
    in_hook = 0;
}
```

# Example: tracing memory allocations

- It works!

```
$ LD_PRELOAD=./libmemhookv2.so ls
==== MEM DEBUG ====
Allocated 1032B at 0x55da077342d0 (current 1032B, peak 1032B, cumulative 1032B)
Target function free loaded at 0x7fa93569e850
==== MEM DEBUG ====
Freed 1032B at 0x55da077346e0 (current 0B, peak 1032B, cumulative 1032B)
Target function malloc loaded at 0x7fa93569e260
==== MEM DEBUG ====
Allocated 472B at 0x55da07734af0 (current 472B, peak 1032B, cumulative 1504B)
==== MEM DEBUG ====
Allocated 120B at 0x55da07734cd0 (current 592B, peak 1032B, cumulative 1624B)
==== MEM DEBUG ====
Allocated 1032B at 0x55da077346e0 (current 1624B, peak 1624B, cumulative 2656B)
==== MEM DEBUG ====
Freed 120B at 0x55da07734cd0 (current 1504B, peak 1624B, cumulative 2656B)
==== MEM DEBUG ====
Freed 1032B at 0x55da077346e0 (current 472B, peak 1624B, cumulative 2656B)
==== MEM DEBUG ====
Freed 472B at 0x55da07734af0 (current 0B, peak 1624B, cumulative 2656B)
...
...
```

- You can iterate on that to make it a **full memory profiler**
- For example give a **backtrace** for every call
- For C++ you can also change `new` and `delete`

# Other applications

---

- You can hook performance critical routines to make performance measurements
- You can hook MPI functions in an HPC code to see all network transfers and measure their speed
- You can hook CUDA library functions to trace the interaction with a GPU
- You can hook any buggy function if you can't run your program interactively in a debugger
- etc...

# Warnings about potential issues

---

- In C++ be careful about name mangling when using libdl
- HPC applications are often multi-threaded, be careful about the reentrancy of the wrapper
- The memory tracker would not work for a multi-threaded application. Do you see why? Can you improve it?
- What about statically linked functions?

# Hooking statically linked functions

---

- This is much harder: the target function code is part of the main program binary and its calls are **hardcoded jumps**
- You can `LD_PRELOAD` a replacement, but it will just be ignored as calls don't go through the dynamic linker
- You could imagine “de-linking” the binary, and then statically link it again with the replacement function
- De-linking is non-trivial because of relocations, but it is not impossible and practised in RE
- But here we'll do something else...

# Static function hack: strategy

---

- Remember: when a program runs **its code is loaded somewhere in memory** (in the text segment)
- We can do the following
  1. **Load a replacement** with `LD_PRELOAD`
  2. Find the address of the target function, and **overwrite its first instructions with a jump** to the replacement
- Step 2. will need to be executed **before main** starts, and the jump will have to be **coded directly in binary**

# Static function hack: strategy

---

## Normal execution

```
main:  
  
    ; some code  
  
    call func  
  
    ; more code  
  
func:  
    push rbp  
    mov rbp, rsp  
    sub rsp, 0x10  
  
    ; function code  
  
    ret
```

mov r10, new\_func  
jmp r10

```
ctor:  
  
    ; hack the function  
    ; in main program  
  
new_func:  
  
    ; replacement code  
  
ret
```

# Static function hack: strategy

## Normal execution

```
main:  
    ; some code  
    call func  
    ; more code  
  
func:  
    push rbp  
    mov rbp, rsp  
    sub rsp, 0x10  
  
    ; function code  
  
    ret
```



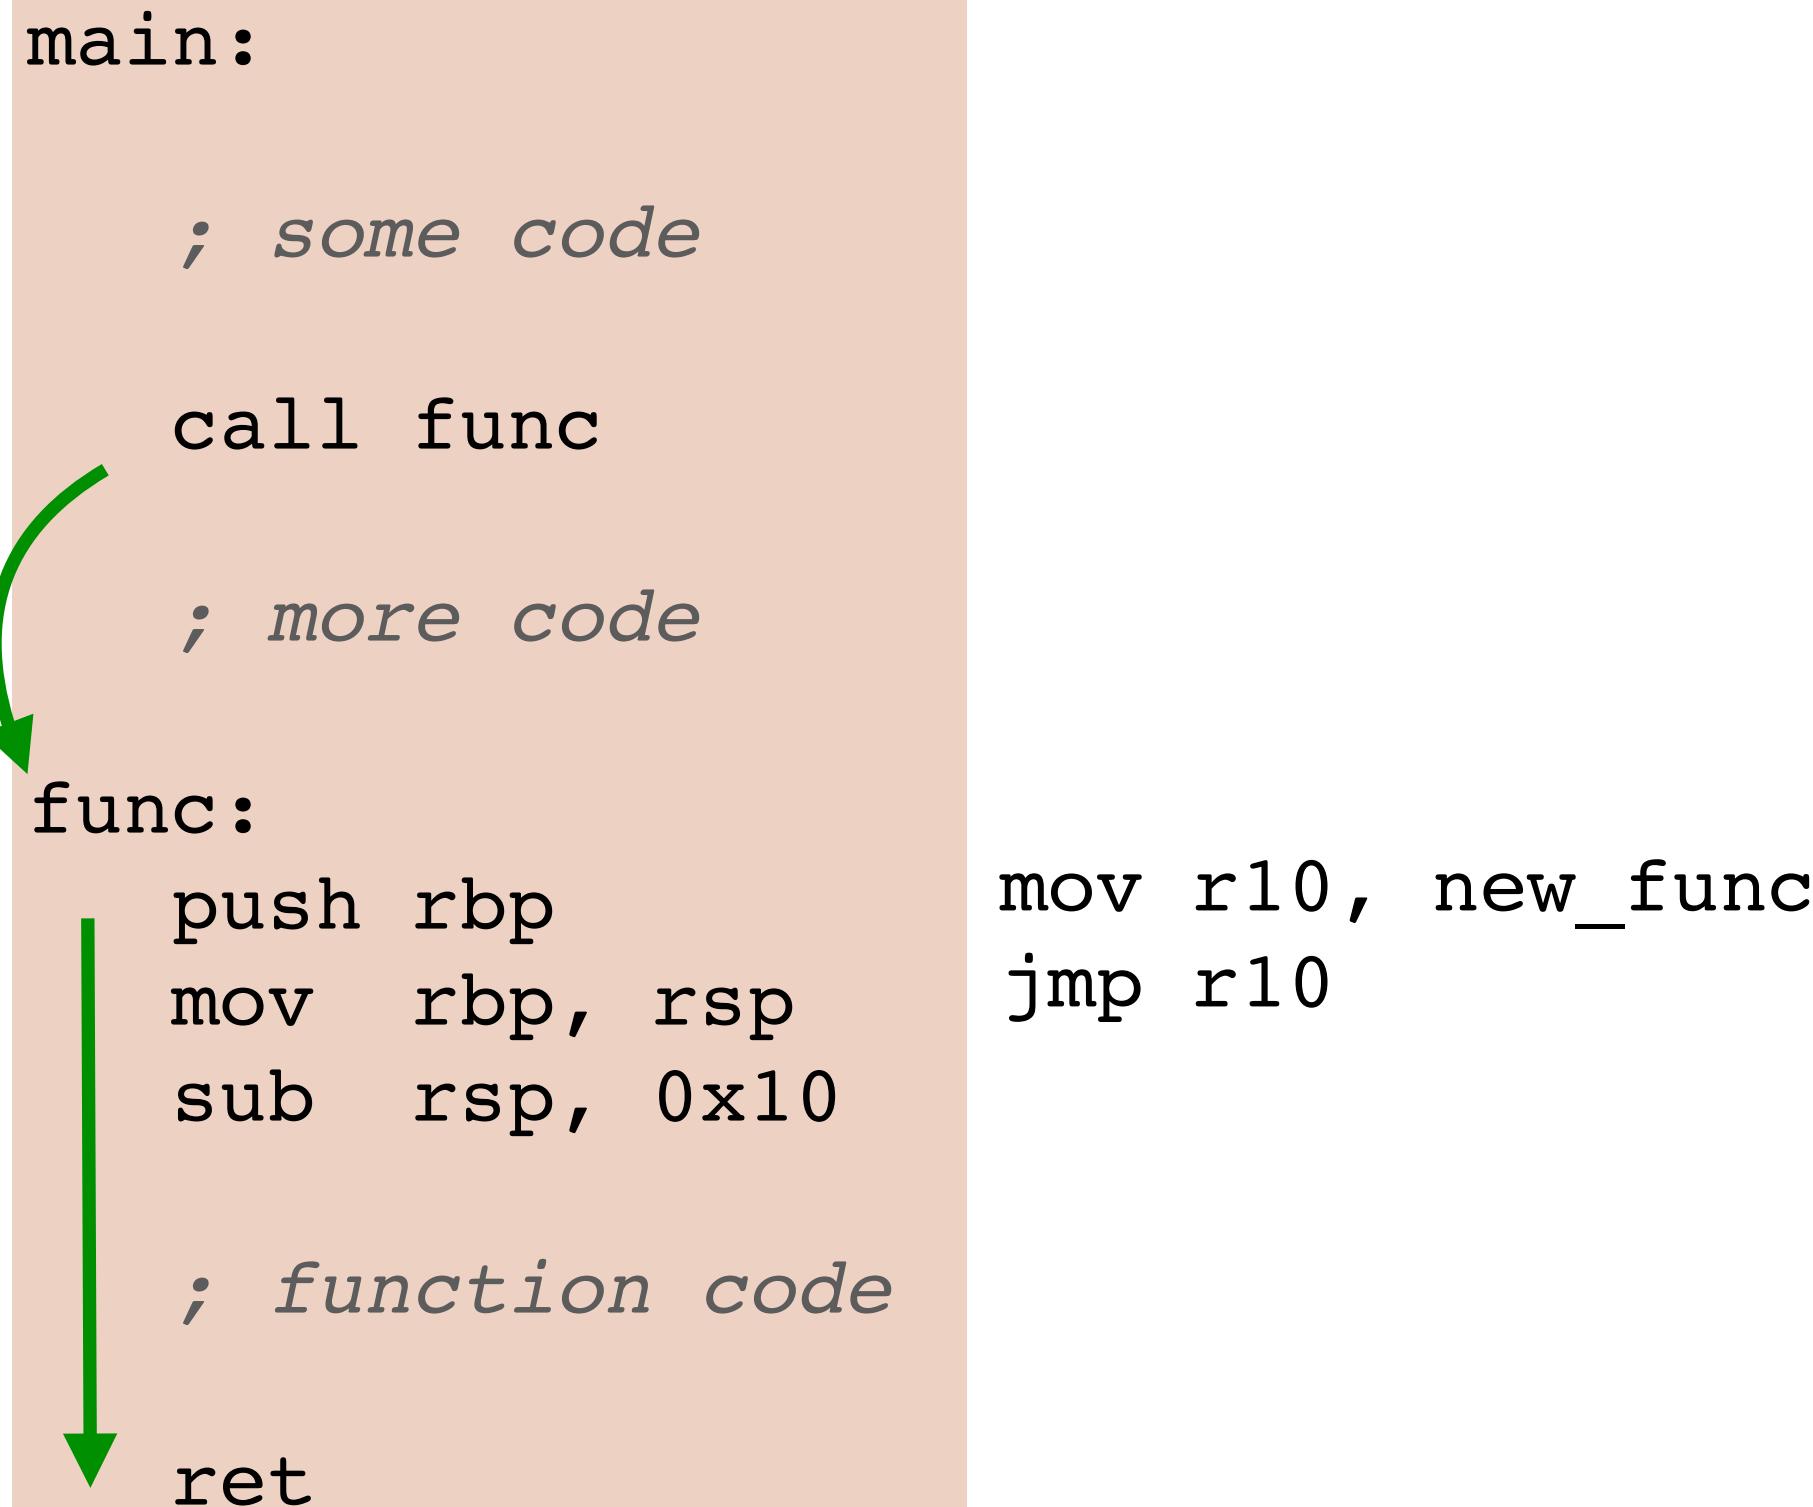
```
        mov r10, new_func  
        jmp r10
```

```
ctor:  
    ; hack the function  
    ; in main program  
  
new_func:  
    ; replacement code  
  
    ret
```

# Static function hack: strategy

## Normal execution

```
main:  
    ; some code  
    call func  
    ; more code  
  
func:  
    push rbp  
    mov rbp, rsp  
    sub rsp, 0x10  
    ; function code  
    ret
```

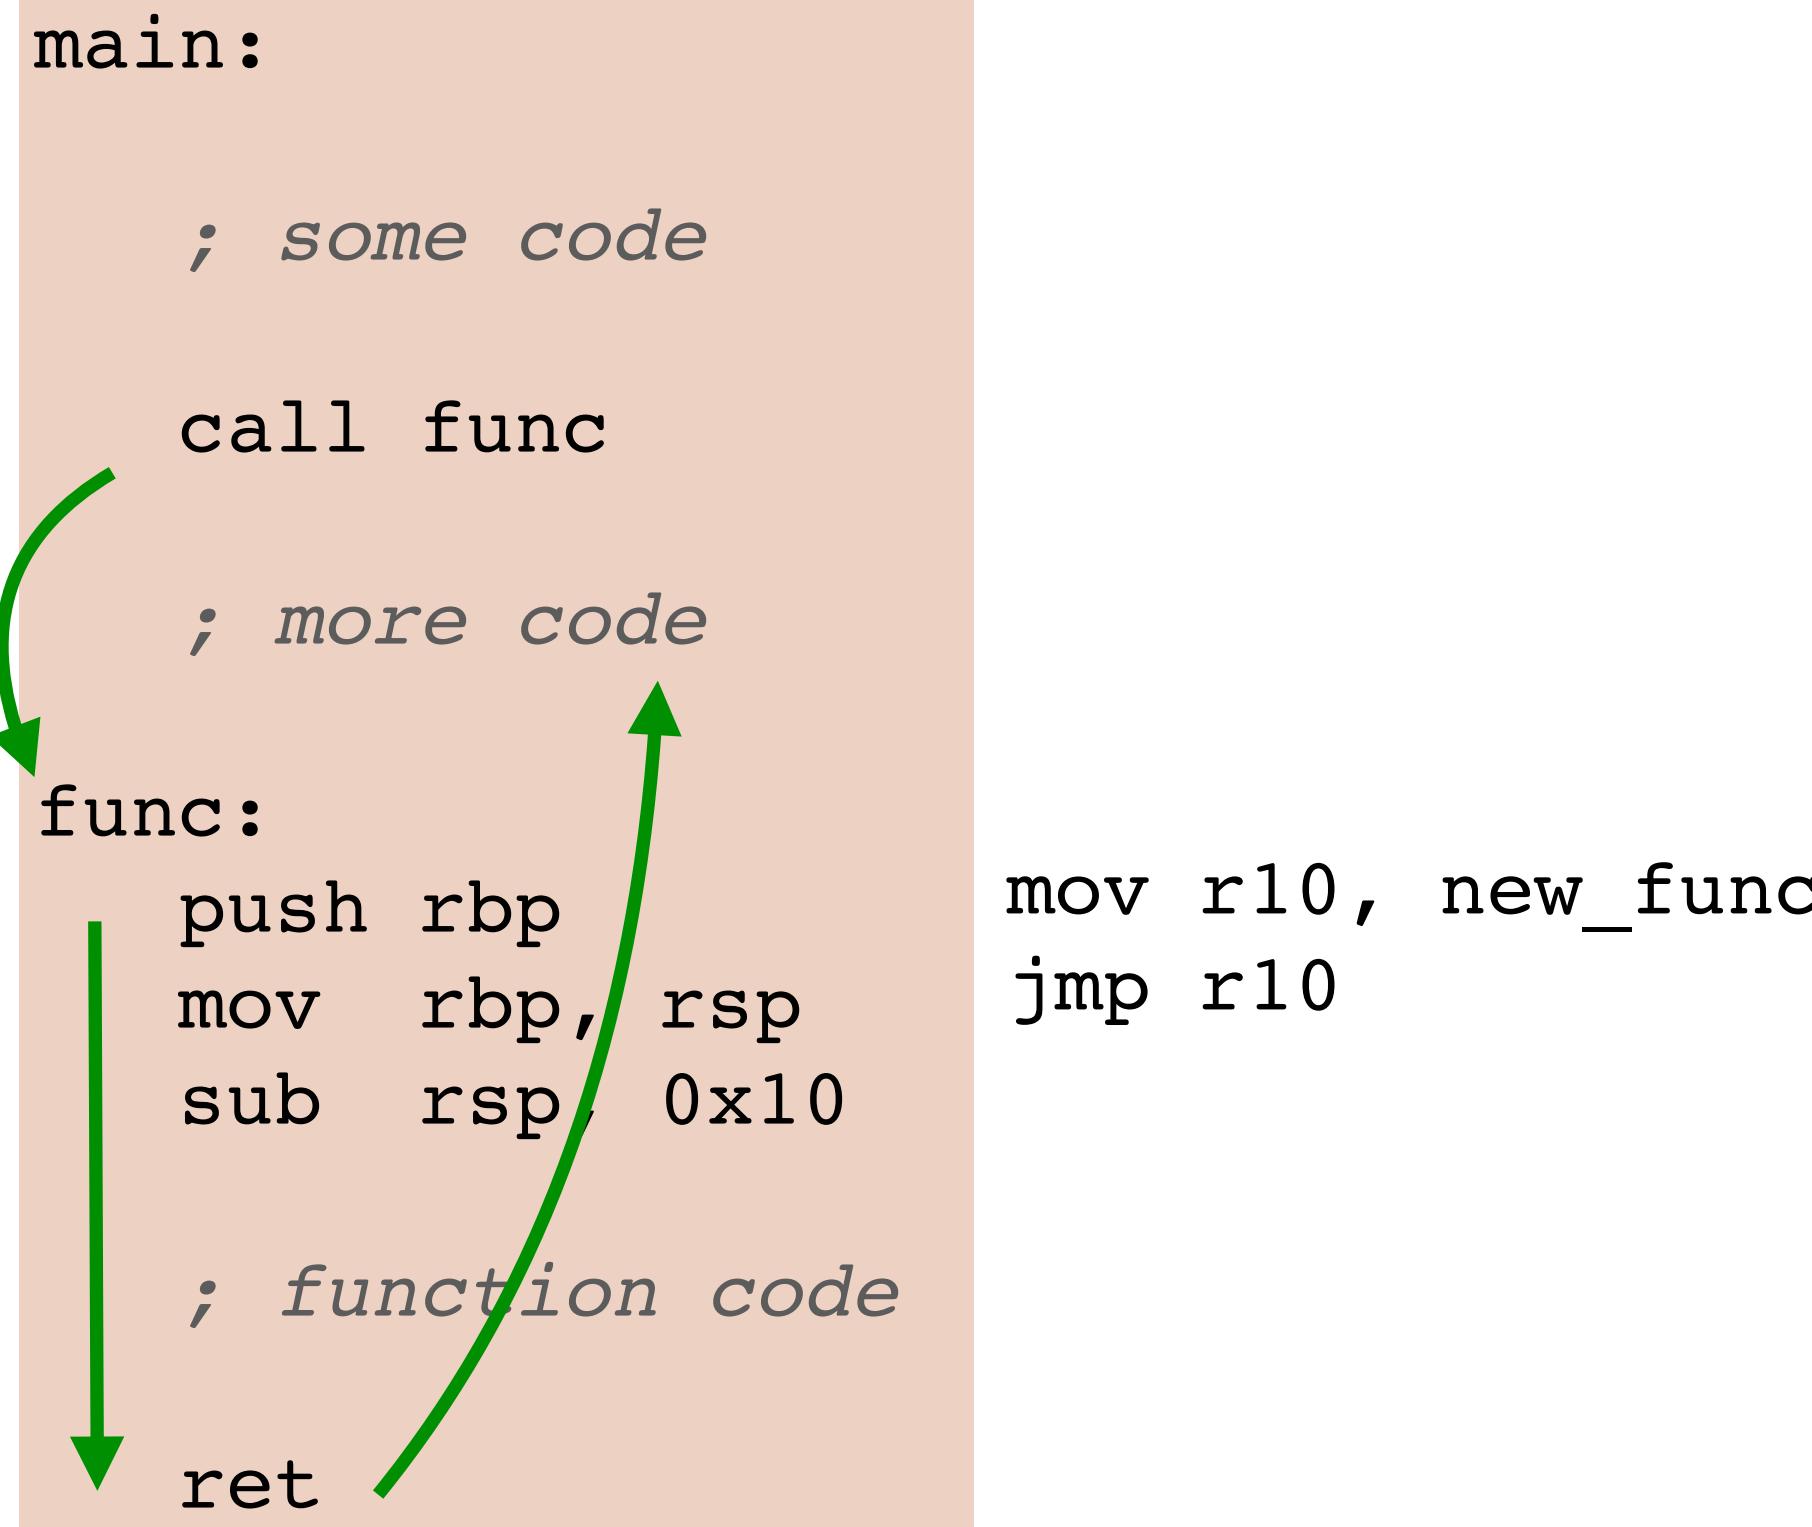


A green arrow starts at the 'call func' instruction in the main section, points down to the start of the func section, and then loops back up to the 'ret' instruction in the func section.

```
ctor:  
    ; hack the function  
    ; in main program  
  
new_func:  
    ; replacement code  
    ret
```

# Static function hack: strategy

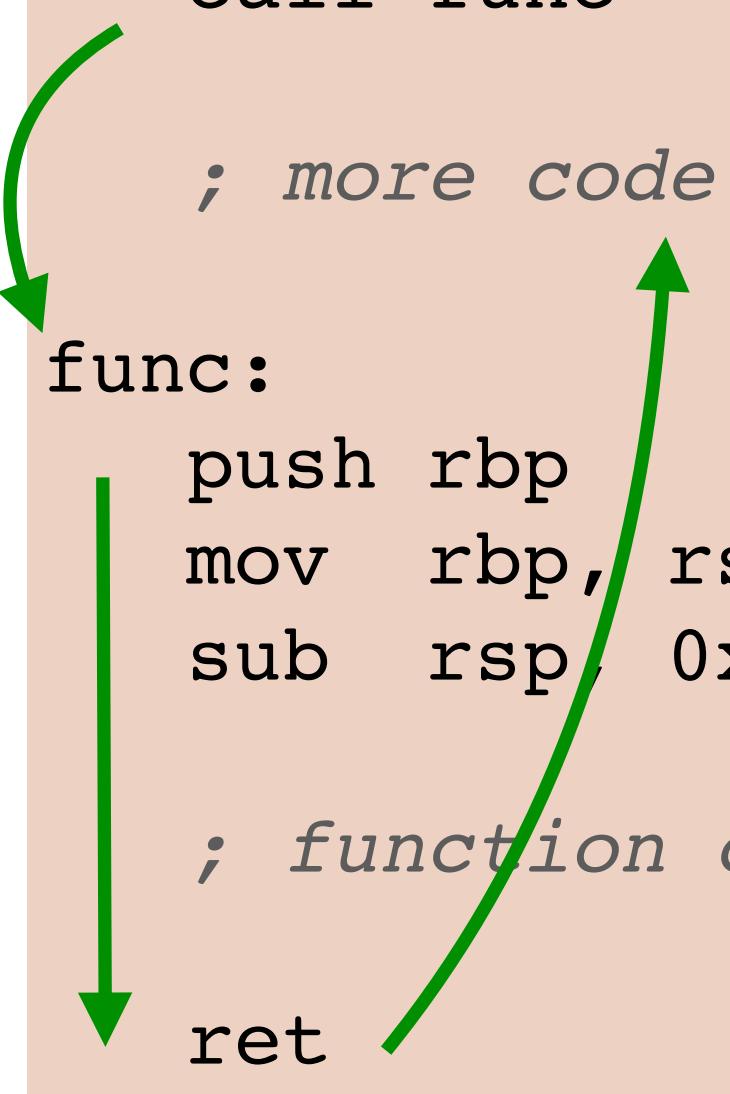
Normal execution



```
ctor:  
    ; hack the function  
    ; in main program  
  
new_func:  
    ; replacement code  
  
ret
```

# Static function hack: strategy

```
main:  
    ; some code  
    call func  
    ; more code  
  
func:  
    push rbp  
    mov rbp, rsp  
    sub rsp, 0x10  
    ; function code  
    ret
```

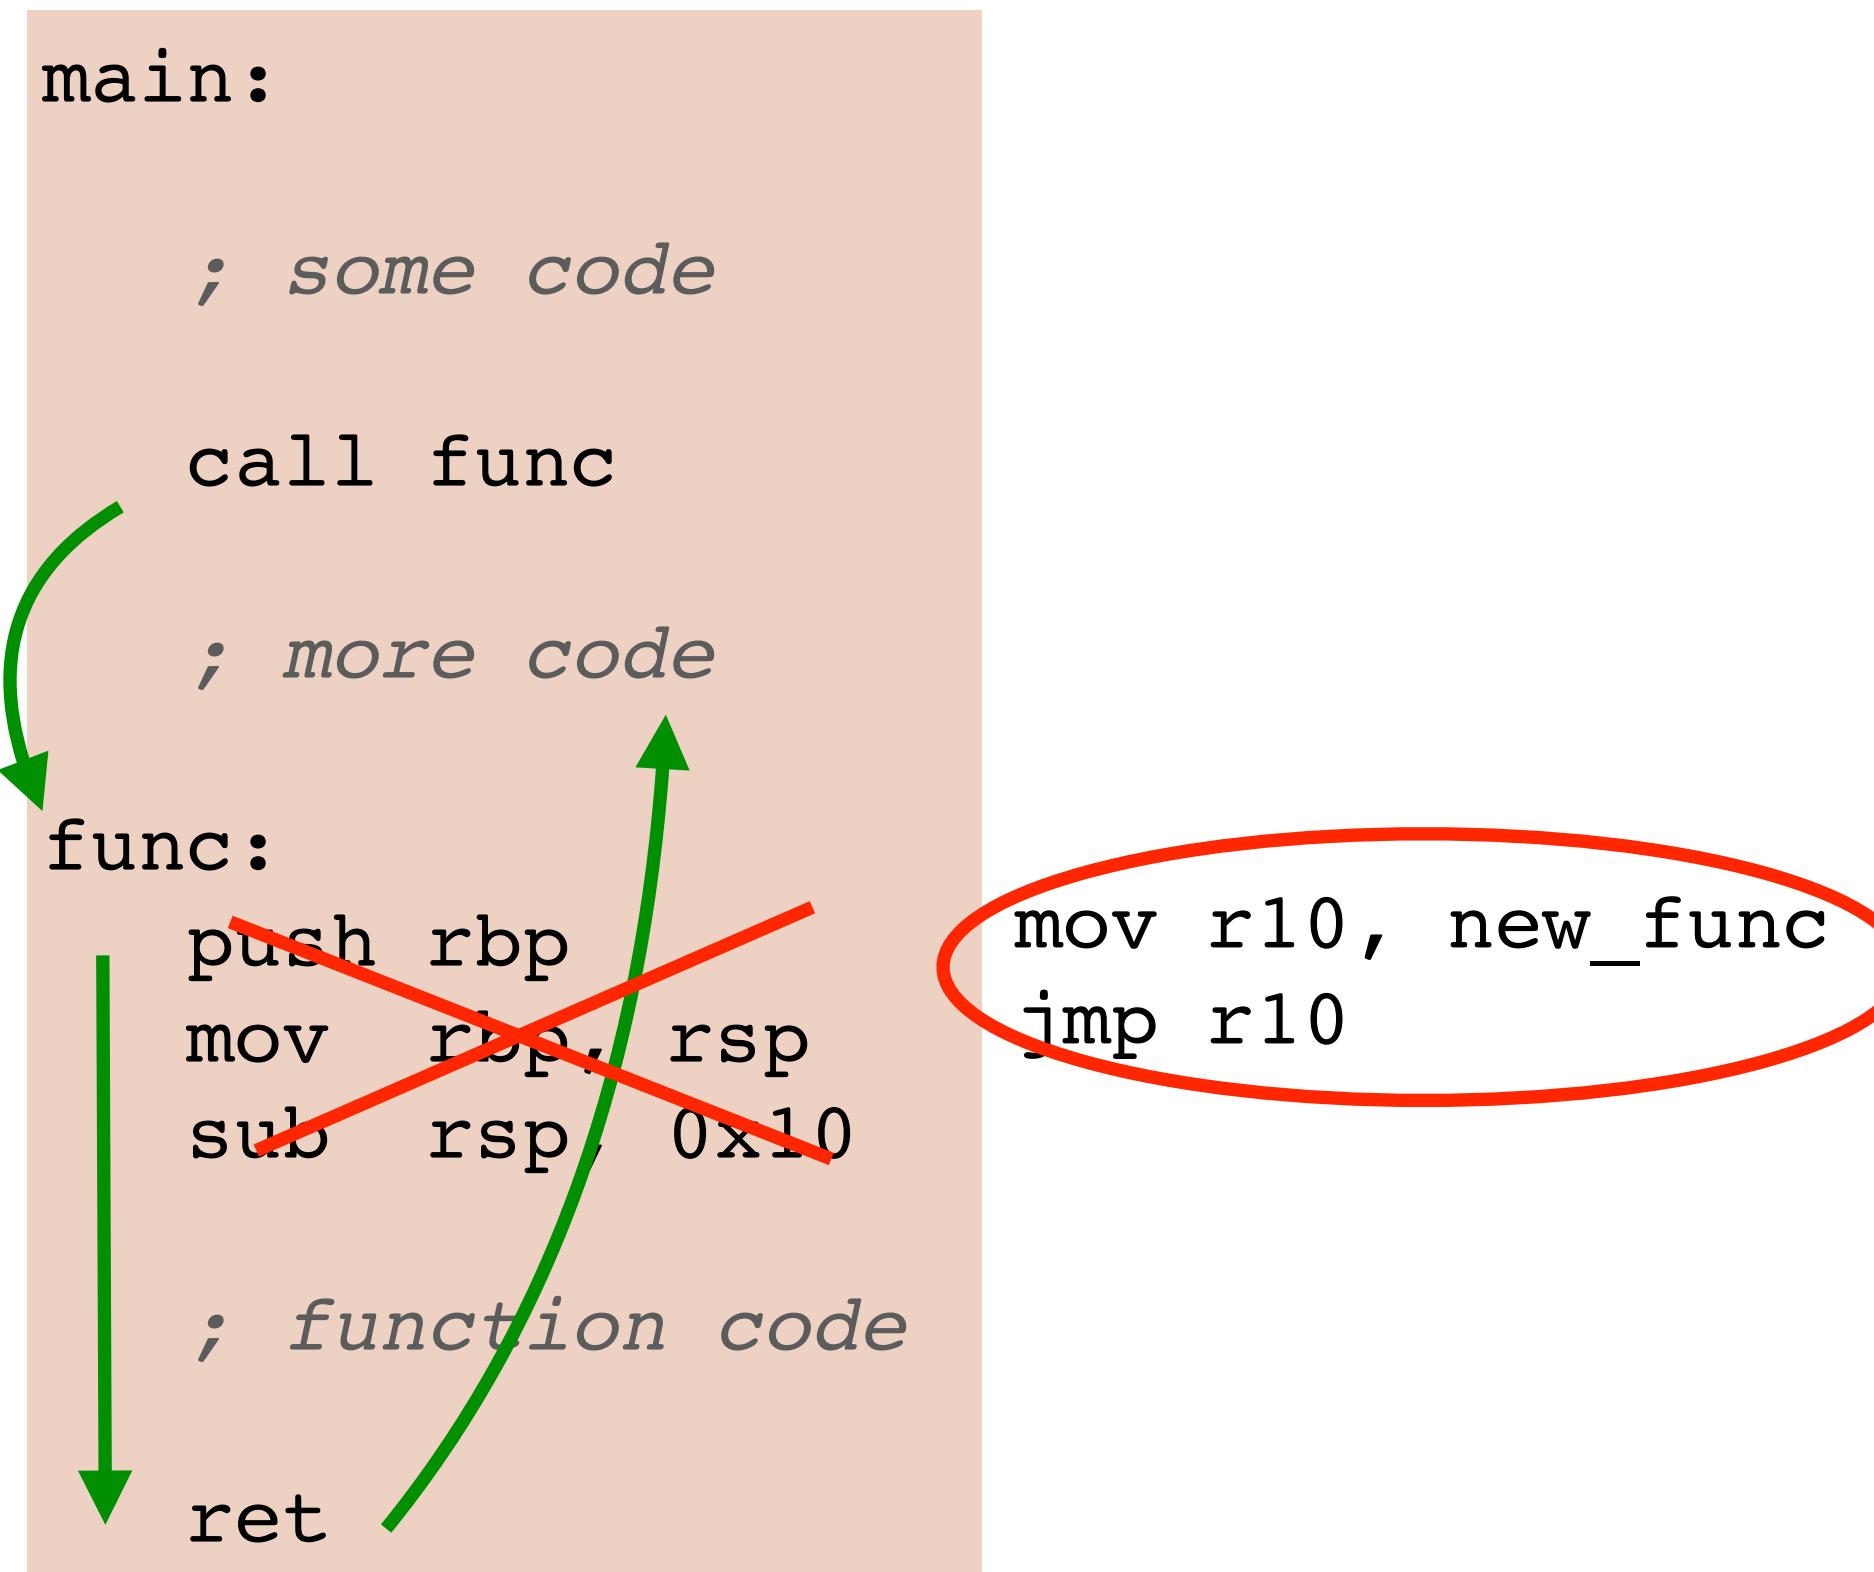


```
    mov r10, new_func  
    jmp r10
```

Modified execution

```
ctor:  
    ; hack the function  
    ; in main program  
  
new_func:  
    ; replacement code  
    ret
```

# Static function hack: strategy



Modified execution

Before main

```
ctor:  
    ; hack the function  
    ; in main program
```

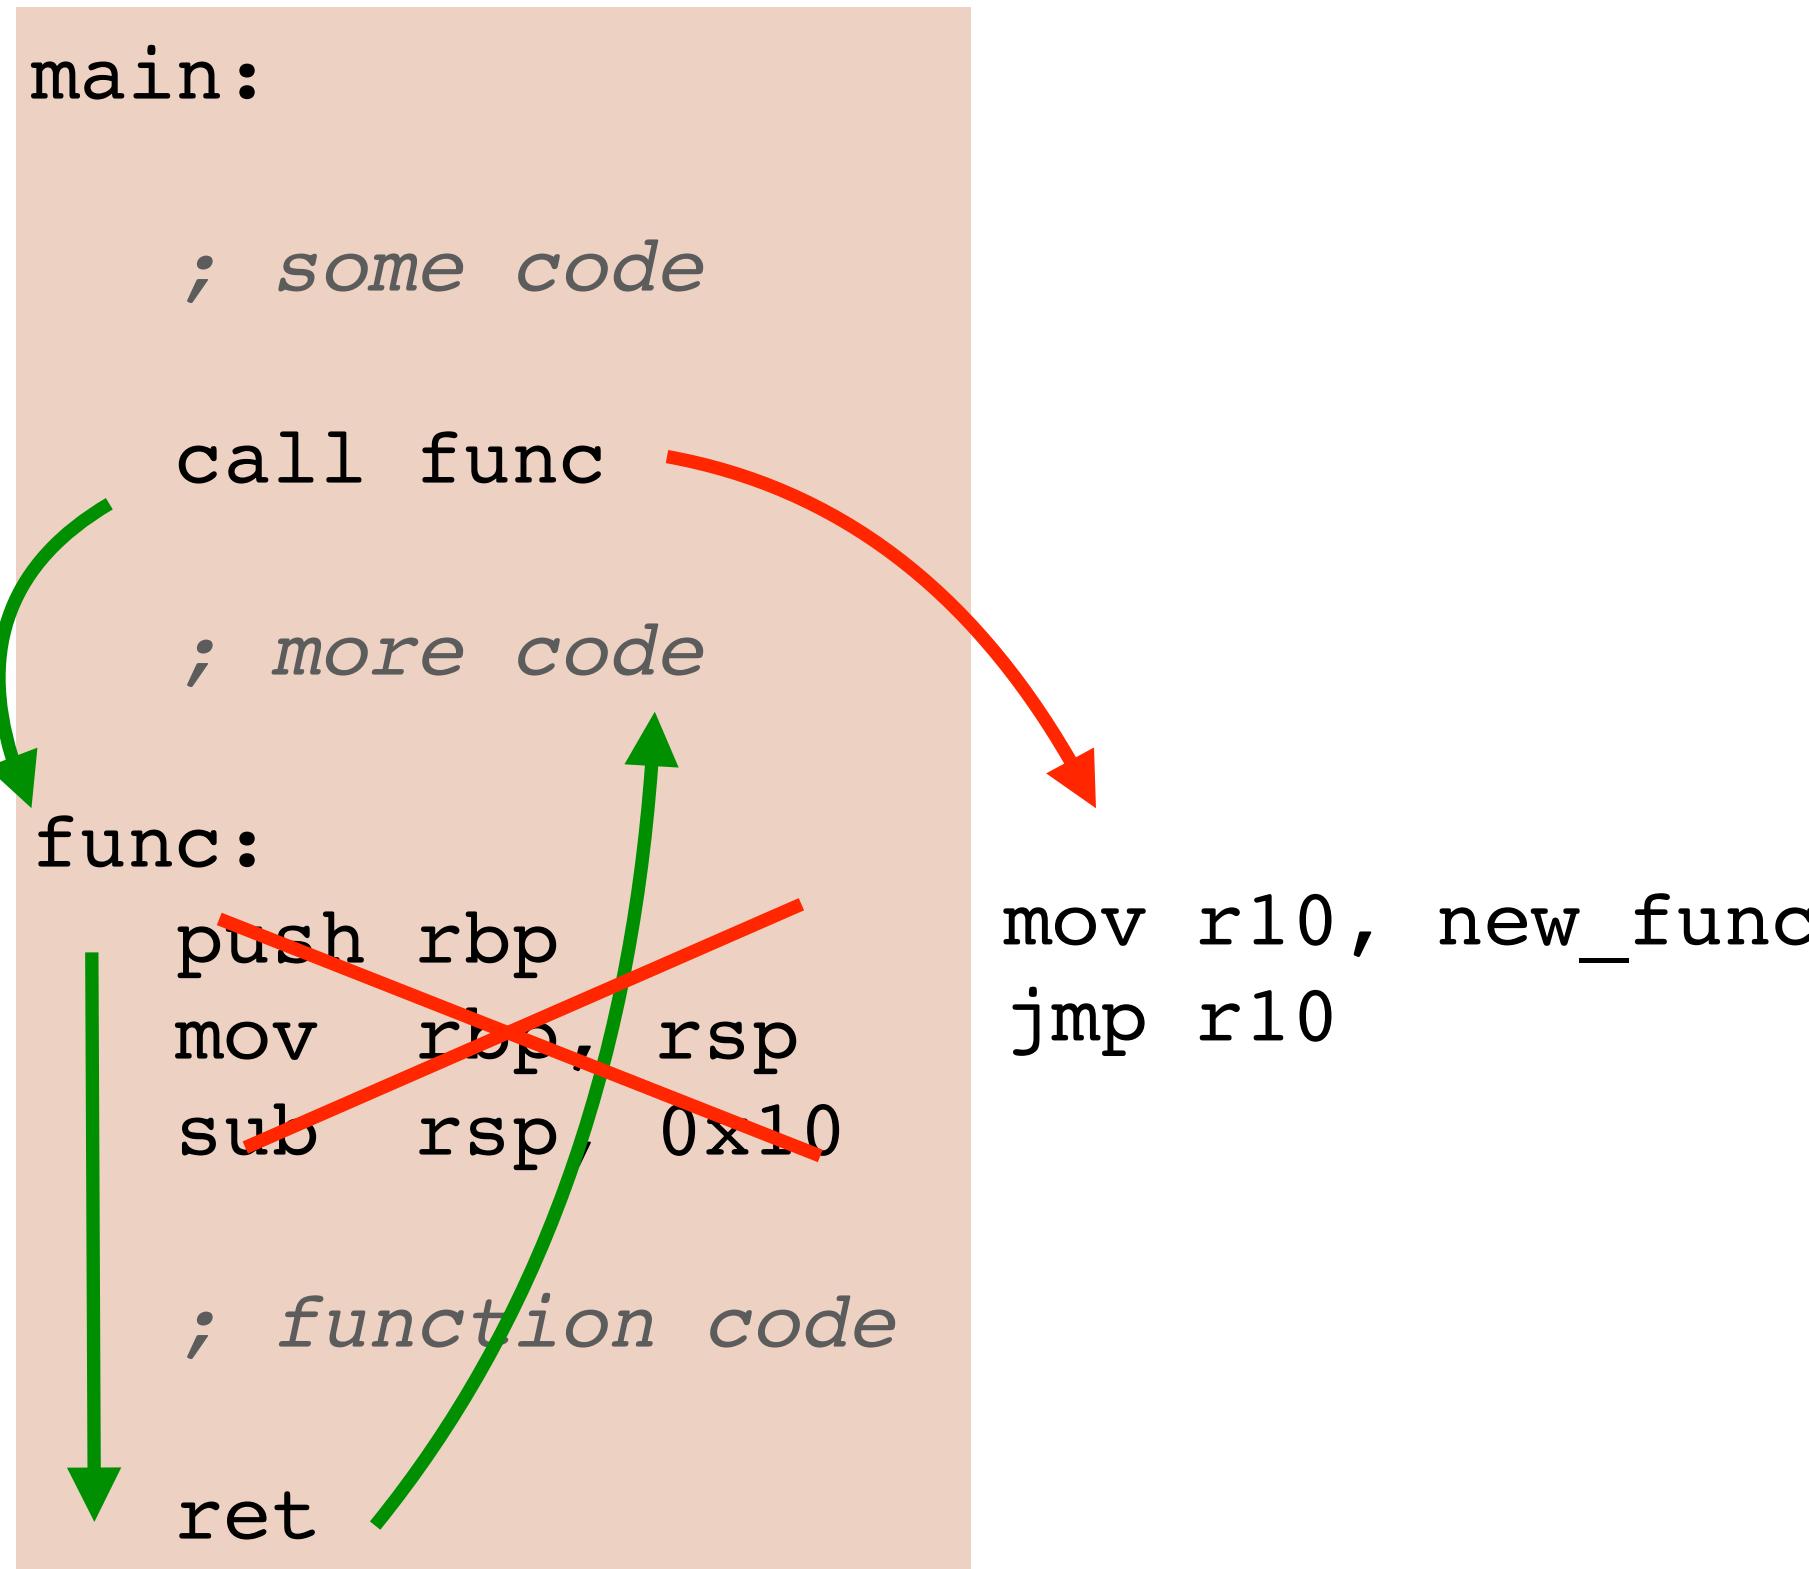
```
new_func:  
    ; replacement code
```

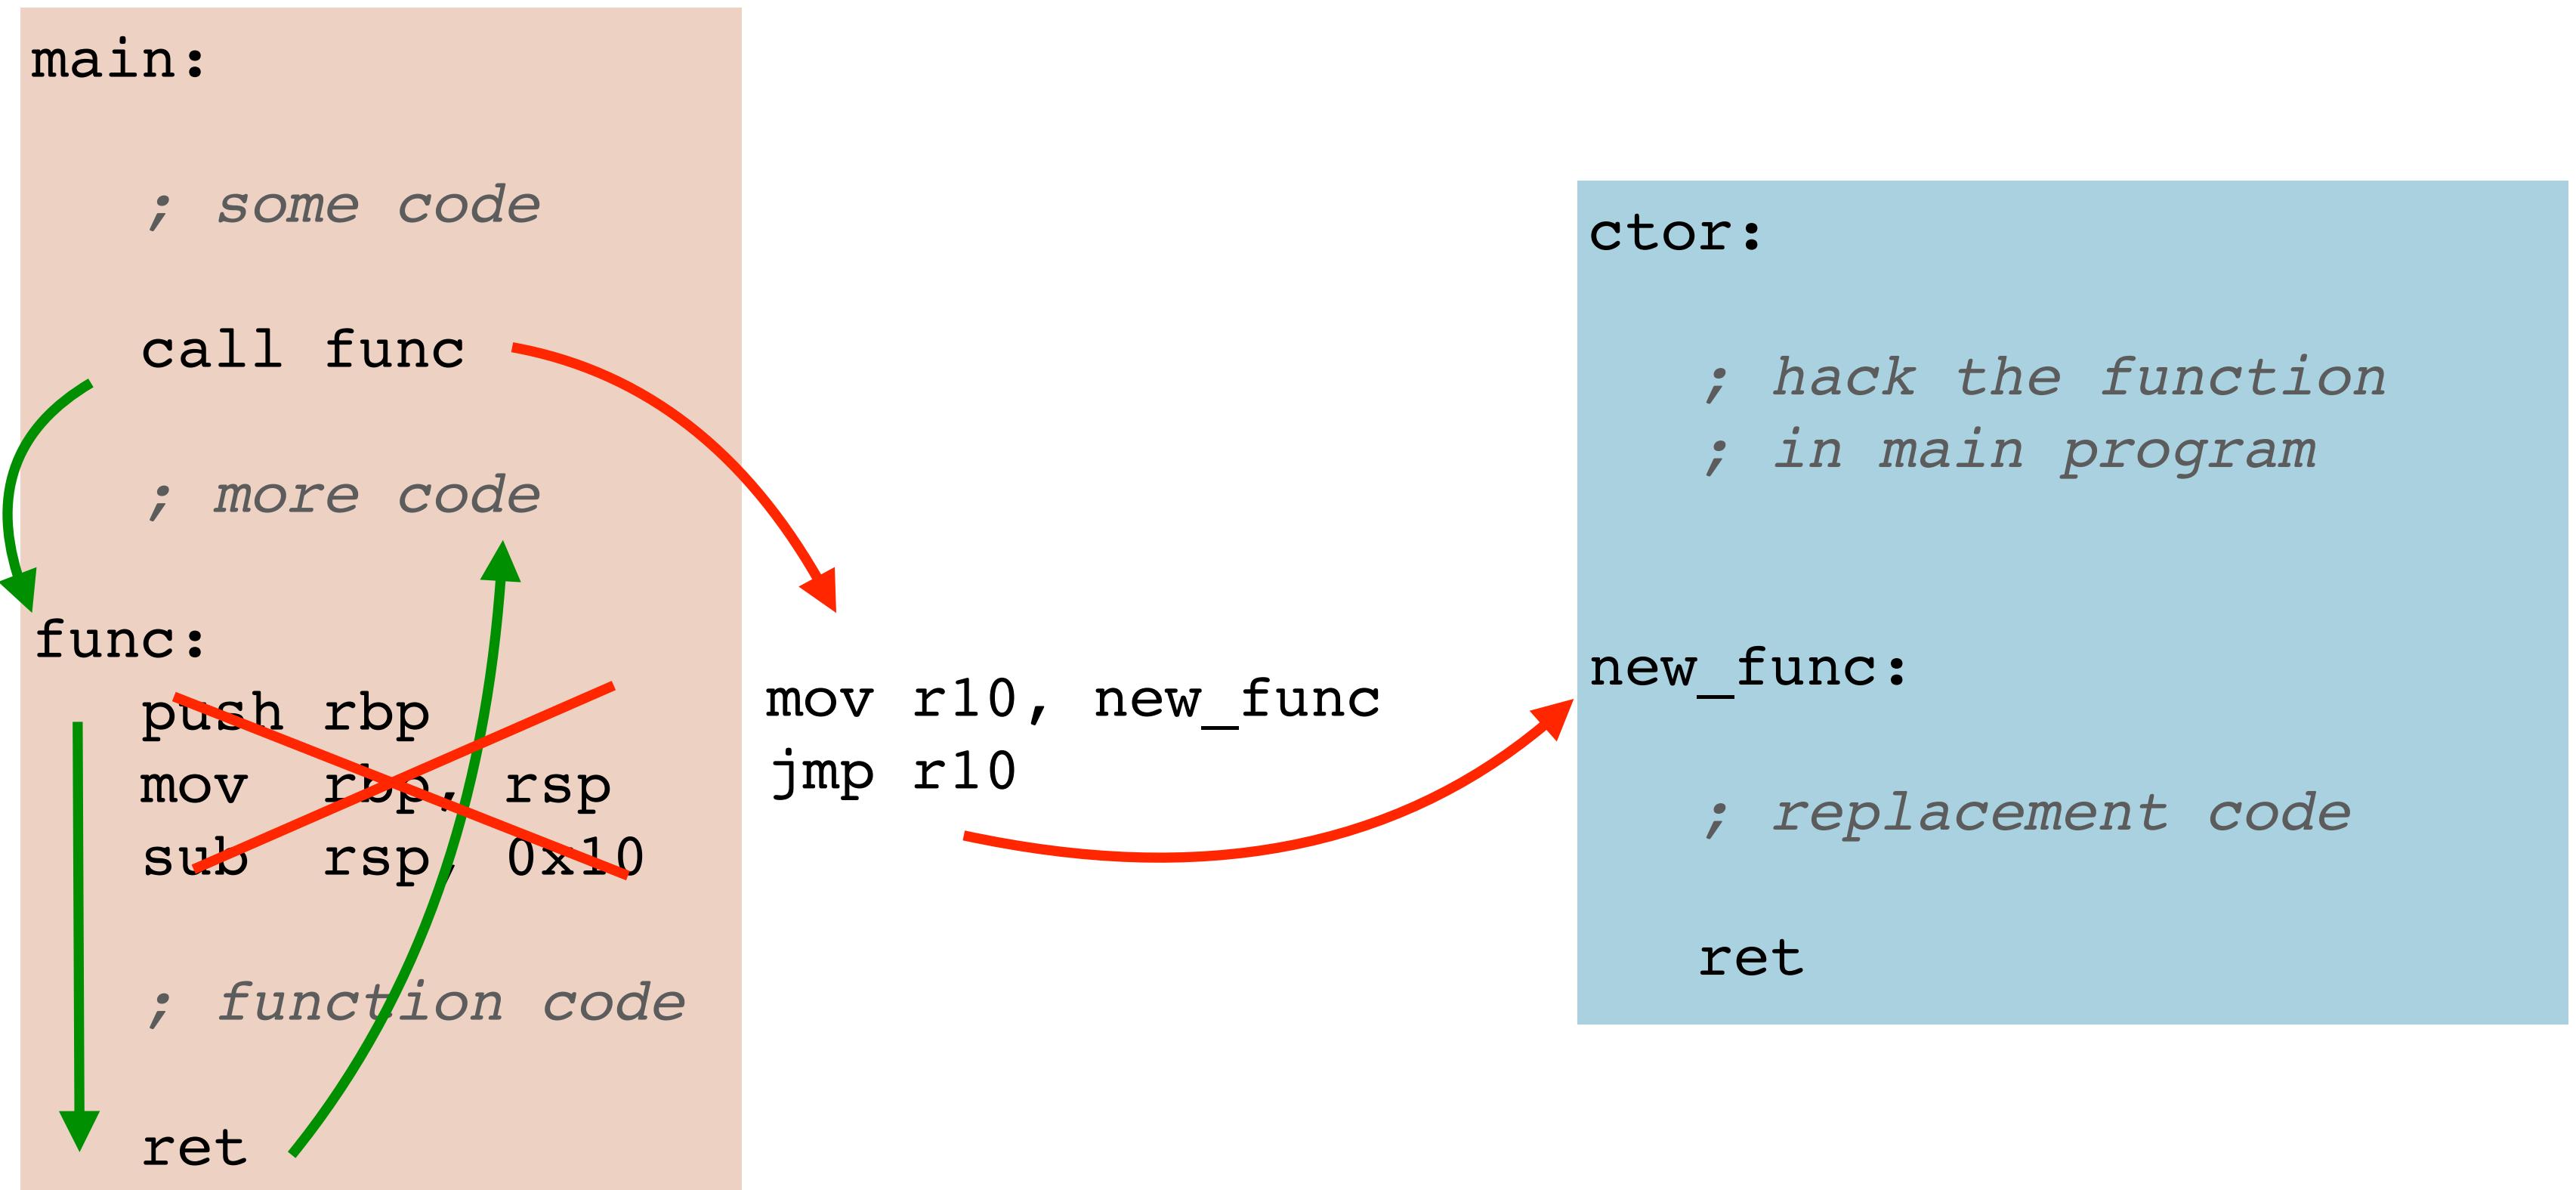
```
ret
```

# Static function hack: strategy

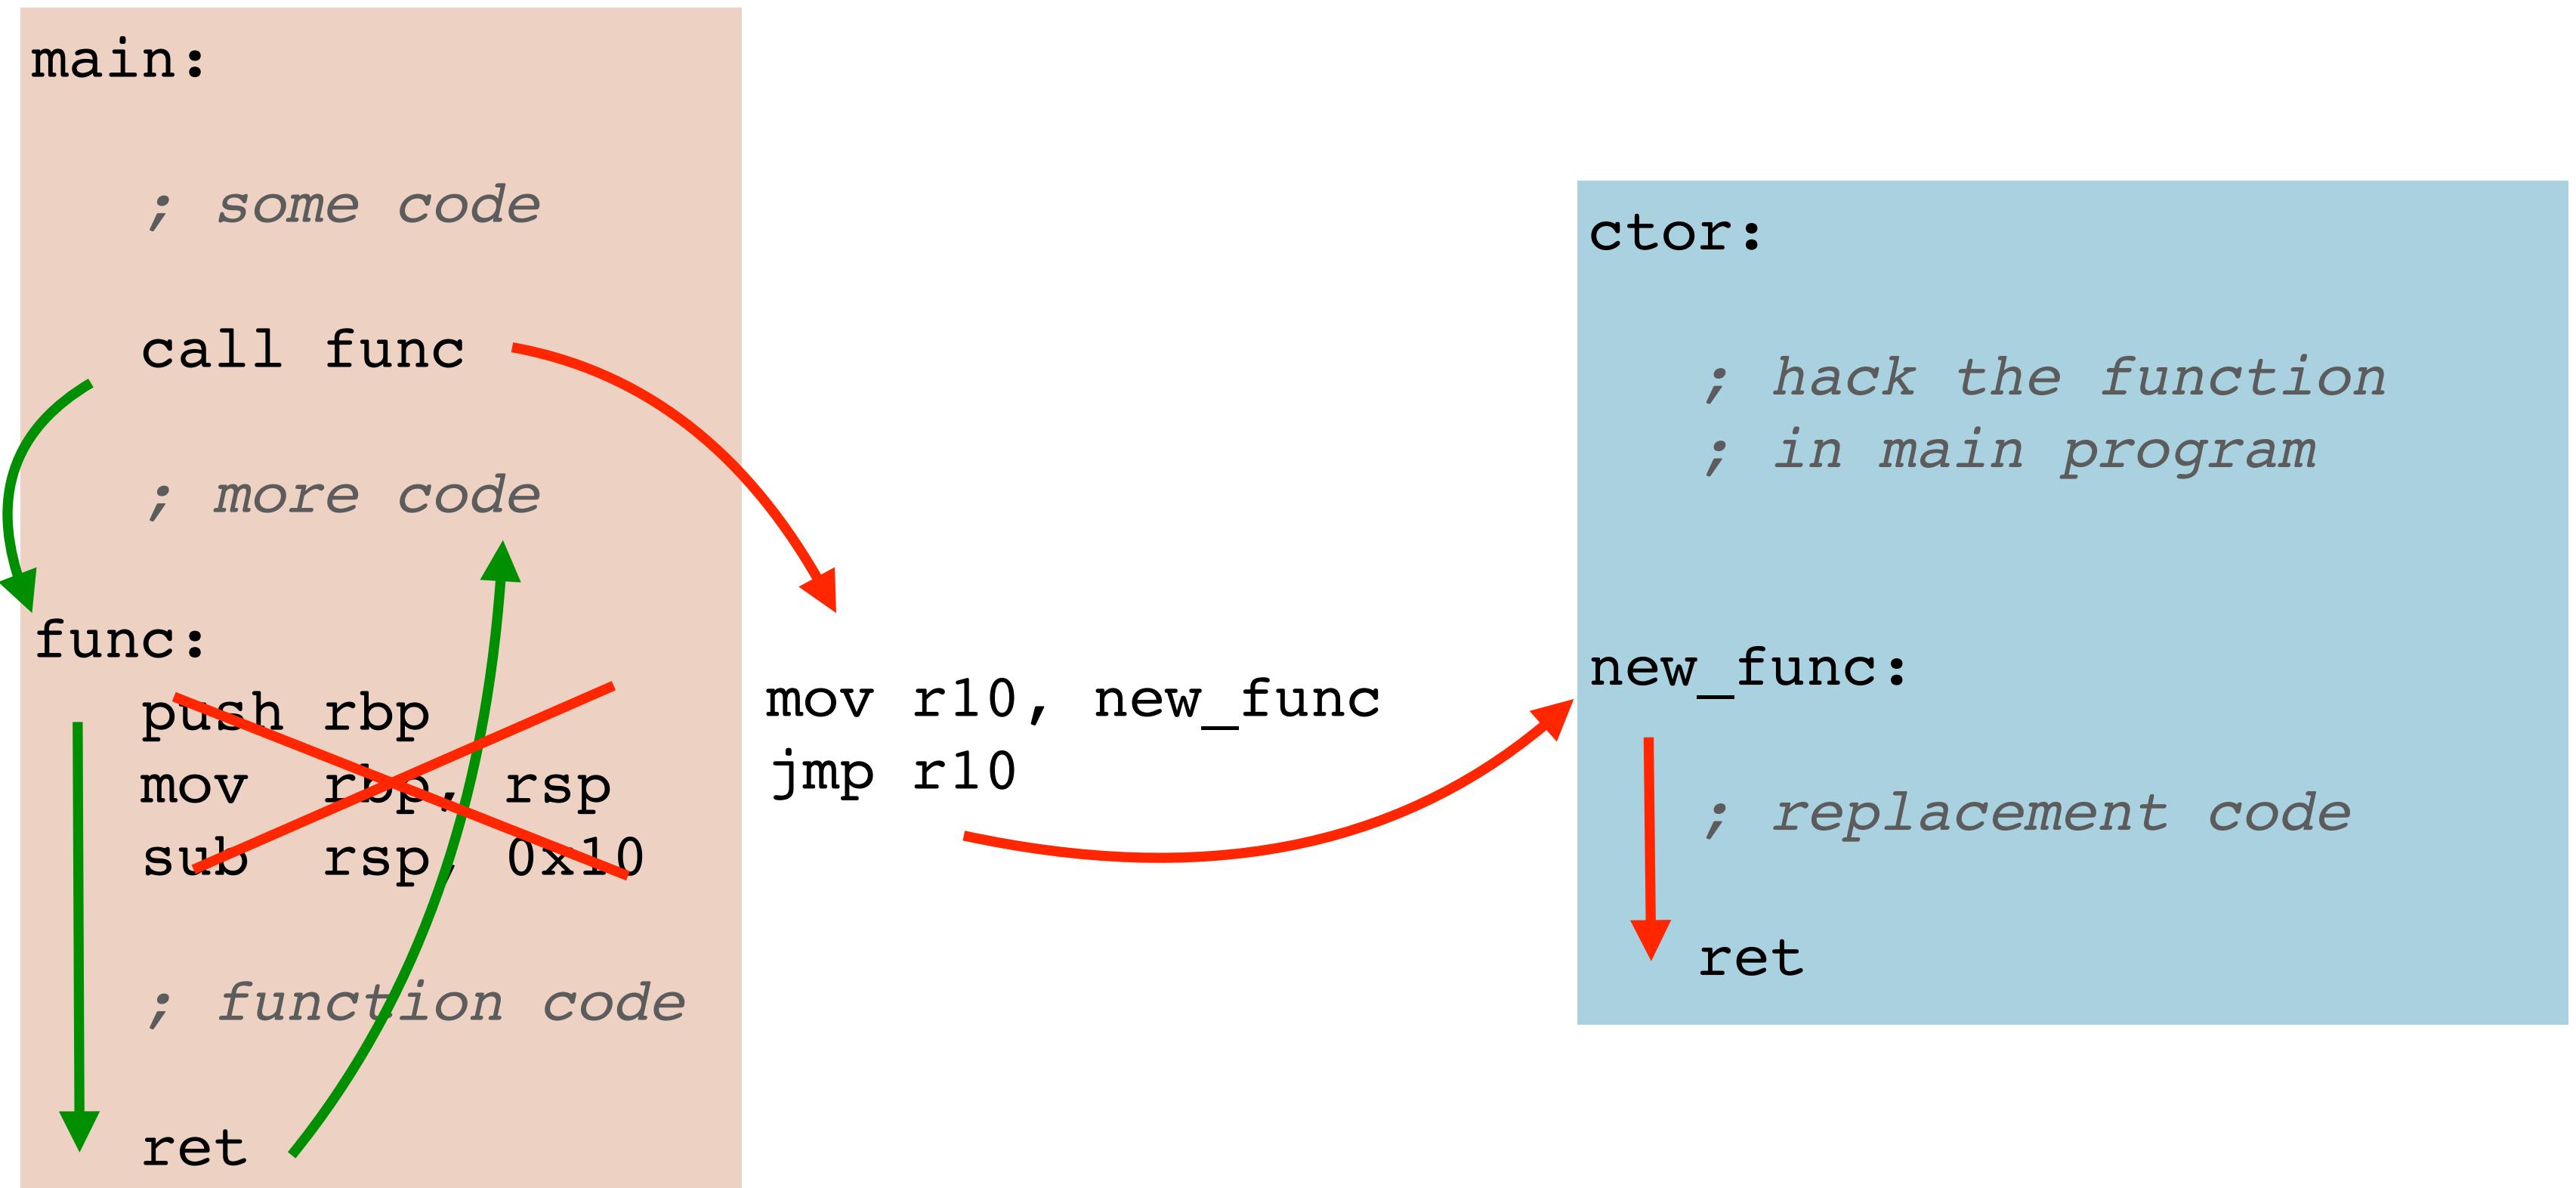
Modified execution



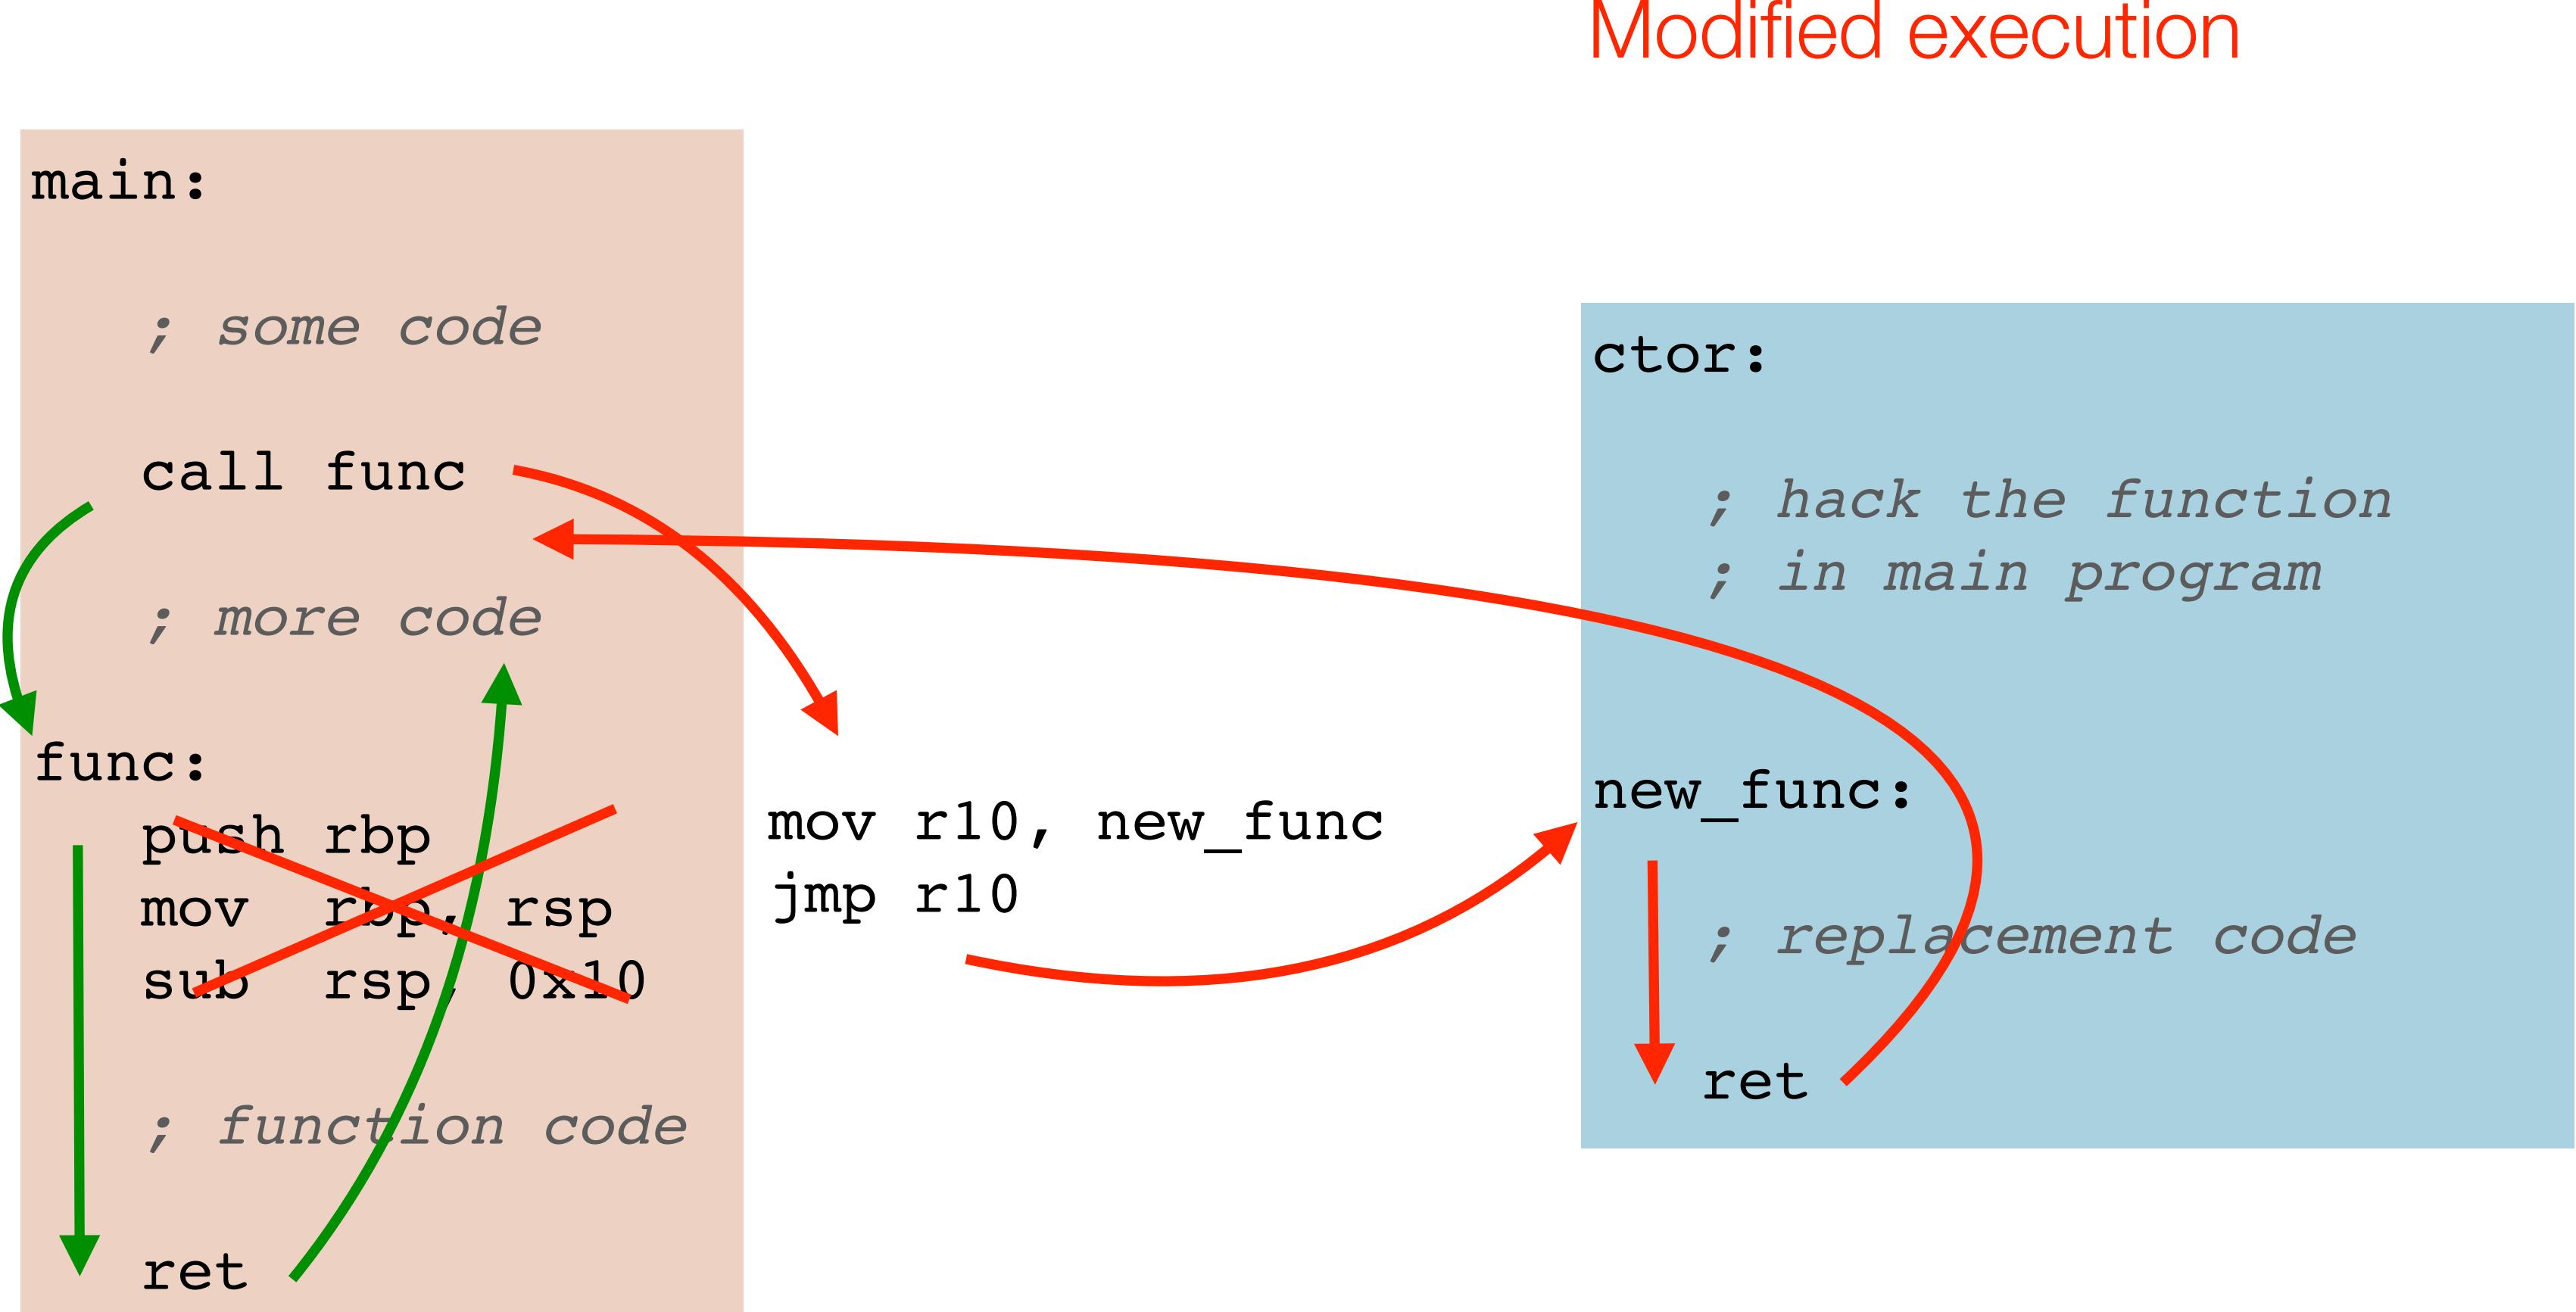
# Static function hack: strategy



# Static function hack: strategy



# Static function hack: strategy



# Static function hack: making the shellcode

---

- Assembly:

```
1 mov r10, 0aaaaaaaaaaaaaaaa ; address is just a 64 bits placeholder  
2 jmp r10
```

Address is a placeholder for the target function address, it will need to be changed at runtime

- Shellcode:

```
$ nasm -f elf64 jump.asm -o jump.o  
$ objcopy -O binary jump.o jump.bin  
$ hexdump -v -e '"\\\"x" 1/1 "%02x" "' jump.bin  
\x49\xba\xaa\xaa\xaa\xaa\xaa\xaa\xaa\xaa\x41\xff\xe2
```

13B long

# Static function hack: code modification

```
1 __attribute__((constructor)) static void ctor(void)
2 {
3     // get pointers to the original and new functions and calculate the jump offset
4     // only works if the main program exports its symbols! (cf. v2 for an alternative)
5     void *main_program_handle = dlopen(NULL, RTLD_NOW);
6     int64_t *target_func      = dlsym(main_program_handle, "func");
7     if (!target_func)
8     {
9         fprintf(stderr, "%s\n", dlerror());
10        exit(EXIT_FAILURE);
11    }
12    int64_t *new_func = (int64_t*)&replacement_func;
13
14    // skeleton for shellcode
15    unsigned char shellcode[] = "\x49\xba\xaa\xaa\xaa\xaa\xaa\xaa\xaa\xaa\x41\xff\xe2";
16    size_t code_size = sizeof(shellcode) - 1;
17
18    // copy address into shellcode and inject
19    memcpy(shellcode + 2, &new_func, sizeof(uint64_t));
20    memcpy(target_func, shellcode, code_size);
21 }
```

- The `constructor` attribute allows execution before `main`

# Static function hack: first attempt

- Ufffff.....

```
1 $ clang -g callfunc.c -o callfunc -ldl -Wl,--export-dynamic
2 $ clang -g -fPIC -shared funchookv0.c -o libfunchookv0.so -ldl
3 $ LD_PRELOAD=./libfunchookv0.so ./callfunc
4 [1]    97703 segmentation fault (core dumped) LD_PRELOAD=./libfunchookv0.so ./callfunc
```



# Static function hack: first attempt

- The debugger is, as usual, more helpful

```
$ lldb ./callfunc
(lldb) target create "./callfunc"
Current executable set to '/mnt/hgfs/binary-hack/lectures/4-hack/statichook/callfunc' (x86_64).
(lldb) env LD_PRELOAD=./libfunchookv0.so
(lldb) r
Process 94869 launched: '/mnt/hgfs/binary-hack/lectures/4-hack/statichook/callfunc' (x86_64)
Process 94869 stopped
* thread #1, name = 'callfunc', stop reason = signal SIGSEGV: address access protected (fault address: 0x401135)
  frame #0: 0x00007ffff7f426ff libc.so.6`__lldb_unnamed_symbol1092$$libc.so.6 + 143
 libc.so.6`__lldb_unnamed_symbol1092$$libc.so.6:
→ 0x7ffff7f426ff <+143>: mov    qword ptr [rdi + rdx - 0x8], rcx
  0x7ffff7f42704 <+148>: mov    qword ptr [rdi], rsi
  0x7ffff7f42707 <+151>: ret
  0x7ffff7f42708 <+152>: mov    ecx, dword ptr [rsi + rdx - 0x4]
(lldb) p/a 0x401135
(int) $0 = 0x00401135 callfunc`func + 5 at callfunc.c:5
```

- It actually makes sense, nowadays in essentially all Unix OSs the **text segment is read-only**
- But as the user running the program **you own the memory** and have the right to change that!

# Static function hack: writing on text segment

- Find the first page of the target's code and make it writable using **mprotect**

```
1 // make the memory containing the original function writable
2 // the address passed to mprotect must be aligned on page boundary
3 size_t    page_size  = sysconf(_SC_PAGESIZE);
4 uintptr_t start      = (uintptr_t)target_func;
5 uintptr_t end        = start + 1;
6 uintptr_t page_start = start & ~page_size;
7 int status = mprotect((void *)page_start, end - page_start, PROT_READ | PROT_WRITE | PROT_EXEC);
8 if (status)
9 {
10    fprintf(stderr, "Error unprotecting page\n");
11    fprintf(stderr, "Value of errno : %d\n", errno);
12    fprintf(stderr, "Error opening file: %s\n", strerror(errno));
13    exit(EXIT_FAILURE);
14 }
```

- Pages are generally at least 4kB so one page is more than enough for our 13B shellcode

# Static function hack

---

- It works!

```
$ clang -g callfunc.c -o callfunc -ldl -Wl,--export-dynamic
$ clang -g -fPIC -shared funchookv1.c -o libfunchookv1.so -ldl
$ LD_PRELOAD=./libfunchookv1.so ./callfunc
=====
INJECTION SUCCESSFUL
Injected code   : \x49\xba\xb0\xe1\x78\x61\x62\x7f\x0\x0\x41\xff\xe2
Target function : 0x401130
New function    : 0x7f626178e1b0
=====
calling variant with id 15431
```

- ... but there is a catch: did you notice the linker option **-Wl,--export-dynamic**?
- It exports the symbols of the main program for linking, this is the only way to get the address of the target with **dlsym**

# Static function hack: for real, i.e. without `dlsym`

---

- If the export option is dropped

```
$ clang -g callfunc.c -o callfunc -ldl
$ LD_PRELOAD=./libfunchookv1.so ./callfunc
./callfunc: undefined symbol: func
```

- But the symbol is of course here, it is just not visible to the dynamic linker anymore.
- The BFD library, part of [GNU binutils](#), can browse and ELF image and give a symbol's address
- Have a look at the v2!

# Limitations and extensions

---

- *What if the target is less than 13B?*

The injected code must be written in a new “relay” function, this can be done by adding executable memory with `mmap` (cf. “Hello world” shellcode example)

- *What if we just want to wrap around the target?*

The issue is that the jump overwrites code in the target. This can be solved by creating a “trampoline” function, which execute the replaced code and resume the execution of the target right after the injected jump.

# Is that safe?

---

- **Very**, all code injections example here rely on the fact that the program is unprivileged and belongs to the user
- Everything was done through calls to system functions
- Attack vectors using code injection are generally designed around **buffer overflow exploits**
- Remember the Fibonacci sequence going crazy through buffer overflow?
- Think about it, if you keep overflowing a stack buffer you will eventually **corrupt the return address** of the function...